

Autonomous Computing Systems

Neil Joseph Steiner

Dissertation submitted to the Faculty of the
Virginia Polytechnic Institute and State University
in partial fulfillment of the requirements for the degree of

Doctor of Philosophy
in
Electrical Engineering

Peter Athanas, Chair
Mark Jones
Cameron Patterson
Gary Brown
Djordje Minic

March 27, 2008
Bradley Department of Electrical and Computer Engineering
Blacksburg, Virginia

keywords: Autonomic, Hardware, Information, Interaction, FPGA

Copyright © 2008, Neil Joseph Steiner. All Rights Reserved.

Autonomous Computing Systems

Neil Joseph Steiner

Abstract

This work discusses *autonomous computing systems*, as implemented in hardware, and the properties required for such systems to function. Particular attention is placed on shifting the associated complexity into the systems themselves, and making them responsible for their own resources and operation. The resulting systems present simpler interfaces to their environments, and are able to respond to changes within themselves or their environments with little or no outside intervention. This work proposes a roadmap for the development of autonomous computing systems, and shows that their individual components can be implemented with present day technology.

This work further implements a proof-of-concept demonstration system that advances the state-of-the-art. The system detects activity on connected inputs, and responds to the conditions without external assistance. It works from mapped netlists, that it dynamically parses, places, routes, configures, connects, and implements within itself, at the finest granularity available, while continuing to run. The system also models itself and its resource usage, and keeps that model synchronized with the changes that it undergoes—a critical requirement for autonomous systems. Furthermore, because the system assumes responsibility for its resources, it is able to dynamically avoid resources that have been masked out, in a manner suitable for defect tolerance.

Xilinx graciously supported this research by granting the author exceptional access to confidential and proprietary information. The reader is advised that Xilinx software license agreements do not permit reverse engineering of the company's intellectual property in any form.

Acknowledgements

Thanks to Dr. Peter Athanas, my reason for joining Virginia Tech and the Configurable Computing Lab. It has been a privilege, and a tremendously rewarding experience. May Skynet never find me.

Thanks to Dr. Mark Jones, Dr. Cameron Patterson, Dr. Gary Brown, and Dr. Djordje Minic. Just as the strength of a conference rests on its program committee, the strength of a dissertation rests on its advisory committee.

Thanks to Mr. John Rocovich and the Bradley Foundation for funding three years of my research as a Bradley Fellow. This forward-looking work would have been impossible to undertake without the independent source of funding.

Thanks to Xilinx Research Labs for your tremendous support of both this work and prior related work. I am deeply indebted and very grateful, and I remain a true believer in your products.

Thanks to Carlos Lucasius and Sun Microsystems, Inc., for allowing me to use Java SE for Embedded (PowerPC) beyond the normal trial period.

Thanks to Jeff Lindholm at Xilinx, Inc., who passed away last year. Jeff was one of the good guys.

Thanks to my past mentors: Maya Gokhale, Justin Tripp, Brandon Blodget, and James Anderson.

Thanks to Matt Shelburne and Tim Dunham for your assistance in helping me debug bitstream and EDK issues. The issues were small, but the time saved was enormous.

Thanks to my many CCM Lab friends, and Stephen Craven in particular. It has been a real pleasure working with so many of you, and I will remember this place with many fond memories.

Thanks to my many other friends: Eric, Ben, Tullio, Rick, J.D., James G., Amy, Jeremy, Stella, Tim, Phil, Lauren and Lauren, Carla, Alex, Valérie, and more.

Thanks to my parents, who somehow find the time to worry about when I'll get married; to Bern and Camaleta, Stephanie, Eric, and Luke; and to Norine and Andrew, Jordan, Jackie, and Blaise.

Et un très grand merci aussi à Johannes, à Véronique, et à Lili pour votre amitié. Peut-être trouverai-je enfin le temps de rétablir notre contact.

Preface

0.1 Overview

The long-term underlying question that prompted this work concerns the very nature of computation [1]: How exactly do hardware and information interact? Information is both the fundamental ingredient of computation and its end product, and without information, computation would serve no purpose. On the other hand, information is incapable of doing anything without an underlying hardware substrate.

Computation is so pervasive in modern society that one might easily assume all of the foundations to be well understood, but while we very well understand how to *use* computation, we are harder pressed to explain what it *is*. That hardware and information do interact is well established however, and that interaction can be observed in both microscopic¹ and macroscopic behaviors.

0.2 Microscopic Behavior

The microscopic behavior of hardware and information is governed by the fundamental nature of matter, energy, and information, and by the interactions permitted by the laws of physics [2]. Understanding of the underlying behavior does remain a long term interest of this work, but at present the fundamental questions still lie beyond the reach of physics.

According to the quantum mechanical formalism, the wavefunction ψ of a system encodes all of the information about the physical state of that system,² so there is at least some known relationship between information and the wavefunction. If one can then express a relationship between matter or energy and the wavefunction, it should be possible by transitivity to derive a relationship between matter, energy, and information: Perhaps some analogue or extension to Einstein's $E = mc^2$ which

¹The term *microscopic* is relaxed here to simply refer to fundamental behavior, with no assumption that this behavior be directly observable with a microscope.

²Some physicists are careful to speak of *physical* information [3], perhaps tacitly allowing that information in the wavefunction might not be the most fundamental form of information.

also encompasses information. Such a relationship would effectively represent governing physical laws, and would provide a rigorous foundation for research into the nature of computation.

Unfortunately, these seemingly simple questions about relationships are in fact tremendously difficult to address. The reported successes in quantum mechanics, quantum cryptographic key exchange, and quantum teleportation, and the publicized search for answers in cosmology, string theory, and the Theory of Everything, can convince the lay person that fundamental understanding is within reach, when there is in fact an enormous amount of work yet to be done:

- Physicists are not yet agreed on the nature of information. The Landauer view of information as physical [4], suggests that information is actually a form of matter or energy. The Wheeler view of information as the origin of matter and energy [5], suggests that information is a superset of matter and energy. In both of these cases, it is reasonable to ask about the constituent relationships, but nobody seems to be investigating these questions at present.
- Physicists are also grappling with the fundamental nature of matter and energy. Quantum mechanics explains how particles behave as waves, and quantum field theory explains how waves behave as particles, but nobody is sure how matter and energy relate to the wavefunction. The feeling is that string theory may suggest some geometric origin for degrees of freedom in matter, but “at the moment nobody understands what string theory is” [6].

The present inability of physics to provide answers reflects the deep complexity of nature and the difficulty of unraveling its mysteries, as well as the difficulty of formulating questions in a manner conducive to answers [7]. While the most efficient way of studying hardware-information dynamics would probably come from such answers, those answers are not available at this time.

0.3 Macroscopic Behavior

The macroscopic behavior of hardware and information governs larger scale interactions in computational circuits and systems, and is built upon the microscopic behavior, though the relationships are presently obscured from us. On this level, it is interesting to classify systems according to the dynamics of their state spaces. One may ask whether a system’s state space can change, and if so, how? The underlying argument is that the dynamics of the state space have significant implications for the flexibility and capabilities of the system.

This work implicitly acknowledges that the dynamics of the system may include hardware changes, whether actual or virtual. At present it is possible to achieve virtual hardware changes in configurable devices such as FPGAs. Actual hardware changes will likely become possible through continued research in nanotechnology and molecular self-assembly [8]. The meaning of the distinctions between virtual and actual changes will be presented in more detail in later chapters of this work.

Does it make sense to study macroscopic behavior if one is actually trying to understand microscopic fundamentals? Perhaps so. In physics, it is common to study gravitational lenses, supernovae, black holes, and cosmology—all of which are interesting in their own right—while seeking to understand exceedingly small particles and phenomena. Certain weak physical interactions only become observable at high energies, and to study them it is often necessary to use cosmological sources or large accelerators. This work pragmatically posits that any approach is valid if it can help shed light on the underlying fabric of nature, particularly when other means are not available.

Another reason for studying large systems is that they tend to exhibit emergent behaviors that are difficult or impossible to infer from fundamentals. Biology has been acquainted with Avogadro-scale systems for at least as long as recorded history, and as computation progresses towards such scales, there is an increasing effort to glean ideas and approaches from living systems. Efforts are currently underway by Ganek and Corbi at IBM to develop *autonomic systems* that are self-aware, self-configuring, self-optimizing, self-healing, and self-protecting [9], and by the DARPA Information Processing Technology Office (IPTO) to develop *cognitive systems* that are able to reason, able to learn, able to communicate naturally, able to respond to unexpected situations, and able to reflect on their own behavior [10].

While this work does not focus on artificial intelligence, it does consider the infrastructure that might be necessary to support autonomy on large computational systems. In anticipation of Avogadro-scale computing, it is sensible to ask how such systems might be designed, and it is worthwhile to learn from biological systems, which do in fact grow, learn, adapt, and reason.

Table of Contents

Preface	v
0.1 Overview	v
0.2 Microscopic Behavior	v
0.3 Macroscopic Behavior	vi
Table of Contents	vii
List of Figures	xv
List of Tables	xvii
Glossary	xx
1 Introduction	1
1.1 Overview	1
1.2 Autonomy	1
1.3 Objectives	3
1.3.1 Roadmap	3
1.3.2 Implementation	3
1.4 Organization	4

2 Background	5
2.1 Overview	5
2.2 Digital Design Implementation	5
2.3 Virtex-II Pro Architecture	12
2.3.1 User and Device Models	12
2.3.2 Configuration and State Planes	15
2.3.3 State Corruption Considerations	15
2.3.3.1 LUT RAM	16
2.3.3.2 SRL16	16
2.3.3.3 BRAM	16
2.3.4 Unsupported Slice Logic	17
2.4 XUP Board	17
2.5 ADB	19
2.5.1 Routing	20
2.5.2 Unrouting	20
2.5.3 Tracing	21
2.6 Device API	21
2.7 Linux on FPGAs	22
2.7.1 University of Washington: EMPART Project	23
2.7.2 Brigham Young University: Linux on FPGA Project	23
2.7.3 University of York: DEMOS Project	23
2.8 Related Work	24
2.8.1 Autonomy	24
2.8.2 Placers and Routers	25

2.8.3 JHDLBits	26
3 Theory	27
3.1 Overview	27
3.2 Dynamics	27
3.3 Configurability	28
3.4 Scalability	30
3.5 Modeling	32
3.6 Domain	34
3.7 Observations	34
4 Roadmap	36
4.1 Overview	36
4.2 Levels of Autonomy	38
4.2.1 Level 0: Instantiating a System	38
4.2.2 Level 1: Instantiating Blocks	38
4.2.3 Level 2: Implementing Circuits	39
4.2.4 Level 3: Implementing Behavior	40
4.2.5 Level 4: Observing Conditions	40
4.2.6 Level 5: Considering Responses	41
4.2.7 Level 6: Implementing Responses	42
4.2.8 Level 7: Extending Functionality	42
4.2.9 Level 8: Learning From Experience	43
4.3 Issues	43
4.3.1 Policies	43

4.3.2 Security	44
4.3.3 Testing	44
Implementation	45
5 Hardware	47
5.1 Overview	47
5.2 Board Components	49
5.3 FPGA Components	49
5.4 Internal Communication	51
5.5 Framebuffer	51
5.6 Sensor Stubs	55
5.6.1 PS/2 Stub	55
5.6.2 Video Stub	56
5.6.3 Audio Stub	56
5.6.4 Unused Stubs	57
5.7 Interface Stubs	57
5.8 Custom Dynamic Cores	58
5.8.1 Video Capture Core	58
5.8.2 FFT Core	59
5.8.3 Other cores	59
5.9 Floorplanning	59
5.10 Summary	61
6 Software	62
6.1 Overview	62

6.2 Linux	63
6.2.1 crosstool	63
6.2.2 mkrootfs	64
6.2.3 Kernel 2.4.26	64
6.2.4 Packaging the Kernel	65
6.2.5 BusyBox	65
6.2.6 Other Utilities	65
6.2.7 Filesystem	66
6.3 Custom Drivers	66
6.3.1 AC '97 Driver: /dev/ac97	67
6.3.2 Framebuffer Driver: /dev/fb	67
6.3.3 FPGA Driver: /dev/fpga	68
6.3.4 GPIO Driver: /dev/gpio/*	69
6.3.5 SystemStub Driver: /dev/hwstub	69
6.3.6 ICAP Driver: /dev/icap	69
6.3.6.1 Errors in Low-Level hwicap_v1_00_a Driver	69
6.3.6.2 CRC Calculation Errors	70
6.3.7 PS/2 Driver: /dev/ps2/*	71
6.3.8 Stub Driver: /dev/stub	71
6.3.9 Video Driver: /dev/video	71
6.4 Java	72
6.4.1 Unusable Java Runtime Environments	72
6.4.2 Java SE for Embedded (PowerPC)	73

7 System	74
7.1 Overview	74
7.2 Capabilities	75
7.2.1 Protocol Specification	75
7.2.2 Architecture Support	77
7.2.3 Library Support	79
7.3 Autonomous Controller	80
7.4 Autonomous Server	82
7.4.1 Circuits	86
7.4.2 Protocol Handler	86
7.4.3 Cell Semantics	87
7.4.4 Model	87
7.4.5 Libraries	88
7.4.6 EDIF Parser	89
7.4.7 Placer	90
7.4.7.1 Abandoned Algorithm: Simulated Annealing	91
7.4.7.2 Resource Map and Filtering	93
7.4.7.3 Constraints	94
7.4.7.4 Inferring Cell Groups	94
7.4.7.5 Placer Algorithm	95
7.4.8 Mapper	97
7.4.9 Router	99
7.4.10 Configuration	100
7.4.11 Connections	101

7.4.12 XDL Export	102
8 Operation	104
8.1 Overview	104
8.2 Preparation	106
8.3 Startup	107
8.3.1 Linux	107
8.3.2 Drivers	107
8.3.3 Server	108
8.3.4 Controller	109
8.4 Operation	109
8.4.1 Interaction	109
8.4.2 Events	110
8.4.3 Robustness	110
8.4.4 Time-Limited IP	111
8.5 Validation	111
8.5.1 Level 0: Instantiating a System	112
8.5.2 Level 1: Instantiating Blocks	112
8.5.3 Level 2: Implementing Circuits	112
8.5.4 Level 3: Implementing Behavior	112
8.5.5 Level 4: Observing Conditions	112
8.5.6 Level 5: Considering Responses	112
8.5.7 Level 6: Implementing Responses	113
8.5.8 Actual Demonstration	113

8.6 Analysis	113
8.6.1 XDL Export	114
8.6.2 Performance	114
8.6.3 Inspection	117
8.6.4 Overhead	118
9 Conclusion	122
9.1 Claims	122
9.2 Narrative	122
9.3 Roadmap	123
9.4 Implementation	124
9.5 Evaluation	125
9.6 Reflections	126
9.7 Future Work	127
9.8 Conclusion	129
A Slice Packing	130
A.1 Introduction	130
A.2 Rules	130
B Design Proposal	135
B.0 Preamble	135
B.1 Overview	135
B.2 Component Availability	135
B.3 Proposal	137
B.4 Platform Selection	138

B.5 Tasks	139
B.5.1 Incremental Placer	140
B.5.2 Incremental Router	140
B.5.3 Virtex-II Pro Bitstream Support	141
B.5.4 System Model	141
B.5.5 Platform Support	142
B.5.6 Complete System	143
B.6 Optional Tasks	143
Bibliography	145

List of Figures

2.1	Configurable logic example: Virtex-II slice.	7
2.2	8-bit adder: behavioral VHDL.	8
2.3	8-bit adder: technology-independent schematic (with I/O ports).	8
2.4	8-bit adder: technology-mapped netlist.	10
2.5	8-bit adder: placed but unrouted design	11
2.6	8-bit adder: fully routed design	11
2.7	Virtex-II slice subcells	14
2.8	Virtex-II unsupported slice logic	18
2.9	XUP Virtex-II Pro development system.	19
4.1	Levels of autonomy.	37
5.1	Demonstration system: high-level view.	48
5.2	Framebuffer schematic (1 of 2).	52
5.3	Framebuffer schematic (2 of 2).	53
5.4	Demonstration system floorplan	60
6.1	Software block diagram.	63
7.1	Server block diagram.	75

7.2	Autonomous controller screen capture.	81
7.3	Resource usage at startup.	83
7.4	Resource usage after interactive resource masking.	84
7.5	Resource usage after placement of video capture and FFT cores.	85
7.6	System model components.	88
7.7	Simulated annealing algorithm.	92
7.8	Server placer algorithm.	96
7.9	Top-level net.	101
8.1	Startup and operation.	105
8.2	Configured slice sample (XDL view).	118
8.3	Configured slice sample (FPGA Editor view).	119
8.4	Placed and routed sample.	120

List of Tables

5.1	SXGA video timing parameters	54
8.1	Dynamic test circuits	115
8.2	Dynamic circuit performance comparison	115
8.3	Embedded dynamic circuit performance	116
8.4	Embedded dynamic circuit comparative analysis	117
B.1	Platform considerations	139

Glossary

ADB: A wire database for Xilinx Virtex/-E/-II/-II Pro FPGAs. ADB can be used to route, unroute, or trace designs.

API: Common abbreviation for Application Programming Interface. A collection of functions and/or data which can be used to access underlying functionality in a standard way.

Arc: A connection between two wires or *nodes*. Borrowed from graph theory.

ASIC: Application-Specific Integrated Circuit. An integrated circuit fabricated to perform a specific function. ASICs typically offers the best size, power, and speed characteristics, but cannot be modified once they have been fabricated, and are generally very expensive to fabricate.

Bitgen: The configuration bitstream generator provided by Xilinx for its configurable devices.

Bitstream: Generally refers to a *configuration bitstream*. The configuration data in a form suitable for download into a configurable device. Before a configuration bitstream is applied to a configurable device, the device is effectively blank, and performs no function at all.

Bitstream Generator: Generally refers to a *configuration bitstream generator*. A tool that converts a fully placed and routed design into a configuration bitstream suitable for download into a configurable device.

Cell: A block of logic. Generally refers to a block of configurable logic in an FPGA, but can also refer to a logic cell in an EDIF netlist, or to a logic cell in a user design.

CAD: Computer-Aided Design. Applied to software tools that process or transform design information. In this context, the term includes synthesizers, mappers, placers, routers, bitstream generators, and more.

Cell Matrix: A configurable architecture with very fine granularity and massively parallel reconfigurability. Simulation only—not presently available in silicon.

Clock: A periodic signal used to synchronize logic.

CMOS: Complementary Metal-Oxide-Semiconductor. The prevalent integrated circuit technology in use. Provides very good noise immunity and very low static power dissipation.

EDIF: Electronic Design Interchange Format. A netlist format with widespread industry acceptance.

EDK: Embedded Development Kit. The Xilinx tool set used to generate processor-based designs that can be embedded inside FPGAs. The EDK provides support for both PowerPC and MicroBlaze processors, as well as PLB, OPB, LMB, and other busses, and a large number of IP cores.

Flip-Flop: A sequential device that can retain state information. A single flip-flop can implement a single bit of memory in a design. Flip-flops are generally controlled by a clock signal.

FPGA: Field-Programmable Gate Array. A programmable logic device containing an array of logic elements that can be configured in the field. FPGAs were originally used primarily for glue logic and for prototyping, but have become quite common in final products. One interesting property that FPGAs may have is the ability to be partially configured, from the inside, an unlimited number of times.

HDL (Hardware Description Language): A high-level textual language that can be used to describe and synthesize hardware. Common examples are VHDL and Verilog.

IP: Intellectual Property. Used to refer to logic cores developed by vendors. These cores are circuits that can be implemented in configurable logic.

ISE: Integrated Software Environment. The Xilinx data and tool set used to generate configuration bitstreams for their FPGAs.

JRE: Java Runtime Environment. The JVM and API necessary to run Java programs.

JVM: Java Virtual Machine. The virtual machine used for execution of compiled Java programs.

LUT: Lookup-Table. A common configurable circuit that selects an output based on the values of its inputs. A LUT can be configured to perform an arbitrary operation on its inputs. LUTs are used extensively in FPGAs, and a 2-input LUT can implement any of 16 functions on its input data, including AND, OR, NOT, NAND, NOR, XOR, XNOR, 0, 1, and so on.

Makefile: A file readable by the *make* utility, describing targets and rules that facilitate large build processes. Makefiles are frequently used for software builds, but can just as readily be used for hardware builds or other kinds of processes.

Mapper: A tool that maps generic logic gates to the specific logic primitives available on the target architecture. In general, this includes mapping of AND, OR, and NOT gates into primitives like LUTs and flip-flops.

Net: An electrical node. In an ASIC a net is a single physical node. In a configurable device, a net may also be a virtual node, formed by virtual connections between independent electrical nodes.

Netlist: A circuit described in terms of logic components and the connections between them. Permits interchange between dissimilar tools. EDIF is the predominant netlist format.

NoC: Network-on-Chip. A coarse-grained communication method between blocks within the same chip.

Node: A wire. Borrowed from graph theory.

OPB: On-Chip Peripheral Bus. An IBM CoreConnect bus designed to connect peripherals to each other and to the PLB bus. *Can instead connect to the LMB bus in other architectures, but the LMB is not described here.*

Passthrough: A routing connection that *passes through* a logic cell, taking advantage of unused logic to increase routability.

Placer: A tool that allocates mapped primitives to specific resource instances within the target device. Before placement, a design can be considered to target a particular architecture, but after placement, the design targets a specific device.

PLB: Processor Local Bus. An IBM CoreConnect bus designed to connect the PowerPC processor to other local devices such as memory controllers and OPB bridges.

PLD: Programmable Logic Device. An integrated circuit whose logic can be configured to perform some function. See FPGA.

Router: A tool that determines how to make the necessary connections between all of the placed primitives in a target device. Routing is a difficult problem, and is typically the lengthiest part of the design flow. In particular, it can be very difficult to route millions of nets and still ensure that each one meets timing requirements.

SRAM: Static RAM. As opposed to DRAM or Dynamic RAM. SRAM is much faster than DRAM, but cannot be packed as densely, and is therefore more expensive and difficult to make in large sizes.

Standard Cell: A predefined logic block used in integrated circuit design. Larger and less efficient than a *full custom* equivalent, but carefully designed and analyzed, and easy to instantiate.

Synthesizer: A tool that converts hierarchical and behavioral HDL code into an implementation consisting of logic gates.

Tool Flow: A set of design tools that consecutively process parts of a design from some original format into the desired target format. In the context of FPGAs, the tool flow consists primarily of synthesis, mapping, placing, routing, and configuration bitstream generation.

VHDL (VHSIC Hardware Description Language): A formal and complex HDL that includes powerful simulation and abstraction capabilities. Similar in feel to ADA. c.f. Verilog.

Verilog: A lightweight HDL that is less cumbersome than VHDL. Similar in feel to C. c.f. VHDL.

VLSI: Very Large Scale Integration. Refers to very high density integrated circuits.

Xilinx: Leading manufacturer of programmable logic devices, and FPGAs in particular. The newly introduced flagship family is the Virtex-5, followed by the Virtex-4, Virtex-II Pro, Virtex-II, Virtex-E, and Virtex.

XUP: Xilinx University Program. Used to refer to the Xilinx University Program Virtex-II Pro development board, manufactured by Digilent, Inc.

Chapter 1

Introduction

1.1 Overview

Autonomy in living systems is the ability to act with some measure of independence, and to assume responsibility for one's own resources and behavior. This in turn makes it possible for systems to function in the absence of centralized external control. While these capabilities are useful and well established in living systems, they also turn out to be desirable in complex computing systems, for reasons that will be examined in further detail in this work. This introduction takes a closer look at autonomy, at the objectives of the work, and at the organization of the remaining chapters.

1.2 Autonomy

One may informally describe an *autonomous computing system* [11] as a system that functions with a large degree of independence, much as a living system does. This requires that the system assume responsibility for its own resources and operation, and that it detect and respond to conditions that may arise within itself or its environment.

The programmable machines of today are quite flexible, but they are generally also dependent on specific programming or configuration from external sources. By comparison, an autonomous system would be able to function with more abstract or generic programming, and might in some cases be able to infer and apply the necessary programming on its own, without external control.

Autonomy is desirable in part because internal and external conditions do change in real systems, and when that occurs, it is desirable for the system to adjust gracefully. Failure to adjust not only limits the usefulness of the system, but also accelerates its obsolescence. The existence of software operating systems means that software updates can usually be applied without rebuilding

an entire system, but that kind of convenience is only now emerging in the hardware domain [12]. Furthermore, even in the absence of changes, two identical systems may find themselves in dissimilar environments, in which case adaptability becomes important. With the growing complexity of hardware systems, the notion of a hardware operating system [13, 14] becomes more and more appealing, particularly if it affords the system some measure of autonomy.

A complexity argument in favor of autonomous systems comes from the growing productivity gap: The fact that the number of transistors on a chip is growing more rapidly than the designer’s ability to use them. There is much talk about the need for better design tools, but at large enough scales, the systems ought to begin carrying some of their own weight. If a system assumes more responsibility for itself, then the designer may be able to work at a higher level of abstraction, and may be able to focus on smaller self-contained components, without having to worry as much about the infrastructure of the system. The higher level of abstraction also means that designs become more portable to other devices and even architectures.

An economic argument in favor of autonomous systems comes from the serious impact of shrinking feature size on semiconductor manufacturing yield. The industry depends upon perfect and identical devices, whether they are general purpose processors or memories or ASICs or FPGAs, and even minor manufacturing defects on a device can sometimes render it entirely useless. That might not be the case if a device could instead oversee its own programming, and could avoid defective areas in the same way that a disk drive masks out faulty sectors. If the industry can relax its demand that devices be perfectly identical, by letting the system manage its own resources, the effective manufacturing yield may perhaps increase. And as the industry continues to approach nanometer feature sizes, where defect-free manufacturing is not achievable, the ability to tolerate defects and faults becomes an absolute necessity.

There are many reasons to ask for systems that are easier to develop, operate, and maintain, but that requires transferring more of the complexity into the systems themselves. Systems should take responsibility for the things that they can do on their own, and should ideally be able to learn for themselves. In some ways, this is the “thinking machine” advocated by Hillis [15]. The present work, however, is interested in autonomous systems because of the capabilities that they provide, and not simply because they can effectively address a particular problem domain.

But autonomous computing systems also have drawbacks. There is an area overhead required by the autonomous logic. There is a performance penalty for implementing functionality on top of configurable logic. There is a need to develop distributed and incremental versions of tools and algorithms that are more efficiently suited to autonomous systems. There is the added complexity of testing devices that are not all identical, nor even static over time. And there is probably also a necessary paradigm shift in the way that hardware and software designers think and work. Many objections and similar issues arose with assemblers and compilers and operating systems and object oriented languages at one time or another, but these tools are used extensively today and their benefits are unquestioned.

The close connection between living systems and what is desirable in autonomous computing systems makes it natural to personify them. But far from reflecting a belief that such systems will spontaneously come to life, it simply reflects the fact that *a system acting on its own no longer appears to be completely inanimate*. And while this work does not concern itself directly with artificial intelligence, it does consider what infrastructure might be necessary to support autonomy on large computational systems. In anticipation of Avogadro-scale computing, it is sensible to ask how such systems might be designed, and it is worthwhile to learn from biological systems, which do in fact grow, learn, adapt, and reason.

Autonomy in this work denotes the right to self-govern, the ability to act, freedom from undue external control, and responsibility for actions taken.

1.3 Objectives

This work has two objectives: 1) To propose a roadmap for the development of autonomous computing systems, and 2) to implement a proof-of-concept system that demonstrates autonomy. Both the roadmap and the implementation are further detailed in subsequent chapters.

The reader should note on one hand, that although the implementation presented here is based upon FPGA technology, the concepts extend much more broadly, and on the other hand, that although mainstream technology offers no way for computing systems to *physically* extend themselves, that capability is the subject of credible large-scale research. These points are revisited in greater detail in chapters 3 and 4.

1.3.1 Roadmap

The roadmap section considers the components and functionality required to support autonomy in computing systems, and proposes a collection of levels that provide that functionality. The levels are conceptually straightforward, and provide a roadmap to guide the implementation phase of the present work and of larger efforts.

1.3.2 Implementation

The implementation section tests the validity of the roadmap, and surpasses any known prior work by demonstrating autonomy at a very fine granularity inside a running system. The system is able to detect conditions in itself or its environment, and respond to them in some suitable manner, whether by implementing new hardware functionality within itself, or by invoking or enabling software functionality.

Although the relevant details will be discussed at length in subsequent chapters, it is worth pointing out that all implementation tools that the system requires will have to run *inside* the system that they are modifying. In that sense, the system will be not *closed*, but at least partly *self-contained*.

1.4 Organization

This dissertation is organized into six major sections, supported by background and conclusion chapters:

Chapter 2 - Background: “*What are we talking about anyway?*”

Reviews configurable computing, available tools, and related work.

Chapter 3 - Theory: “*What kinds of changes are possible?*”

Discusses the dynamics that may be present in computing systems.

Chapter 4 - Roadmap: “*Where are we trying to go?*”

Discusses the steps or components necessary to support autonomy.

Chapter 5 - Hardware: “*How is the hardware structured?*”

Discusses the hardware design and implementation.

Chapter 6 - Software: “*How is the software structured?*”

Discusses the software design and implementation.

Chapter 7 - System: “*How about the entire system?*”

Discusses the server and the larger system.

Chapter 8 - Operation: “*But does it work?*”

Discusses the server and system operation.

Chapter 9 - Conclusion: “*What’s the bottom line?*”

Summarizes results and discusses future work.

Chapter 2

Background

2.1 Overview

The roadmap and implementation phases of this work depend upon a wide range of hardware and software tools and subsystems, and this background chapter may consequently be somewhat disjoint. The information is offered more for completeness than for narrative value, and the reader should feel free to consider each section or subsection independently.

The major sections review digital design implementation, the hardware demonstration platform, the critical software subsystems, and prior or related work.

2.2 Digital Design Implementation

Because most engineers and scientists are unfamiliar with the concepts involved in digital system design and reconfigurable computing, this section presents an overview to help prepare the reader to better understand what is discussed in subsequent chapters.

Complementary Metal-Oxide-Semiconductor (CMOS) technology forms the basis of the majority of modern digital integrated circuits. CMOS technology is mature and well understood, and can be used to create microprocessors, memories, and a wide variety of other standard and non-standard devices, including Application-Specific Integrated Circuits (ASICs). An ASIC can be designed to perform some arbitrary function, and to do so with good size, power, and speed characteristics. But complex integrated circuits are difficult and costly to design, fabricate, and test, and they are only able to perform their intended predetermined function. In other words, the transistors and connections that constitute the circuit are fixed, and cannot be changed after fabrication.

Although it is not yet possible to physically extend devices, it is possible to introduce another layer of functionality into them, to allow some manner of configurability: at the cost of increased size and power and of decreased speed, it is possible to implement circuitry like that in Figure 2.1 that can be dynamically configured. These devices are broadly known as *Programmable Logic Devices*, with the largest and most powerful class being the Field-Programmable Gate Array (FPGA). As the name implies, an FPGA consists of an array of gates or more complex digital logic, that can be configured “in the field” instead of merely being pre-fabricated in specialized facilities.

Digital logic can be reduced to logic gates, and these gates can be implemented with CMOS transistors or with FPGA logic. Allowing for differences in size, cost, and performance, a design can then be implemented in either technology, and it is a common practice to prototype and debug designs in FPGAs before implementing the final product in ASICs. It is also increasingly common to simply implement the design with FPGAs in the final product, particularly in low volume designs or designs that require reconfigurability.

But digital design is not usually done at the gate level—systems are much too complex for that. Instead, designers generally use higher level Hardware Description Languages (HDLs) such as VHDL, Verilog, or SystemC, whose appearance deceptively resembles that of high-level software programming languages. An important property of HDLs is that they allow designers to make use of hierarchy and of behavioral coding. The hierarchy provides a way to design very complex systems in a manageable fashion. The behavioral coding provides a way to specify complex functionality without having to consider its underlying logic implementation.

Behavioral coding allows a designer to decide that they want to add two numbers for example, without immediately having to think in terms of the gates that will perform the addition. Thus the designer may construct an 8-bit adder with the VHDL code shown in Figure 2.2, knowing that the simple `x <= a + b` statement in Line 14 will actually be *synthesized* into circuitry like that of Figure 2.3. The synthesized circuitry must then be *mapped* to the logic resources of the target architecture, and must finally be *placed* in available resources and *routed* in order make the necessary connections between logic resources.

The separation between the described functionality and the underlying logic is reconciled with the help of design tools that perform the synthesis, mapping, placement, and routing:

Synthesizer: Translates a behavioral description into a structural description [16]. Gajski likens this to the compilation of high-level software languages into assembly language. The resulting gates and connections form a *netlist*, that is essentially an unconstrained graph of the circuit.

Mapper: Maps the netlist gates to the primitives of the underlying technology. In the case of an ASIC, the gates would be mapped to transistors or to predefined blocks called *standard cells*. In the case of an FPGA, the gates would be mapped to the logic primitives available on its particular architecture, including *lookup tables* (LUTs) and *flip-flops*. The resulting mapped primitives form a *technology-mapped netlist*, as shown in Figure 2.4.

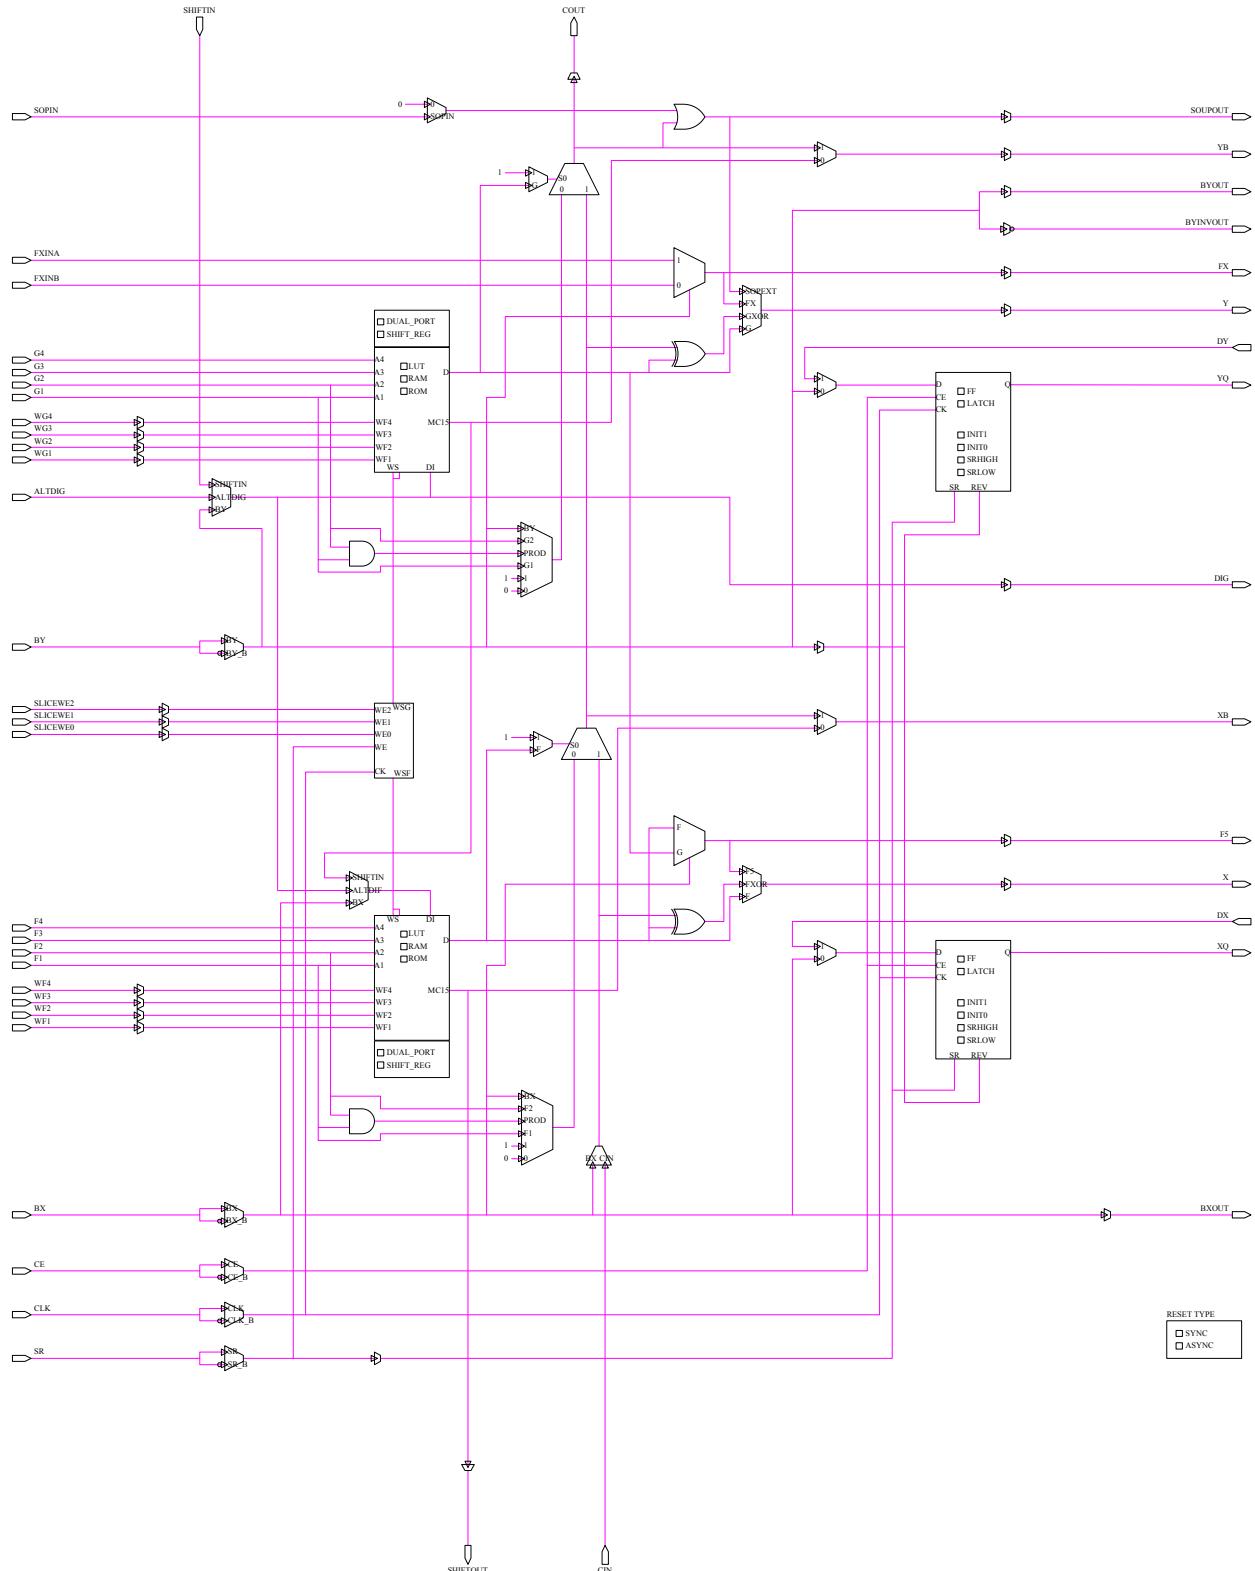


Figure 2.1: Configurable logic example: Virtex-II slice.

```

1 library IEEE;
2 use IEEE.STD_LOGIC_1164.all;
3 use IEEE.STD_LOGIC_ARITH.all;
4 use IEEE.STD_LOGIC_UNSIGNED.all;
5
6 entity adder is
7     port ( a : in std_logic_vector(7 downto 0);
8             b : in std_logic_vector(7 downto 0);
9             x : out std_logic_vector(7 downto 0));
10 end adder;
11
12 architecture behavioral of adder is
13 begin
14     x <= a + b;
15 end Behavioral;

```

Figure 2.2: 8-bit adder: behavioral VHDL.

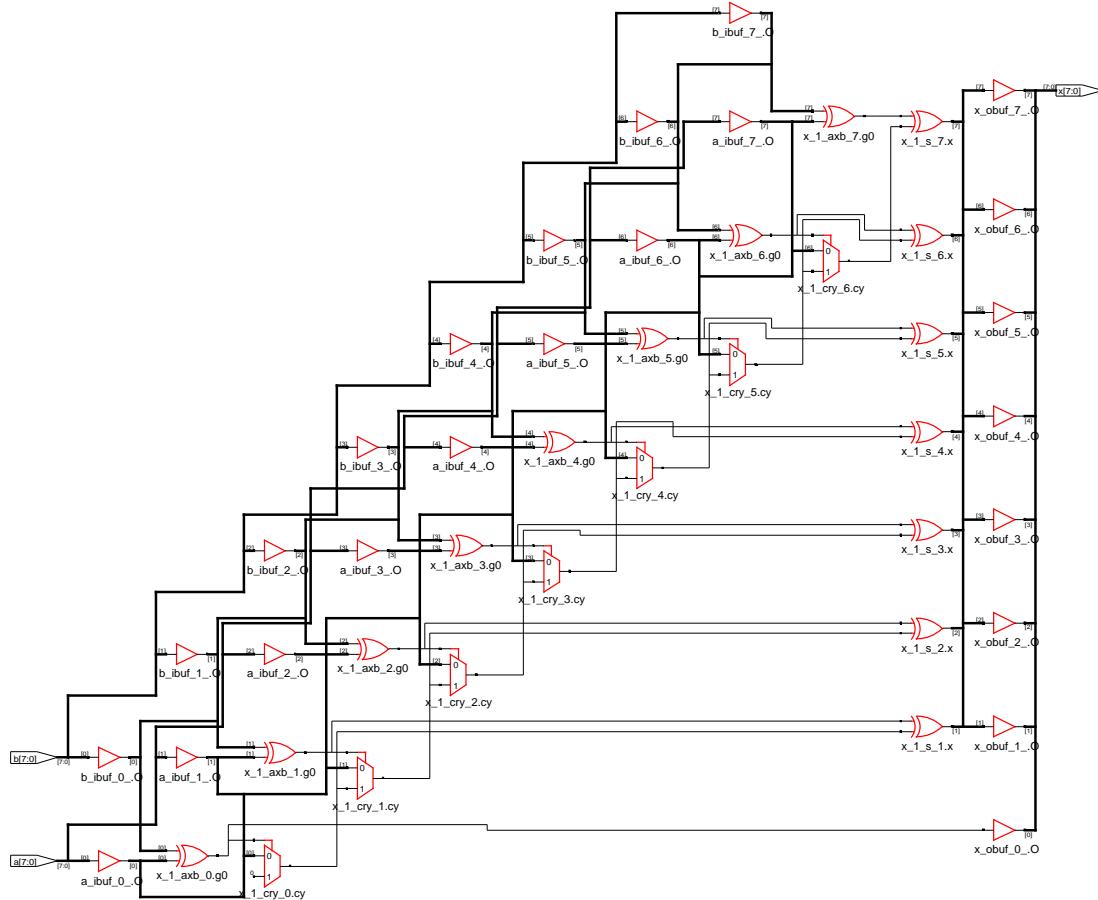


Figure 2.3: 8-bit adder: technology-independent schematic (with I/O ports).

Placer: Takes the technology-mapped netlist and assigns a position in the target device to each primitive in the netlist. The resulting netlist is no longer an abstract mathematical graph, but an allocation map of the necessary resources in the target device. The placed but unrouted design is shown in Figure 2.5.

Router: Determines how to connect each of the placed primitives. Depending on the granularity of the primitives, this may comprise millions or even billions of connections. The routing and placing are often performed jointly to minimize congestion and wiring delays, and it is important to note that routing is generally the most time consuming step in the tool flow. The fully routed design is shown in Figure 2.6.

This process is significantly complicated by the various constraints that the design may impose. This is particularly true with timing constraints and signal propagation times in high speed designs, where the speed of light starts to become a non-negligible factor.

Once the placing and routing are finished, a designer targeting an ASIC would generate a layout and a set of fabrication masks, usually at significant cost. A designer targeting an FPGA would instead generate a *configuration bitstream* that could be sent to the *configuration port* of an off-the-shelf FPGA, in order to implement the intended design.

Of particular interest to the present work are FPGAs that can be configured 1) partially, 2) dynamically, 3) internally, and 4) repeatedly. This means that it is possible to 1) change a subset of the FPGA, 2) while the rest of it continues to operate unchanged, 3) from inside the FPGA, 4) as many times as desired. These properties, collectively described as *partial dynamic reconfigurability*, provide the basic foundation for an exploration of autonomous computing.

Of the large-scale commercially available FPGA architectures, only the Xilinx Virtex-II, Virtex-II Pro, Virtex-4, and Virtex-5 families presently support partial dynamic reconfiguration. Other important FPGA families are available from manufacturers such as Actel, Atmel, Lattice, and particularly Altera, but those families are unusable in the current work for one of two reasons: Either they are too small to support the necessary overhead and infrastructure for autonomy, or they do not support partial dynamic reconfiguration.¹

¹Partial dynamic reconfiguration still holds greater research interest than commercial interest at present, and the lack of support for it in competing devices should not be taken to mean that those devices are without merit.

```

1 | (edif adder
2 | :
3 | :
4 | (library work
5 |   (edifLevel 0)
6 |   (technology (numberDefinition ))
7 |   (cell adder (cellType GENERIC)
8 |     (view behavioral (viewType NETLIST)
9 |       (interface
10 |         (port (array (rename a "a(7:0)" 8) (direction INPUT))
11 |           (port (array (rename b "b(7:0)" 8) (direction INPUT))
12 |             (port (array (rename x "x(7:0)" 8) (direction OUTPUT))
13 |           )
14 |         (contents
15 |           (instance x_1_axb_1 (viewRef PRIM (cellRef LUT2 (libraryRef VIRTEX)))
16 |             (property init (string "6"))
17 |           )
18 |           :
19 |           (instance GND (viewRef PRIM (cellRef GND (libraryRef UNILIB))))
20 |             (net (rename a_0 "a(0)") (joined
21 |               (portRef (member a 7))
22 |                 (portRef I (instanceRef a_ibuf_0))
23 |               )))
24 |               :
25 |               (net (rename x_1_1 "x_1(1)") (joined
26 |                 (portRef O (instanceRef x_1_s_1))
27 |                   (portRef I (instanceRef xobuf_1))
28 |                 )))
29 |               )
30 |             (design adder (cellRef adder (libraryRef work))
31 |               (property PART (string "xc2v40cs144-6") (owner "Xilinx")))

```

Figure 2.4: 8-bit adder: technology-mapped netlist.

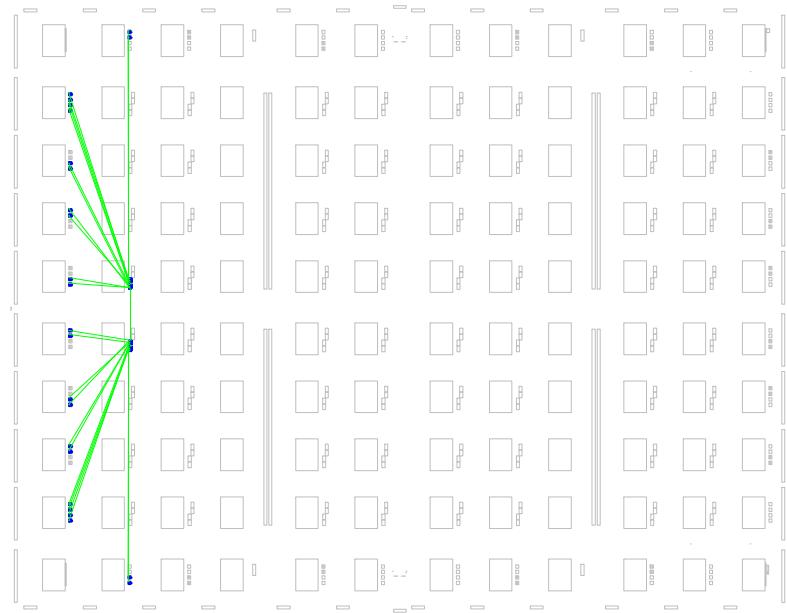


Figure 2.5: 8-bit adder: placed but unrouted design in XC2V40-CS144—FPGA Editor view.

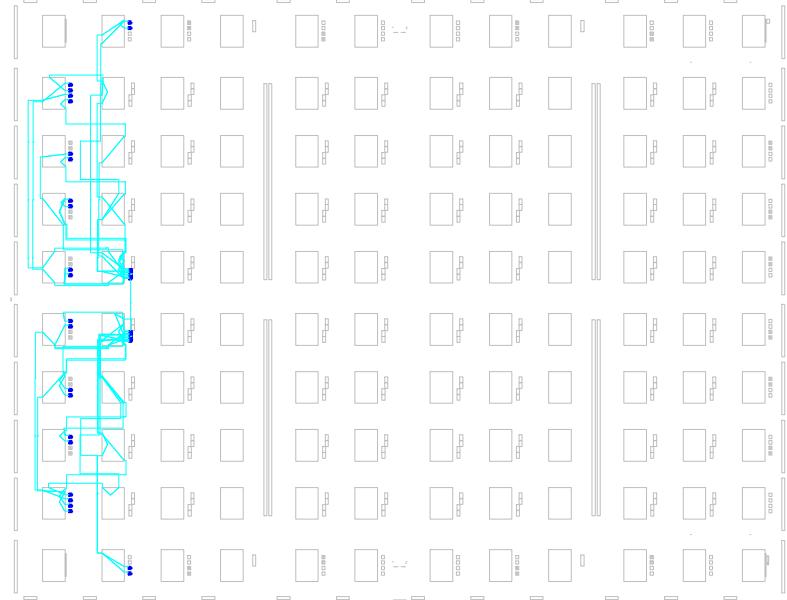


Figure 2.6: 8-bit adder: fully routed design in XC2V40-CS144—FPGA Editor view.

2.3 Virtex-II Pro Architecture

The Virtex-II Pro architecture was introduced in 2002 as a new family of Xilinx FPGAs [17, 18], retaining all the capabilities of the prior Virtex-II architecture [19], and adding embedded PowerPC 405 processors and gigabit transceivers. Although these architectures have since been surpassed by the Virtex-4 and Virtex-5 families, they still provide a large amount of processing power, configurability, flexibility, and compatibility with I/O standards.

In most of this dissertation, the formal Virtex-II and Virtex-II Pro names will be set aside in favor of Virtex2 and Virtex2P. Furthermore, because the Virtex2P is an extension of the Virtex2 architecture, any statements that are made about Virtex2 capabilities, should be understood to apply to both Virtex2 and Virtex2P devices.

Many of the capabilities of the Virtex2P pertain to more mundane usage, and are thus not closely aligned with the scope of this work. Those capabilities are well documented in the Virtex-II Pro Platform FPGA Handbook [18], and in pertinent application notes published by Xilinx. This section will instead describe the Virtex2P from a vantage point relevant to this work, and will emphasize issues pertaining to partial dynamic reconfiguration, particularly when they are not well documented elsewhere.

One commonly used facet that nevertheless remains relevant here is the Xilinx Unified Libraries [20], a collection of hundreds of common logic elements or logic primitives that are used by designers and/or by synthesizers and mappers. The Unified Libraries are included in Xilinx's standard Integrated Software Environment (ISE), along with all of the other implementation tools.

The reason for the relevance of the Unified Libraries is the fact that netlists generated by synthesizers describe circuits in terms of these library elements. The implementation phase of this work includes the ability to process EDIF netlists, and it must therefore know how to instantiate, use, and configure the library elements that the netlists reference.

2.3.1 User and Device Models

Many hardware designers are able to treat FPGAs as abstractions, concerning themselves mainly with the fact that the devices can implement digital circuitry so as to perform desired functions. Others may delve a little deeper into the *user model* of the devices, as presented by FPGA Editor or by other ISE tools. But only a small number will ever need to understand the underlying *device model* that is used inside some of the Xilinx tools.

The *user model* of the Virtex2P consists primarily of logic cells and of routing resources. The FPGA is internally composed of *tiles* of varying types that form an irregular two-dimensional grid. These tiles in turn contain the logic and wiring resources that make up cells, wires, and switches, but the tiles themselves are generally not exposed in the user model.

Logic cells come in 27 different types, more fully described in Section 7.2.2. Many users are familiar with only a handful of these—SLICEs, TBUFs, IOBs, RAMB16s, and MULT18X18s—either because of their abundance or because of their scarcity. SLICE cells are by far the most abundant cells in a device, and are the primary resource used to implement digital functionality, so common in fact that the type name is generally rendered uncapitalized as *slice*.

The Virtex2P slice, previously depicted in Figure 2.1, contains two Look-Up Tables (LUTs), two flip-flops, a variety of multiplexers, as well as dedicated logic to support fast adders, wide functions, shift registers, and distributed memory. Many of these slice components are referenced directly or indirectly as Unified Library elements.

The slice is also logically and practically divided into four semi-overlapping quadrants, centered around the LUTs and the flip-flops, as further illustrated in Figure 2.7. The lower and upper LUTs are named F and G respectively, and the lower and upper flip-flops are named X and Y respectively. There is a close relationship between the F LUT and the X flip-flop, and between the G LUT and the Y flip-flop, but each of these groups and their elements can also be used separately.

As for the routing resources, many users are peripherally aware of the existence of clock trees, and of *longs*, *doubles*, and *hexes*, but little more. The routing resources in fact far surpass the logic cells that they serve, when measured in terms of die area² or complexity. The routing, unrouting, and tracing problem was addressed at length in prior work [22], and enters only indirectly into the present work. This work specifies cell pins as endpoints, and simply invokes the router, unrouter, or tracer to do the heavy lifting and deal with the full complexity of wiring resources in the appropriate family.

The *device model* directly relates to the *user model*, but not always in one-to-one fashion. Stated in a different manner, the *apparent* device model that the software presents to the designer is in fact different than the *actual* device model. This discrepancy makes conversions between the two models a lossy process, and the cause of significant extra complexity in the low-level tools. This is also part of what thwarts most attempts to reverse-engineer configuration bitstreams. It is generally possible to extract a considerable amount of information through some tuning and trial-and-error, and many researchers have attempted to do just that, but most would-be reverse-engineers fail to appreciate how complex the user/device model mappings actually are.

Tiles take on much greater prominence in the device model, by virtue of the fact that they more closely represent the VLSI layout of the physical die. Configuration bits that control logic cells and routing resources are no longer necessarily confined to the tiles that the resource appears to reside in. In the underlying device data, tiles are in fact collections of configuration bits, rather than collections of logic and routing resources, and are therefore of greater importance to bitgen³ than to FPGA Editor.

²A Xilinx Field Application Engineer once tried to explain the economics of FPGAs—quantifying die area usage for logic versus routing—in the following manner: “We sell you the routing and give you the logic for free.” [21]

³Bitgen is the bitstream generation tool provided with ISE.

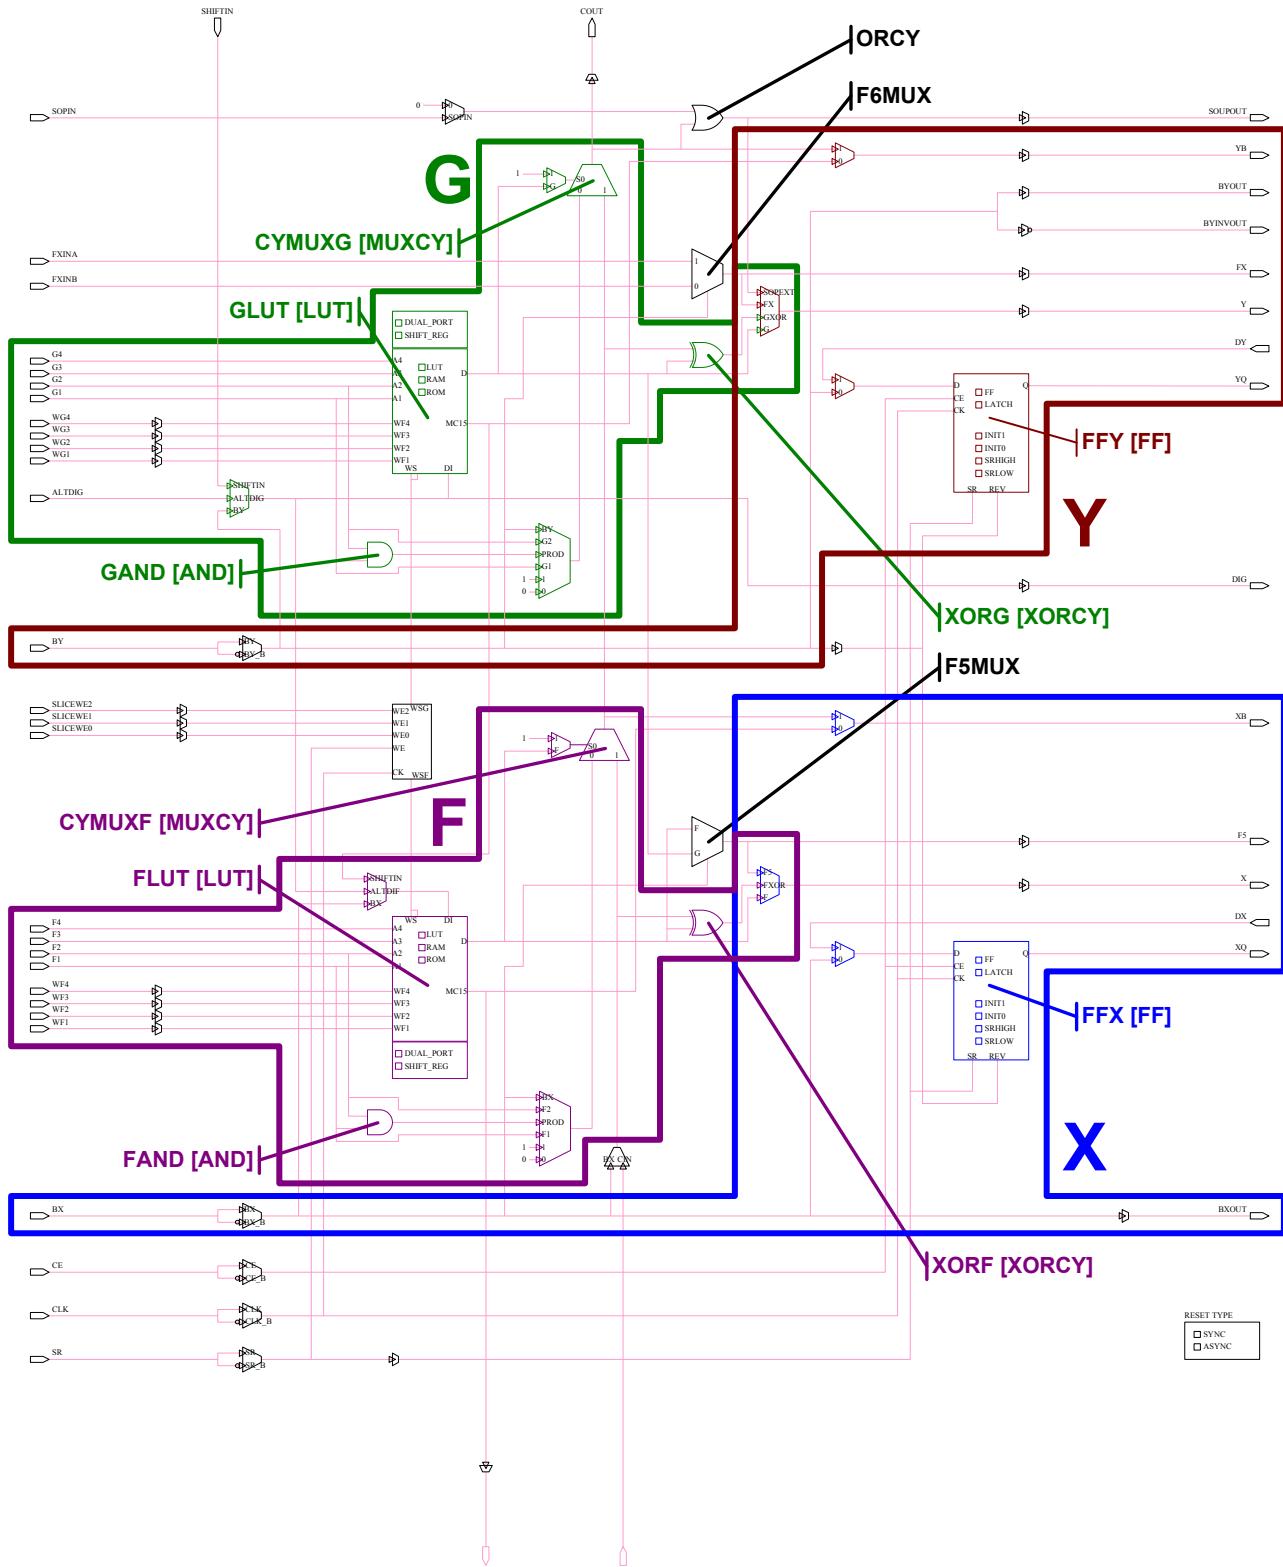


Figure 2.7: Virtex-II slice subcells: G and Y in upper half, and F and X in lower half; callouts indicate subcells that are referenced directly or indirectly by library components.

2.3.2 Configuration and State Planes

In the context of partial dynamic reconfiguration, there are a few aspects of the device model that become relevant to the current discussion. Although designers tend to think of FPGAs as collections of logic cells and wiring resources—which is precisely the abstraction that they are designed to present—it is equally appropriate to think of FPGAs as collections of configuration bits. Bitgen treats the device *bitmap*⁴ as a collection of configuration frames—where a frame is the smallest independently reconfigurable portion of a device—generally a vertical column spanning the entire height of the device or the height of a clock domain.

The situation is further complicated by the fact that some configuration bits in the bitmap define the configuration of the device, while others define the state of the user’s design. In practice the device behaves as though it has two parallel planes—a *configuration plane* and a *state plane*⁵. Within the actual device, all of these bits exist in the form of SRAM memory.

Configuration data that is written to or read back from the FPGA always accesses the configuration plane. To push data from the configuration plane up to the state plane, it is necessary to generate a reset.⁶ To pull data from the state plane back down to the configuration plane, it is necessary to generate a *capture*.⁷

The difficulty which the configuration and state planes introduce with respect to partial dynamic reconfiguration is that they are not as cleanly separated as one would expect. Configuration information always resides on the configuration plane, but state information does not always reside solely on the state plane.

2.3.3 State Corruption Considerations

To avoid duplicating memory bits for certain kinds of resources, the Virtex2 architecture keeps some of its state information on the configuration plane. That means that updates to the configuration plane can corrupt information that logically belongs to the state plane.

⁴A *bitmap* is a representation of configuration frames in memory, while a *bitstream* is an encoding of those configuration bits into packets suited for the device’s configuration controller.

⁵Xilinx does not use the *configuration plane* or *state plane* terminology.

⁶In the case of a full configuration bitstream, the device’s Global Set/Reset (GSR) signal is asserted when configuration completes, in order to preload all flip-flops with their default settings. In the case of an active partial configuration bitstream, when a portion of the device must be reconfigured while the rest of it continues to run, the GSR option is unacceptable. In such cases, it is possible to toggle the SRINV setting two times in a row for all slices and IOBs—while temporarily forcing the flip-flops into ASYNC mode if no clock is present—in order to achieve the same effect. If the circuit or circuits that were reconfigured contains a reset input, it may be easier to simply assert that signal to achieve that same effect.

⁷The capture functionality can be triggered from inside the FPGA, from a configuration bitstream, or from the JTAG interface.

2.3.3.1 LUT RAM

The LUT resources in a Virtex2 slice can be used in one of three modes: Look-Up Table, ROM, or RAM. In both Look-Up Table and ROM modes, the function of the LUT is controlled by configuration bits whose settings do not change. In RAM mode, however, those same configuration bits change whenever the user’s design performs a write operation to the LUT RAM.

The reconfiguration time for a single frame on the Virtex2P XC2VP30 is on the order of $10\ \mu s$, while a typical clock period on the same device is on the order of 10 ns. That difference means that the contents of a LUT RAM can change nearly a thousand times during a single frame reconfiguration, so the read-modify-write scheme is unsuitable unless writes to the LUT RAM in question are suspended during reconfigurations. Suspending circuits during reconfiguration is undesirable however, because it runs contrary to the general compact that hardware is always online—some circuits can never tolerate being suspended. Functionality that can tolerate being suspended can probably be implemented more easily in software, and can therefore be supported without the more extensive infrastructure presented in this work.

Because of the potential for state corruption, LUT RAMs are forbidden within any configuration frames that dynamic circuits may use.⁸

2.3.3.2 SRL16

SRL16s are LUT RAMs configured in shift-register mode, and are therefore vulnerable to the same state corruption risks as general LUT RAMs. These elements are therefore forbidden in dynamically usable regions.

2.3.3.3 BRAM

Virtex2 BRAMs⁹ are 18 Kb dual-ported static memory blocks, distributed throughout the configurable fabric. When these devices are used in single-ported mode, there is no risk of state corruption.¹⁰ However, in dual-ported mode, one of the two ports doubles as the block’s interface to the configuration controller. This means that if a circuit and the configuration controller both try to access the same BRAM simultaneously, even if the configuration controller is merely trying to perform a readback operation, it may interfere with the circuit’s access to the BRAM, and thus cause state corruption.

⁸It is important to understand that *every* LUT within a frame is modified whenever that frame is rewritten. This means that there is no way to selectively avoid updating LUTs that happen to be configured in RAM mode.

⁹The actual cell type is RAMB16.

¹⁰The absence of state corruption risk reflects a *best understanding*, and should not be taken as an authoritative statement.

It is worth pointing out that 18 Kb BRAMs are relatively large resources in an FPGA. To guarantee that no contention could occur, the FPGA would have to include shadow memory on the configuration plane, effectively doubling the amount of memory bits required by BRAMs. Considering that many FPGA users never require readback, this largely unnecessary replication of memory bits would be difficult to defend.

2.3.4 Unsupported Slice Logic

To avoid state corruption issues, all LUT RAM and shift-register functionality in slices is forbidden for dynamic circuits. The wires and logic subcells shown in red in Figure 2.8 are thus never used in dynamic circuits in the present work.

Fortunately, configuration bits for LUT RAMs and BRAMs *never* overlap with configuration bits for routing resources in the Virtex2 architecture. The existence of such resources in non-dynamic circuitry—the base system—therefore does not prevent the system from safely routing or unrouting anywhere that it pleases in the entire device. This point will be revisited in Section 5.9.

2.4 XUP Board

The implementation phase of this work uses the Virtex-II Pro Development System [23], designed by the Xilinx University Program, and manufactured by Digilent, Inc. [24]. This work will refer to the board more informally as the *XUP board*.

The XUP board is a wonderful development platform, built around a Virtex-II Pro XC2VP30 FPGA, and its two embedded PowerPC 405 cores. The board has a slot for up to 2 GB of DDR SDRAM, a CompactFlash connector, dual PS/2 connectors, an RS-232 port, a 10/100-capable Ethernet connector, AC-97 audio in and out connectors, an SXGA-capable video DAC with connector, along with switches, buttons, LEDs, and the ability to accept daughter cards. The board is shown in Figure 2.9.

Development was done with a -6 speed grade XC2VP30, 512 MB of DDR SDRAM,¹¹ a 1 GB CompactFlash MicroDrive, and a Digilent VDEC1 video decoder daughter card. Plugged into the board were a PS/2 keyboard, a PS/2 mouse, a 1280 × 1024 SXGA monitor, an NTSC camera, a DVD player, and Ethernet and RS-232 connections.

¹¹Experience shows that without the proper timing parameters for the DDR controller, memory modules will not work on the board. Modules not listed as *Qualified SDRAM Memory Modules* listed in [23] are unlikely to work properly.

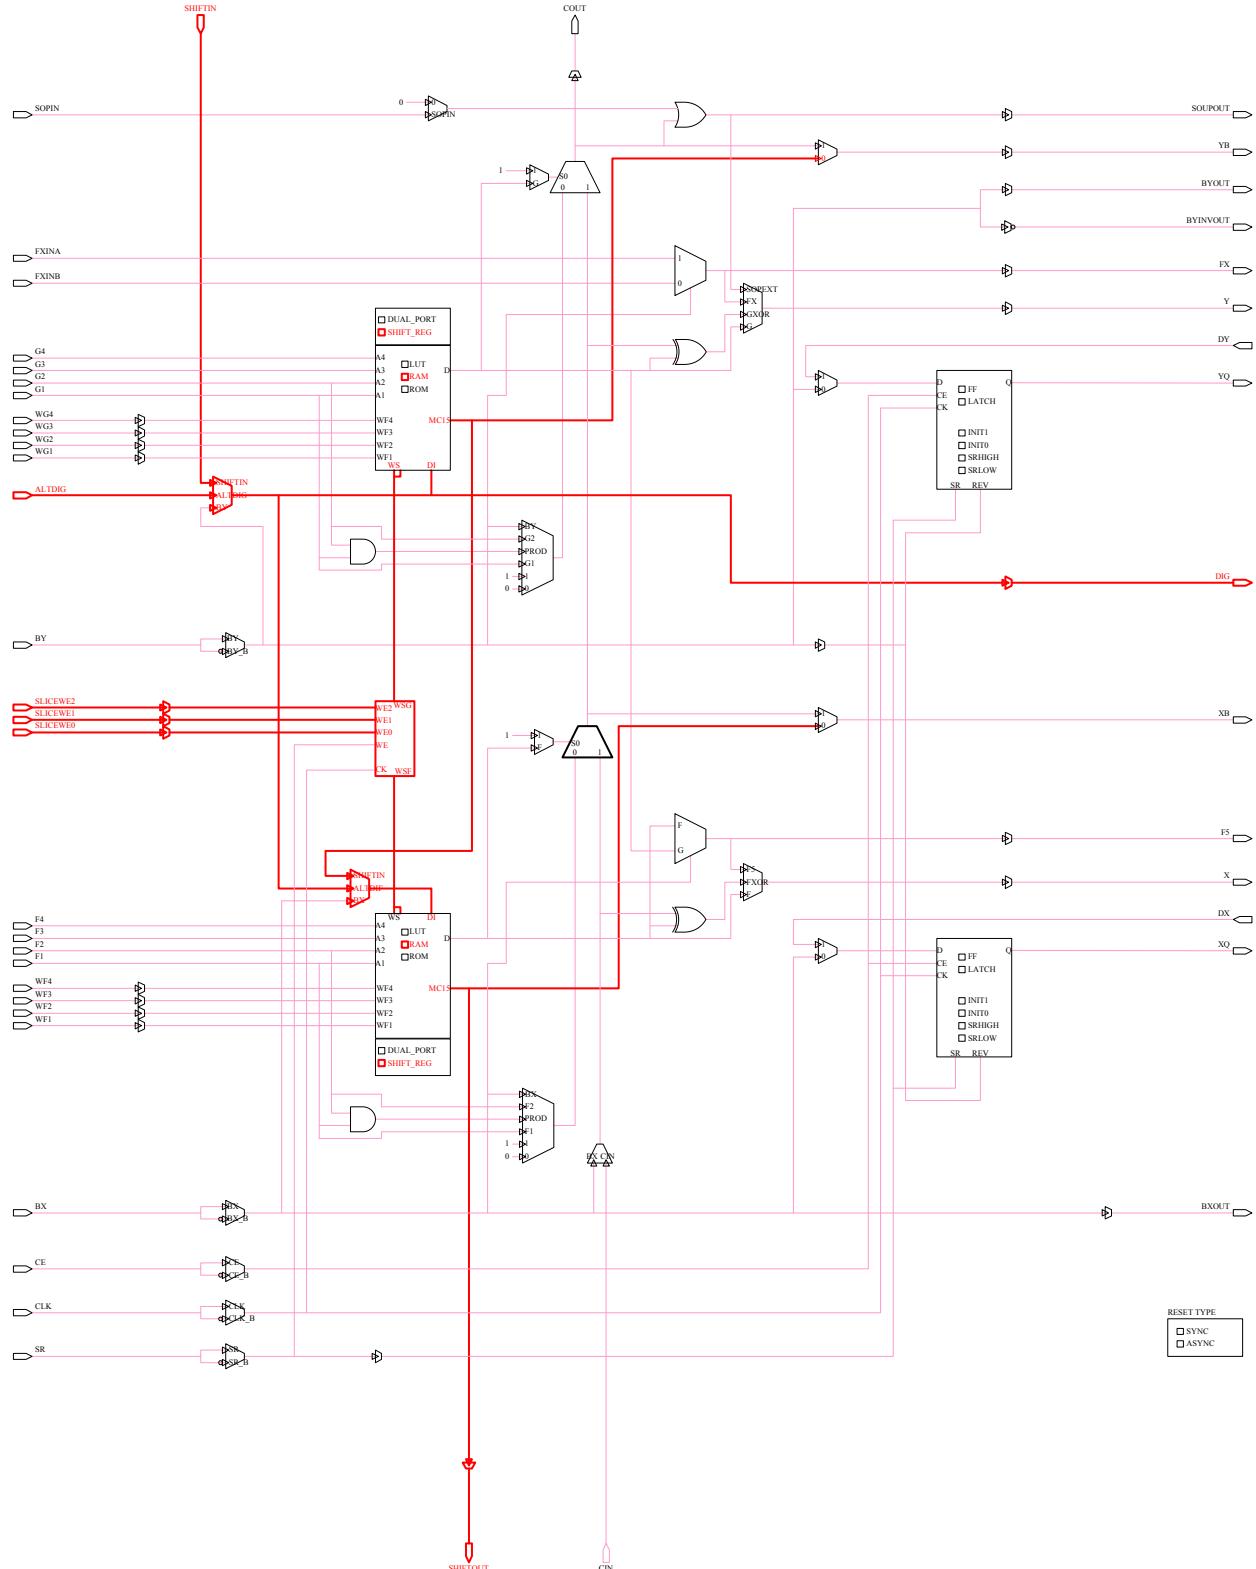


Figure 2.8: Virtex-II unsupported slice logic: shift-registers and LUT RAMs are disabled.

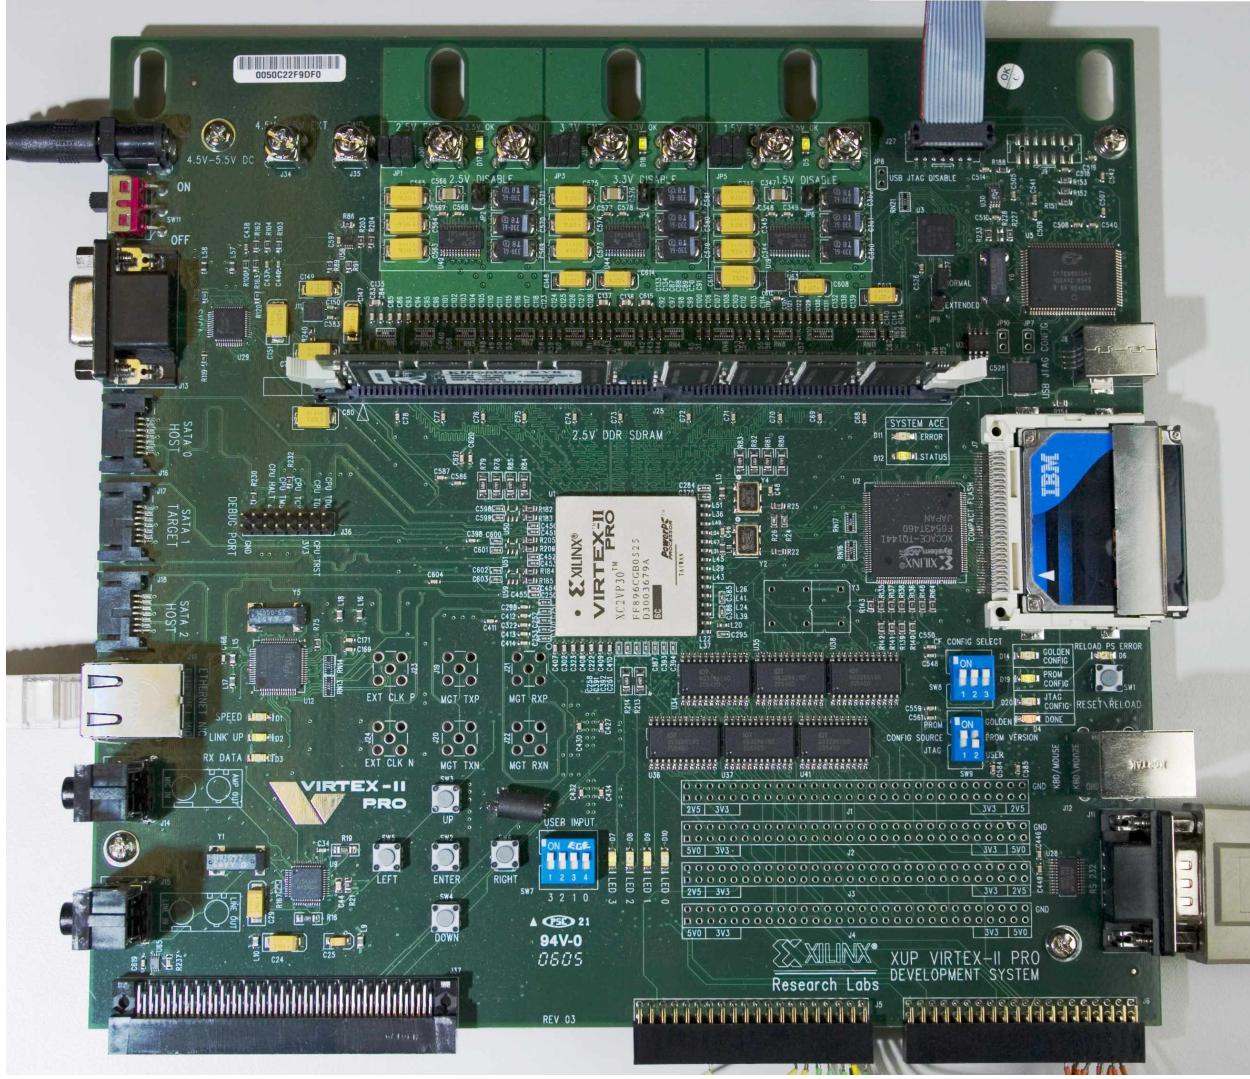


Figure 2.9: XUP Virtex-II Pro development system.

2.5 ADB

A system that generates hardware dynamically must also be able to perform dynamic routing and unrouting. While some systems only require trivial routing capabilities, others require a much more complete infrastructure. The Alternate Wire Database (ADB) developed in prior work [22] provides complete routing, unrouting, and tracing services for the current work.

Although ADB supports a number of Xilinx FPGA families, the present work employs it in support of the Virtex2P XC2VP30 that is present on the XUP board. The underlying database provides

exhaustive wiring information for the device and its 1,433,460 wires. The one notable exception is the absence of timing data, unfortunate because it means that ADB can neither attempt to meet timing requirements while routing, nor determine the maximum delay of any route that it has routed or traced. For commercial purposes, the absence of timing data would be unacceptable, but for the sake of research, ADB is otherwise complete, correct, embeddable, and—most importantly—available.

The low-level details and capabilities of ADB are discussed at length in [22], so they will not be revisited here. This section will instead focus on the capabilities that directly support the implementation phase of the work.

2.5.1 Routing

The router included with ADB can make connections between arbitrary device wires, provided a path exists. Routes can begin from a specified source wire or from an existing net, and can terminate in one or more specified sink wires. This facilitates many different kinds of connections, including the following:

- Routing from a source to a sink in a user circuit.
- Routing from a source to multiple sinks in a user circuit.
- Extending a user net with one or more additional sinks.
- Connecting multiple user nets to form a top-level system net.

Minor extensions to the router and heuristic were developed for the present work, and provide the ability to reserve wires,¹² to change the ordering of multi-sink nets, and to inform listeners of all connections made.¹³

2.5.2 Unrouting

A system that can dynamically implement circuits will probably also need to remove circuits at some point. The ability to remove connections inside user circuits, as well as connections between user or system circuits, is provided by the unrouter. The unrouter relies heavily on the tracer, described next, and can function in a number of ways:

¹²The Virtex2 architecture includes certain bidirectional inputs that can function as convenient waypoints to significantly improve routability, but using the wires in this way ensures that their inputs become unreachable. The ability to reserve inputs that a circuit may need makes it possible to bounce off bidirectional inputs that will not be used, while still guaranteeing that inputs which must be used will still be reachable.

¹³For the sake of debugging and metrics, it was desirable to export all dynamically implemented circuits in a format suitable for inspection with FPGA Editor, and for full timing analysis. The intermediate XDL format required complete *pip* information that would otherwise have been unavailable.

- Unroute everything in the *sinkward* direction from a specified wire.
- Unroute the current branch in the *sourceward* direction from a specified wire.
- Unroute the full net in all directions starting from any wire on the net.

Along with corresponding capabilities for logic cells, the unrouter provides the ability to remove connections and release their resources for use by other circuits, or to modify existing connections dynamically. The router and unrouter together allow the system to arbitrarily add or remove connections without killing the system itself.

2.5.3 Tracing

The tracer is able to look at an arbitrary device wire and identify any sources and sinks that may belong to the same net as that wire. In addition, it is able to trace through configuration bitstreams to infer the wiring usage information, and thus ensure that the router or unrouter may safely modify any part of an existing system without corrupting it. The tracer can function in four different ways:

- Trace everything in the *sinkward* direction from a specified wire.
- Trace to the next branch in the *sourceward* direction from a specified wire.
- Trace to the net source from a specified wire.
- Trace the full net in all directions starting from any wire on the net.

When a configuration bitstream is traced, the tracer iterates through all logic cell pins, and finds all paths that connect a cell output to one or more cell inputs.¹⁴

2.6 Device API

As discussed in Section 2.3.1, the user and device models are quite complex, and may not map simply onto one another. The amount of underlying data alone is impressive [22], and subject to enough exceptions that even clever compression and encoding techniques can only compact the data so far. Earlier research tools, such as JBits [25], were able to rely on assumptions that were largely valid for the architectures that they targeted, but those assumptions no longer hold in more modern

¹⁴Ignoring paths that begin or end in stubs, and do not form a complete path from an output to one or more inputs, is less than ideal because it raises the possibility that stub wires may be marked unused when they are in fact in use. The Virtex and Virtex-E architectures sometimes require stubs that tie unused wires to known levels, to prevent destructive oscillation within the device. Because of this, although the tracer finds all paths that start from a cell output, it discards any that do not reach a cell input. In practice this approach does not cause the problems that one might fear.

architectures. Devices are no longer uniform arrays of configurable logic blocks, and configuration bits no longer necessarily reside in the same part of the device as the resources that they control.

An unreleased Device API tool was developed for the sake of generating configuration bitstreams, while remaining rigorously faithful to the underlying device models and data.¹⁵ This tool should be thought of as a lighter and factored version of the back-end ISE tools, in that it provides complete logic and wiring information, and is able to generate partial dynamic configuration bitstreams.

The Device API tool is unfortunately unadvertised and unreleased, and simply unavailable outside of Xilinx. Nevertheless, the tool does perform as described, and does so very well. If that were not the case, the claims and results presented in this work would be untenable, as anyone who understands low-level Virtex2P internals can confirm.

2.7 Linux on FPGAs

Linux is a modern and highly portable open-source kernel and operating system, running on platforms ranging from embedded systems to supercomputers, on over 20 different processor families [26]. The availability of Linux for the PowerPC family means that it can run on the embedded PowerPC 405 processors inside the Virtex2P XC2VP30, which makes it an ideal operating system for FPGAs.

A lightweight variant of Linux for processors without Paged Memory Management Units (PMMUs) is also available, under the name uClinux [27], but the absence of a PMMU makes it unsuitable for many software packages that are required in the present work, notably Java.

Porting Linux to a new system is generally a fairly involved endeavor, and many users take advantage of its open-source character to re-use as much as possible of other people's work.¹⁶ Much of the work involved in the Linux port used here originates with MontaVista Linux [28].

A number of groups have implemented Linux on the XUP board, and have posted information and instructions to facilitate the process for others. The information posted by the University of Washington (UW) and by Brigham Young University (BYU) have proven especially valuable for the current work, and are gratefully acknowledged.

¹⁵Xilinx graciously supported this research by granting the author exceptional access to confidential and proprietary information. The reader is advised that Xilinx software license agreements do not permit reverse engineering of the company's intellectual property in any form.

¹⁶Specialized or embedded platforms frequently include special hardware, and the availability of drivers for the special hardware can vary. When existing drivers are not available, the person doing the port must develop their own, and in some cases that can be a very ambitious task.

2.7.1 University of Washington: EMPART Project

The Embedded Research Group at the University of Washington is working on a project named EMPART [29], to accelerate and encourage the adoption of *hybrid micro-parallel* architectures on reconfigurable platforms.

The EMPART project cites work from Brigham Young University [30], the University of Illinois at Urbana-Champaign [31], the University of California at Berkeley [32], Wolfgang Klingauf [33], BusyBox [34], crosstool [35], and Paul Hartke.

The information provided guides the reader through the hardware setup in Xilinx's Embedded Development Kit (EDK), creating a cross-compiler, patching, configuring, and building the Linux 2.4.26 kernel, configuring a CompactFlash card that is bootable with the Xilinx SystemACE, and building a root file system.

2.7.2 Brigham Young University: Linux on FPGA Project

The Configurable Computing Laboratory at Brigham Young University has posted their Linux on FPGA work [30], "focusing on the development of software, methodologies, and instructions for porting and using Linux on FPGA-based platforms."

Although UW based much of their Linux work on BYU's work, and did a good job of guiding the reader through the necessary steps, there were a few times when reading through the BYU information helped to supplement the discussion. The ability to compare notes was very helpful in light of the daunting process of getting Linux to run on an embedded FPGA platform.

BYU also provides instructions on building OpenSSL and OpenSSH, for use under Linux on a Virtex2P, though those packages were built without knowledge of BYU's instructions in the present work.

2.7.3 University of York: DEMOS Project

The Real-Time Systems group at the University of York has built and made available a 2.6.17.1 Linux kernel for the XUP board [36]. The 2.6 kernel differs significantly from the 2.4 kernel, and though its use would have simplified some compatibility issues, it would also have given cause to other driver and tool compatibility issues.

In light of the expected compatibility issues, and because the present work is not focused on Linux per se, the DEMOS code was not utilized. But any future work would likely benefit from a more modern version of the Linux kernel, and the DEMOS project would likely be a good starting point.

2.8 Related Work

2.8.1 Autonomy

Although there is significant ongoing research into software autonomy [9, 10], very little has been attempted in the hardware domain. Autonomous systems such as *autonomous vehicles* or *unmanned aerial vehicles* are becoming increasingly sophisticated, but they have no ability to extend or modify their hardware, and are therefore of only limited interest to the present work.

In the domain of hardware autonomy, some of the most interesting work comes from Kubisch and Hecht at the University of Rostock [37, 38], whose research was discovered after the present work was already formulated. Kubisch et al. recognize the benefits of partial dynamic reconfiguration, of systems that observe themselves and their environments, of adaptation through source code modifications, and of realizing all of this from within the system. Their configuration granularity is coarse, and it is unclear how they constrain the configurations that they generate. But they effectively apply *Network-on-Chip* (NoC) expertise to the problem, and propose an interesting idea not considered in the present work: a versioning system that would track changes, and provide the ability to roll back to previous stable configurations should problems arise.

The ability to roll back is of critical importance for Kubisch et al. because they are implementing evolvable hardware, and some of the resulting configurations may consequently be nonfunctional or even invalid. The present work instead aims to make *known, deliberate changes*, where roll-back capability is much less crucial.

The work of Kubisch et al. qualifies as autonomous hardware because their system invokes the full Xilinx tool suite inside itself to modify its hardware. But the burden of running workstation-class tools inside an embedded system is high, and Kubisch et al. feel that their dynamic reconfiguration work is not viable with current FPGA architectures and tools [39], so they are deferring it while focusing on hardwired NoC communications between heterogeneous components inside simulated configurable architectures [40]. The similarities between the ideas of Kubisch et al. and the ideas presented here are made even clearer by an unpublished paper of theirs [38].

Other groups are less interested in dynamic reconfiguration, but still desire rapid tools to support what they call “Just-in-Time FPGA Compilation” [41] with the portability that it affords. These approaches typically evaluate available tools and architectures, but find them lacking for the intended purpose, and are forced to propose alternate configurable architectures. This author may have withheld due credit from these efforts in the past, because they neither target autonomy, nor provide tools that are usable with available architectures, but any work that furthers embeddable implementation tools is at least compatible with the spirit of the present endeavor.

There are other interesting concepts that arise from abstracting the underlying hardware, so that circuits can execute anywhere as if running on a virtual machine [42]. That approach has its place

and value, but doesn't appear to apply particularly well to the present work because it tends to foster conformity rather than autonomy.

One final endeavor that was discovered after most of the present work was formulated, is Wigley's research on operating systems for reconfigurable computers [43].¹⁷ Although Wigley's work lacks autonomy in that the processing does not occur within the target system, that limitation is substantially due to the hardware and tools available at the time. If the present work were to be extended or overhauled, a detailed analysis of Wigley's work would be desirable, to better understand the overlaps and tradeoffs involved. Conversely, the present work demonstrates in hardware much of the functionality that limited Wigley's work, and thus helps to show that the alternate architectures or extensions requested by Lysecky and Wigley are not in fact necessary.

2.8.2 Placers and Routers

The placement and routing problems mentioned in Section 2.2, and more fully described in [46] and other references, become increasingly complex as device and circuit sizes grow. A number of research tools tackle placement and routing, but usually target either abstract architectures or greatly simplified real architectures. While such concessions are normally acceptable within the scope of the research, they fail to address the full complexity required for real devices and circuits. A few relevant tools are worthy of mention.

Shortly after introducing JBits [25], Xilinx added JRoute [47], a router well known in the reconfigurable community. While JRoute did much to facilitate routing in reconfigurable systems, it initially only supported uniform grids of CENTER tiles,¹⁸ later extending that support to include I/O and BRAM tiles. But even with CENTER, I/O, and BRAM tile support, there were many areas and resources in the devices that JRoute was unable to accommodate. That situation worsened with each new FPGA family that Xilinx introduced, as newer architectures departed further and further from the simplistic assumptions ingrained within both JBits and JRoute. In spite of these limitations, JRoute is one of the few reconfigurable tools to have targeted a *real* architecture.

Another well respected research tool is Versatile Place and Route (VPR) [48], developed at the University of Toronto. The tool was at least partly intended for architecture exploration. It is fast, it gives good quality results, and its source code is freely available. Unfortunately, from the perspective of the current work, VPR performs no architecture-specific placement or routing, but deals only with abstract architectures. While that is sufficient for exploration purposes, it is insufficient for the purposes of this work. The code was evaluated with the intent of extending it to support the Virtex2P, but that idea was quickly dismissed as infeasible.

Another placer [49] that specifically targeted incremental design in FPGAs, and whose source was accessible, was dismissed as well, because of pervasive assumptions about uniform grids of

¹⁷Much of the earliest work on hardware operating systems is properly credited to Brebner [44, 13, 45].

¹⁸CENTER tiles in the Virtex architecture contain slices, and are the most common tile type in device.

CENTER tiles, and of circuits being already divided into clusters. In this case too, the effort required to extend the tool would have exceeded the effort of developing a placer from scratch.

Each of these tools, along with others, served a valuable research purpose, and performed fairly well within its intended context. Many aspects of these efforts were innovative and commendable, but it remains that none of them solved the full placement or routing problem for a real architecture. The importance of real architectures is frequently dismissed in the academic community, and understandably so, but the unavailability of a solid and extendable tool infrastructure makes it regrettably difficult to build on top of others' research. In many cases, researchers are limited in what they can support, because it can be nearly impossible to obtain actual device data. However, everything necessary to develop a true router or true placer for Xilinx architectures, is available from the output of the *xdl* tool,¹⁹ and it would be wonderful to see research groups begin to target these architectures in their work.

The present effort consequently had to provide a complete placer, to support every logic cell type in the Virtex2P architecture, and to rely on true device models rather than simplified assumptions about resource availability. The router was developed during prior work, as described in Section 2.5, and benefited from exhaustive wiring support for the Virtex2P and other families.

2.8.3 JHDLBits

An interesting effort that overlapped with this work in some unexpected ways was *JHDLBits* [50], an effort to tie the JHDL [51] front-end HDL capabilities to the JBits back-end bitstream generation capabilities.

JHDLBits functioned as a JHDL netlister that extracted netlist information, mapped referenced library primitives to logic cells or subcells, placed the cells, invoked ADB to route the circuits, and invoked JBits to generate the configuration bitstream. Although no code written for the JHDLBits project was developed in great depth, each of the pieces would have contributed useful functionality to the present work. In particular, the present work needs to extract netlist primitives and map them to device cells and then place those cells, making it unfortunate that the JHDLBits effort was not carried further.

¹⁹Many people are familiar with the *xdl* tool for its conversion modes *-ncd2xdl* and *-xdl2ncd*, but the tool also supports a *-report* mode that provides information on every wire and logic cell in the device.

Chapter 3

Theory

3.1 Overview

Before launching into a discussion of how one might build an autonomous system, it is appropriate to consider some underlying factors: dynamics, configurability, scalability, modeling, and domain. A considerable amount of this chapter is influenced by prior work, more fully presented in [1].

3.2 Dynamics

This work speaks of the *dynamics* of a computational system to denote the ways in which the system may change over time. A system that cannot change is taken to be static by definition.

If a system consists entirely of hardware and information, then any changes that it may undergo will have to affect its hardware or its information or both. Although unpublished prior work takes a much closer look at the nature and properties of information, the reader may safely interpret the term *information* in the Shannon sense, fully reducible to bits, and without regard for any meaning or semantics that it may have [52]. The reader may substitute *software and data* for *information* with little loss of generality. For the sake of discussion, it would be interesting to consider examples of each of the following:

- Information modifying itself
- Information modifying its hardware
- Hardware modifying its information
- Hardware modifying itself

The case of information modifying itself is a generalization of the familiar “self-modifying code” paradigm. Software is of course incapable of acting in any fashion whatsoever without the means of its underlying hardware substrate, so to be a little more precise, this “self-modifying code” should actually be described as “hardware modifying its information under informational control.” Information is likewise incapable of modifying its hardware, without the assistance of the underlying hardware substrate.

Just as information is incapable of independently modifying itself or its hardware, hardware is unlikely to modify itself or its information without some sort of informational control, if one is considering a computational system. In cases where the hardware really does function on its own without *any* informational control, even from incoming signals or stored data, then it will likely be limited to set and clear operations, neither of which is particularly useful in terms of computational capability.

Informally, one may say that *information is unable to act on its own*, and that *hardware is unable to think for itself*. But because the hardware and information function jointly as a system, there is reason to develop a slightly more realistic understanding of the possible dynamics:

- System modifying its information
- System modifying its hardware

It is obvious that the case of a system modifying its information happens all the time in computing systems: it is simply called *information processing*. As mentioned earlier, this is the very reason for the existence of computers. But the case of a system modifying its own hardware is rarely if ever seen, because changes to information are simple and inexpensive, but changes to hardware are complex and much more expensive.

While physicists seek to understand whether information is or is not physical—that is, whether it consists purely of matter and energy—one can at least consider how information may be physically encoded. In many cases, information is encoded with energy, be it through voltage levels in a flip-flop, or electrical charge on the floating gate of flash memory cells. In other cases, information is encoded with matter, including bumps and pits on the surface of an optical disc, or magnetic domain orientation at the surface of a magnetic disk. When matter is manipulated in such storage devices, it generally happens on very small scales, and generally occurs as a perturbation in the vicinity of some equilibrium point. Informally, technology uses energy to make small but detectable changes, but rarely assembles new structure in the sense of molecular self-assembly.

3.3 Configurability

The ability of a system to modify any part of its hardware or information would seem to offer the ultimate in configurability. To quote the creators of the Cell Matrix, “The cell matrix began

with a philosophical discussion on how far the hardware/software line could be pushed towards the software extreme” [53]. The Cell Matrix is a highly configurable architecture, that has pushed the hardware/software line quite far indeed, at least in simulation, but it cannot physically change its hardware. However, such changes will likely become possible in the future—the notion of physical changes is *forward-looking*, but no longer *far-fetched*—and it is worth considering what form they might take, or what form one would like for them to take.

This work considers four levels of hardware change that a system may undergo, in order of increasing flexibility that they afford:

Level 0 - No change: Purely combinational functionality—the system has no state capability.

Example: an analog radio receiver.

Level 1 - Transition within reachable state space: No new capability is added to the system.

Example: a finite state machine, with transitions controlled by data and/or instructions.

Level 2 - Transition to previously unreachable state space: No new capability is actually added, but the system’s state space appears to be larger. *Example: a peripheral link that makes additional states reachable once it has been configured.*

Level 3 - Expansion of existing state space: New actual or virtual capability is added to the system. *Example: molecular self-assembly or allocation of previously unused configurable fabric.*

Levels 0 and 1 are not particularly interesting in terms of configurability. Level 0 describes hard-coded functionality that cannot change. Level 1 describes normal computation, where data or instructions may cause the system to behave in a variety of ways, but with no new capability added.

Level 2 describes computation that includes unreachable states. For example, a desktop system might include a cryptographic coprocessor that can store some small amount of privileged information, and enable functionality that was previously unavailable. In this case, the extra functionality was present all along, even though the system was unable to access it. It is not a matter of functionality being added to the system, but simply of becoming available to the system.

Level 3 describes systems that can be expanded in some fashion. It is important to distinguish between *actual* and *virtual* changes:

Actual. The system physically changes. This requires a change in the properties or structure of the hardware. An example of property changes that extend state is the effect of periodic potentials within a semiconductor crystal lattice. While a single atom has a small number of discrete energy levels available to its electrons, an atom within a lattice instead has energy bands that comprise a near continuum of energy levels. Informally, the periodicity of the crystal affords the system degrees of freedom that its constituent atoms do not possess on their own.

Examples of structural changes to hardware necessarily involve the motion of matter. At a microscopic scale, this may become achievable with the help of nanotechnology and molecular self-assembly¹, with promising ideas and work involving artificial matter [54] or programmable matter—more popularly known as *claytronics* [55]. At a macroscopic scale, this could be achievable with the help of robotics. The mention of robotics really only serves to make a point, and is not particularly desirable in practice. On the other hand, anything that can manipulate or rearrange matter at very small scales is both desirable and likely to become a reality, thanks to aggressive research in key areas.

Virtual. The system may not physically change, but its configuration changes in such a way that its virtual state space is extended. This clearly requires some form of configurable logic, such as an FPGA or a Cell Matrix, and is analogous to a software program instantiating new objects through dynamic memory allocation. Subject to memory footprint restrictions, the program is essentially configuring part of the state space that was available to it all along, but from an internal perspective, the program has just grown or expanded in some fashion.

To prototype systems that undergo physical changes would require a large multidisciplinary effort since the necessary technology is not yet mature. But the system’s view of the changes is similar whether the changes are actual or virtual, because it can abstract the configuration changes away from the underlying mechanisms. It should therefore be possible to study such systems through virtual state space expansion in configurable logic.

The distinction drawn between actual and virtual changes does not mean that the two are mutually exclusive. Even virtual changes require some kind of change in the underlying physical state. But actual changes as discussed here are taken to be much larger in scope, in the sense of building something that did not previously exist.

3.4 Scalability

When considering large systems and architectures, it is important to consider how well they scale, and what obstacles impede the scaling. In the past, it was possible to treat certain components as if they were ideal. For example, it was safe to assume that signal propagation delays were negligible. That assumption is no longer safe in the context of higher clock frequencies and larger die sizes.

Scaling systems to larger sizes is usually achieved by putting more components together and by making the existing components smaller. But both of these facets introduce difficulties that must be understood and addressed. In broad terms, the obstacles to making components smaller, other than the difficulty of fabrication, are noise and reliability issues: smaller components are more

¹The Phoenix project aims to implement programs directly in hardware, and expects CAEN (Chemically-Assembled Electronic Nanotechnology) technology to build reconfigurable hardware with densities on the order of 10^{11} switches/cm² [8].

fragile and more easily disrupted by mechanical or thermal mechanisms or by radiation, and their signals are more easily lost in the underlying noise. The obstacles to putting more components together, are distribution issues: each component generally requires power, data, and clock signals, and generates heat that must be removed to guarantee proper operation and prevent damage. There is also a significant organizational issue in large systems, affecting component placement and connections, and complicating any attempts to optimize them.

One can illustrate these issues with well known examples from current CMOS VLSI technology: Wires can take up more die area than the logic that they serve.² Signal propagation times have more or less halted the continued growth of single-core CPUs. Size and propagation issues limit memory bandwidths, and thus the speed at which computation can occur. Heat generation constrains device geometries, and impedes the desired 3-dimensional layering of structures. Shrinking gate sizes dictate that storage components are more susceptible to radiation effects. In many cases, this requires increasingly complex fixes and workarounds, including multicore CPUs, multilevel cache hierarchies, heat sinks and cooling mechanisms that are far larger than the components that they serve, and pervasive use of error correction techniques. These issues are realities that must be considered in all large systems:

Storage:

- Access times are non-negligible.
- Memory bandwidth is limited.
- Storage elements are susceptible to radiation effects.

Connection:

- Signal propagation times are non-negligible.
- Wires are large and numerous.

Logic:

- Switching times are non-negligible.
- Circuits are susceptible to noise.
- Generated heat must be dissipated.

Some of these limitations can be mitigated to a certain extent, and surprisingly some are even reduced at smaller sizes [56], but the limitations cannot be eliminated entirely. It might be desirable to build an extraordinarily dense supercomputer for everyone's desktop, but scaling limitations instead point in the direction of distributed and decentralized computing.

²It is true that this effect may be reduced or eliminated with the availability of additional metal layers.

3.5 Modeling

To change any system in an intentional manner, it is first necessary to know the system's current state or configuration, or at least the relevant subset thereof. Without access to that information, changes would be no more than undirected perturbations, and would at best be based on guesses or random decisions. Exploration in the presence of limited information can be quite valid, but is not the focus of the present work. The remainder of this section will use the term *model* to refer to the appropriate state or configuration information necessary to support intentional changes.

Models are a much more common part of our everyday lives than we often realize. Consider a person opening a door and walking through it. That person didn't need to be told that the door would swing on hinges to create a spatial opening, because their mental model of the door already contained that information. In the same way, even though they likely saw only the door handle and not its internal mechanism, their model told them how to operate the handle to open the door. Without the benefit of a model, an encounter with a door would require experimentation to successfully pass through it. And without the ability to abstract the information and learn from experience, every single door encounter would be a brand new challenge for our unfortunate candidate.

On the other hand, models need only carry enough information to support the desired actions. To drive a car, most drivers need only a basic understanding of how to start the car, how to steer it, how to apply the gas, and how to apply the brakes. A comprehensive understanding of internal combustion engines and powertrains is not necessary for a trip to the store. But a number of other situations, ranging from maintenance to racing, do require additional detailed knowledge. A more complete model will generally provide more flexibility and support more complex changes, but too much complexity—or at least the inability to work at an appropriate level within that complexity—would make every action an impossibly difficult task: trying to understand the operation of a door at the atomic level, for example, would likely be many orders of magnitude too complex for the collective computing power of every computer ever built.

The correctness of models is just as important as their completeness, because an appropriate decision based on incorrect information may result in unexpected or undesirable changes. Mistaking the gas pedal for the brake can quickly turn a routine action into a major problem. In the case of a system that undergoes changes, the model must change accordingly, to preserve correctness and to remain in sync with the system that it describes. Subsequent changes to any portion of the model that is out of sync should be deferred until the re-synchronization is complete.

When the model resides within the system that it is modeling, correctness and synchronization become even more critical for the continued operation of the system. In particular, a system that changes itself can quickly kill itself if it interferes with its own operation, a point that was clearly demonstrated in the implementation phase of this work.

There are therefore a number of properties and capabilities required of models that support autonomous changes:

Completeness: The model must provide sufficient information about any aspect of the system that is subject to change.

Resolution: The model must provide an appropriate level of detail; too much and the evaluation will be too slow, too little and it will be incomplete.

Correctness: The model must accurately reflect the state and/or configuration of the system.

Synchronization: The model must remain synchronized with the changes that the system undergoes.

It is important to note that system models do not necessarily have to be centralized. A sports team can work towards a common objective, even though each member's information is a subset of the collective information. While a team captain might provide general direction to individual team members, each team member still maintains responsibility for their own fine-grained movement. In the same way, a large distributed system would be likely to work towards a common objective, but smaller subdivisions of that system would most likely assume responsibility for their low-level internal operation. This would be especially sensible for homogeneous groups within heterogeneous systems.

The approach used in this work considers the model to be authoritatively correct, and ensures synchronization by propagating the model's state to the system state. At sufficiently elevated levels of complexity, however, that approach may no longer suffice, and it may become necessary to feed state or configuration information back into the model. This may be necessary because of external constraints or changes that prevent the system from functioning as intended. If a person is walking along a rocky path and a rock unexpectedly shifts under them, their limbs and center of balance may suddenly be out of sync with what was expected, and updated information may need to be fed back into the model to prevent or control a fall. It would be quite humorous for the person to resume their walking motion while lying flat on the ground, simply because they were not conscious of the change.

Feedback may also be necessary in cases where the model is incomplete or where the system must interface with some part of its environment. An animal generally does not know the exact distances to each object in its field of view, but instead has some broad estimates which are adjusted by successive approximation as it moves around. In the absence of feedback, changes must be based on exact data, but the availability of feedback makes a system much more flexible and able to function on more generalized information.

3.6 Domain

Although hardware and software are distinctly different domains, they do share some important similarities, and there are many cases where a given computation can be performed in either one. Deciding how to partition a computation between hardware and software is a difficult problem, and the subject of considerable research. In general, that decision is only simplified when performance constraints or complexity clearly favor one domain over the other.

Software offers significant amounts of flexibility and portability, and provides a framework for breaking down complex computations into algorithms consisting of simple instructions. As such, software is a familiar and well understood domain to many designers, but there are also some things that software cannot do well or cannot do at all. One must generally turn to hardware whenever there are stringent performance or timing or concurrency requirements, or whenever it is necessary to interface with other hardware—as previously explained, software necessarily requires a hardware substrate.

Part of the appeal of implementing autonomy in the hardware domain is that nothing precludes the layering of software autonomy on top of it. Hardware autonomy can therefore augment the software-centric work of Ganek and Corbi, and the International Conference on Autonomic Computing [9, 57]. In contrast, software autonomy by itself encompasses only a subset of the flexibility and capability of combined hardware-software autonomy.

A related consideration is the granularity at which to implement a computation. In the software domain, that range extends from machine code through high-level languages. In the hardware domain, it extends from transistor level design to behavioral hardware description languages. The lower levels provide enhanced performance and compactness, at the expense of flexibility, maintainability, portability, and design simplicity. The tradeoffs are complex, and often lead to the use of finer granularity for critical components, while relying on coarser granularity for the bulk of the system. This ties back into the modelling discussion, but the point of interest is that a fine-grained hardware implementation does not impede use of software within the same system, and the partitioning can in fact retain some fluidity.

3.7 Observations

One long term goal of this work is to provide in autonomous systems the same kind of performance that would be obtained with current design flows. The intent is to change the way that hardware systems are designed, from an offline *monolithic* approach, to an online *component* approach, but to do so with minimal performance and size penalty.

To obtain such high quality results, it will be necessary to provide tools comparable to the current ones, and to integrate them within the target systems. This will require access to comparable

algorithms and the same data as existing tools. In particular, the resulting systems will need access to their timing data, so that placement and routing can meet real world constraints, a requirement that no serious hardware designer can do without. There is little theoretical obstacle to these goals, but a large amount of practical development will be necessary.

In retrospect, the intention is really not to make “a thinking machine” as Hillis proposed, but simply to develop a machine that can take responsibility for its own capabilities and resources, and that can absorb as much of its complexity as possible behind a simpler and more flexible interface. Supposing the ability to manufacture such machines, one can envision a large distributed network of configurable systems that could function independently or collaborate with each other, each one doing so autonomously.

Although this research is interested in the possibility of a system developing and learning over time, the system should learn in a deliberate and purposeful manner in response to actual conditions that arise. This differs from Ganek and Corbi’s philosophy that a system should always be trying to optimize itself.

The systems discussed here can be considered autonomous, and can even be considered self-aware in that they model themselves, but that does not make them conscious. They are still machines that follow rules, even if that sometimes includes newly deduced rules. These machines in many respects fit DARPA IPTO’s notion of cognitive systems, and it seems probable that the remaining aspects could be developed on top of this work.

Chapter 4

Roadmap

4.1 Overview

The steps leading to the development of hardware autonomy are surprisingly logical. And although each one is within the reach of current technology in one form or another, integrating the various pieces, and embedding them inside the system that they target can be difficult. The roadmap presented here therefore aims to show how we might wrestle autonomous systems into existence.

This chapter considers autonomous hardware systems in a very general sense, with few assumption about their underlying technology or architecture. Although Chapter 2 discussed tools and capabilities such as synthesis, mapping, placing, and routing, in the context of FPGAs, the reader is reminded that almost every complex electrical system is designed with such tools.

Any design that begins at a behavioral level must be translated to an appropriate circuit description, and mapped to the kind of building blocks that the target architecture provides. These steps are typically called *synthesis* and *mapping*, whether targeting FPGAs or ASICs or board components or nanotechnology. The form of the circuits may differ drastically, but the design tools would nevertheless perform similar kinds of operations. The same is true for *placing* and *routing*. In the most general sense, these steps simply decide where to put the building blocks, and how to connect them—steps once again essential for everything from circuit boards to nanocircuits. And a *configuration port* is simply an access point to deliver a design to a configurable device.

Even the simplest electrical designs require synthesis, mapping, placing, and routing, even when they only take place in the designer’s mind. Furthermore, analogous principles apply to processes from quantum computing to biology, points that are corroborated in [1]. The reader is therefore encouraged to dissociate the terms *synthesis*, *mapping*, *placing*, *routing*, and *configuration port* from the FPGA context of Chapter 2 during the remainder of this chapter. The question at hand is this: How would we design an autonomous hardware system?



Figure 4.1: Levels of autonomy.

4.2 Levels of Autonomy

This chapter presents a hierarchy of autonomy levels—from low to high—where each level builds upon the capabilities of previous levels. Each level represents a logical subdivision of the larger picture, introducing some distinct capability that underlying levels do not possess, while still resulting in a fully functional system.

Figure 4.1 depicts a span ranging from Level 0 systems that have no autonomy, to Level 8 systems that are able to learn from their own experience. This range is quite broad, but each of the levels that span it are attainable, and collectively constitute a roadmap to guide both the present work and the development of future autonomous systems.

4.2.1 Level 0: Instantiating a System

Level 0 systems are not autonomous. They are unable to adapt to changing conditions in themselves or in their environment. While the designer may be able to update them, and perhaps even do so remotely [58], the systems are entirely passive with respect to the update. They do not know what is happening, and they have no control over it.

Scenario: A device is remotely reconfigured with a full or partial bitstream.

Requirements: a configuration port.

Only Level 0 systems accept full configurations. In all subsequent levels, applying a full configuration is analogous to erasing a person’s memory and reprogramming it. The person would probably not survive the shock, and would in any case have lost any autonomy in the matter. *Full configurations are antithetical to autonomous systems.*

A possible exception is the identity configuration, which reconfigures the device exactly as it was, and does not violate autonomy, but is more properly likened to sleep or unconsciousness. A device can deliberately save and restore its full configuration if it so chooses, as it might do if it detects that its power source has failed and that it is operating in battery backup mode.

4.2.2 Level 1: Instantiating Blocks

Level 1 systems play a minor role in their own reconfigurations. They have some notion of their own utilization and free space, typically divided into slots. They have the ability to accept a partial configuration, to decide which slot to place it in, and to make simple connections that may be required. The configuration may have been provided remotely, or may have been retrieved by the system from its own storage.

Scenario: A device reconfigures itself with a partial bitstream.

Requirements: an allocation map and a primitive router.

4.2.3 Level 2: Implementing Circuits

Level 2 systems maintain an internal model of themselves that is more sophisticated than the slots of the Level 1 allocation map. Instead of blindly accepting a pre-generated configuration, they accept netlists that are compatible with their architecture and already mapped to it. Level 2 systems are responsible for placing and routing the netlists, and for updating their internal model. They have much more control over their own use of resources, and the netlists that they accept are likely to be compatible across an entire family of devices, instead of only being compatible with a specific device.

Scenario: A device extends itself by instantiating a netlist.

Requirements: an internal model, a placer, a router, and a configuration generator.

The quality of the placer and router will significantly impact the performance of the updated system. Much research and development has gone into the routing and placing tools used by the industry, and the idea that such tools could fit *inside* a device, as part of its overhead, may seem somewhat unlikely. But if the road map presented here is accepted by the industry and the research community, there will be compelling motivation to port or adapt such tools to the devices that they serve. Furthermore, the overhead costs become more acceptable with increasing device sizes. In cases where there are multiple configurable devices functioning as a single autonomous system, there is good reason to consider developing *distributed* place and route tools, allowing each device to take responsibility for its piece of the puzzle.

It is known that placing and routing are NP-complete problems, so one may begin to imagine significant delays while such activities take place. Fortunately, in the case of an autonomous system, one would frequently be dealing with *incremental* placing and routing in some confined portion of the device, with a well specified area constraint. In short, one may expect a simpler problem than most tool flows have to deal with, because the problem involves only a component rather than the entire device, and because the search space that the tools see is limited by the fact that much of the device may already be in use. The device can continue to function normally while the configuration changes are being prepared, but maintaining glitch-free operation during the actual reconfiguration is a slightly larger challenge, as explained later.

4.2.4 Level 3: Implementing Behavior

Level 3 systems can accept operations or behaviors and synthesize and map them into netlists. Not only are they responsible for deciding where to place components and how to connect them, but they are also responsible for determining how they want to implement those components, possibly depending on current system resources and conditions.

Instead of requiring netlists, Level 3 systems can accept some kind of abstract hardware description. On one hand, this means that these systems can cache and remember configuration steps in a very compact form, possibly re-synthesizing them in the future to account for changing conditions. On the other hand, such hardware descriptions should be suitable for almost any target device, and not just a certain technology family. Clearly this progression of systems leads to greater internal complexity, but also allows greater external simplicity and generality.

Scenario: A device extends itself by implementing a component from a hardware model.

Requirements: a synthesizer and a technology mapper.

The versatility of abstract or generic components allows systems to model themselves in terms of behavior if that proves to be useful. It also paves the way for them to describe themselves to heterogeneous peers, possibly to replicate or transfer themselves before going off-line, as might be desirable in the power failure situation. In a sense, the functionality begins to take a predominant role, and that functionality can theoretically be implemented in any system that has available configurable space, though practical systems would typically be constrained by the resource or network requirements of their clients or servers.

4.2.5 Level 4: Observing Conditions

Level 4 systems have the ability to detect and monitor conditions of interest. These systems require some kind of centralized or distributed controller that can observe or receive notification of changing conditions. The systems may include a variety of distributed sensors, and may also include a prioritizer or classifier that assigns relative importance to conditions that are detected.

Scenario: A system notices that the bit error rate on some incoming link has risen above its norm.

Requirements: a monitor, a collection of sensors, and a prioritizer.

Level 4 systems are able to detect changing conditions without necessarily being able to respond to them. While the mere ability to detect conditions may not appear to be especially useful, it plays a key role in autonomy, because it motivates change. Changes in living systems are not haphazard, but are instead purposeful and directed, which means that living systems change according to internal or external goals and conditions. This work do not view the theory of evolution and its

emphasis on mutation as a useful model for autonomy, nor does it embrace Ganek and Corbi's argument that systems should perpetually try to optimize themselves. Change in an autonomous system should happen for a reason.

Although Level 4 systems can still be designed and analyzed in a traditional manner, it becomes apparent that artificial intelligence can contribute desirable properties, and that will become increasingly true for progressively higher level systems. At the present level, this work foresees artificial intelligence being most useful in the prioritizer or classifier, because it can be difficult to understand the meaning or importance of a condition from raw sensor input. Detecting the fact that a system component has exceeded some critical threshold temperature is not difficult, but trying to determine whether an anomalous rise in error rate from a peripheral constitutes a warning of impending failure may be more difficult.

Just as with living systems, conditions to be monitored may be internal or external to the system. It may be just as important to know that an external link is about to be lost as it is to know that an internal FIFO is nearly full, and determining the severity of the conditions may depend partly on whether the two are correlated. In addition, large systems are probably best served by distributed sensors and monitoring, similar to what one encounters in biological systems. The sensors are the "eyes and ears" of the system, because they allow it to observe itself and its environment.

4.2.6 Level 5: Considering Responses

Level 5 systems include a collection of responses or solutions that can be applied to address the conditions that they observe, as well as some kind of evaluator to help determine whether any particular solution is desirable. These systems merely look for candidate responses, and do not necessarily have the ability to apply the responses.

Scenario: The amount of data coming in from a certain link is approaching the system's ability to process it, and the system looks for an appropriate solution. Possible solutions include dropping packets, forwarding some requests to a neighboring system, re-implementing certain functions with faster circuits at the expense of added area, duplicating the existing circuit and splitting the data between them, informing the sender of the impending problem, running statistical analyses to estimate whether the incoming rate is a temporary surge or a sustained increase, monitoring the situation more closely but doing nothing, and so on.

Requirements: a response library and response evaluator, and possibly a simulator.

A Level 5 system may or may not find a suitable response in its library, and it uses the evaluator to help make that determination. In some cases it may be desirable to include a simulator that can simulate candidate responses, to determine which ones will in fact resolve the problem, so that the evaluator can compare their costs and performances. It is worth noting that because Level 3 and higher systems can synthesize operations and behaviors, the response library can be much more

compact than if it had to store device configuration bitstreams. The drawback is the fact that synthesizing, mapping, placing, and routing require longer execution times than just feeding in a pre-generated bitstream.

4.2.7 Level 6: Implementing Responses

Level 6 systems have the ability to apply responses or to request assistance. This constitutes the first stage of real autonomy. If a suitable response has been identified, the system can apply it by synthesizing and implementing the corresponding behavior. If no suitable response is identified, the system can instead query neighbors or peers or supervisors or other resources to see if any of them can supply a response.

Scenario: Having identified a low memory condition and having determined that some of its unused resources could be used to supplement that memory, a system reconfigures that part of itself and connects the part to itself.

Requirements: an implementer and a requester or notifier.

If a system has identified a suitable response, and has applied that response by synthesizing and implementing the corresponding behavior, it will have autonomously initiated and completed a change to itself.

It may seem strange to consider a request for assistance to be a demonstration of autonomy, but it can in fact be seen as evidence of the system looking for a response and realizing that it has none, and therefore deciding to look outside itself. Biological interactions include many examples of requests for assistance large or small, and in particular, people use this approach extensively. A system that asks for help may actually be able to *learn* by adding to its knowledge, much as a student can learn from a teacher or a tutor or a peer or a book. In this context, the request could be sent to neighbors, to supervisors if there are any, to known external resources equivalent to public libraries, or even to the system designer.

4.2.8 Level 7: Extending Functionality

Level 7 systems have the ability to infer the required behavior from detected conditions, and/or to adapt existing behavioral pieces to the conditions. This level of system falls fairly unambiguously into the artificial intelligence domain: it knows how to make inferences, and it knows how to dissect known behaviors and piece them back together as something that it had no previous knowledge of. This is a clear case of machine learning.

Scenario: A system determines that it can build a necessary behavior out of smaller existing components, and it implements the new behavior.

Requirements: an inference engine.

If the system cannot identify a suitable response, but if it knows enough about the problem, it may try to infer the function on its own, or to piece it together out of other components. If necessary, the system can simulate the problem before implementing it.

4.2.9 Level 8: Learning From Experience

Level 8 systems can analyze the results of implementing changes, and can determine whether their internal libraries should be updated or optimized for future use. If one assumes a strong artificial intelligence capability, then it is possible to envision a system that learns from experience and mistakes: it adds to its knowledge when applicable, and possibly corrects or replaces previously held incorrect or incomplete knowledge.

Scenario: Having re-implemented some of its functionality for performance reasons, a system analyzes the results and makes adjustments to its response library.

Requirements: an assessor.

4.3 Issues

The significant increase in flexibility that accompanies autonomous systems, requires consideration of policies, security, and testing. These issues are complex even in static contexts, but in the dynamic context of the present work they become far too complex to broach, and are best left to true experts in the respective fields.

It is well known that security and testing are much more effective when designed into a system from the ground up, and trying to implement a complex autonomous system without addressing these issues, is a virtual guarantee that the system will fail prematurely.

4.3.1 Policies

Policies or guidelines are important because this work is not trying to study artificial life—it has little interest in how systems evolve if left to themselves, as considered by Kubisch and Hecht [38], or how to continually optimize systems for performance and availability, as considered by Ganek and Corbi [9]. This work instead looks for the systems to perform useful work, and to do so according to priorities and metrics of the user’s choosing. While much research seems interested in autonomy for the sake of studying evolutionary behavior, this work takes a pragmatic approach and asks systems to simplify actual computational and behavioral problems.

4.3.2 Security

Security is just as important an issue, particularly for systems that support behavioral synthesis. The likelihood of denial-of-service attacks grows with system autonomy, and the emergence of hardware viruses is assumed as a given. One can expect autonomous computing systems to be targeted in the same way as autonomous living systems, and in all likelihood systems will eventually require something resembling an immune system to protect themselves.

4.3.3 Testing

Testing will become an even more complex problem than it already is, with systems undergoing internal changes that the designer doesn't know about, and has never even foreseen. Although the individual functions and operations that are provided to the system will presumably have been behaviorally verified by their designers, their internal implementations and collective behavior will not only be untested by formal methods, but will also be subject to change over time.

Implementation

The next four chapters discuss the implementation of a Level 6 autonomous demonstration system. The system is based on the XUP Virtex-II Pro board, running Linux and custom software. Most of the platform selection considerations from the original proposal are presented in Appendix B. Of the options available, it seemed most appealing to use a Virtex2P board, and run Linux on one of its embedded PowerPCs. Any of the candidate architectures would have required a considerable amount of software tool development, and this one was no different.

The resulting system is able to observe external inputs, and to respond to changes by configuring connected devices, or by dynamically implementing or removing hardware functionality. The system displays its resource usage in real-time on a Graphical User Interface (GUI), and also allows the user to interactively mask out resources, thereby injecting simulated defects. Subsequent placements avoids using any masked resources, as a demonstration of defect tolerance in large FPGAs.

The organization of the chapters is as follows:

Chapter 5: Hardware

Describes the static hardware of the system, both inside and outside the configurable logic.

Chapter 6: Software

Describes the underlying system software, including Linux and custom drivers, and Java.

Chapter 7: System

Describes the autonomous server and the autonomous controller, and their interaction with the rest of the hardware and software.

Chapter 8: Operation

Describes the overall startup, operation, and performance of the system.

The static hardware of Chapter 5 provides the Virtex2P Internal Configuration Access Port (ICAP). This port is able to accept configuration bitstreams and forward them to the internal configuration controller, and its availability secures Level 0 capability. Also in this chapter are input sensors that allow the system to observe its environment. When coupled with Chapter 6 and Chapter 7 software, these sensors provide Level 4 capability.

The autonomous server software of Chapter 7 allows the system to model itself, and to parse, place, route, configure, and dynamically implement circuitry, thereby providing Level 2 capability. Finally, the autonomous controller allows the system to respond to conditions that it has detected, typically passing directives to the autonomous server, for Level 6 capability. Level 5 capability is also present, but only in cursory form—to build a more complete Level 5 implementation would be a reasonably simple software development task, but would not add anything truly novel in the immediate context.

Chapter 5

Hardware

5.1 Overview

This chapter describes the static portion of the hardware for the autonomous demonstration platform. It presents an overview of the base hardware and its capabilities, and discusses certain key components in greater depth.

Figure 5.1 depicts selected portions of the system hardware. Surrounding the FPGA itself are audio inputs and outputs, video inputs, keyboard and mouse inputs, and an SXGA display output, as well as the system's main memory. The FPGA encompasses IOBs, two PowerPCs, and a large amount of configurable logic. The configurable logic serves as the substrate for both the base or static portion of the system hardware, and for the dynamic circuits that may be implemented and connected inside of it.

The PowerPC processor communicates with connected components primarily through the PLB bus.¹ Those connected components include a custom combined framebuffer and VGA controller, a DDR RAM memory controller, and a bridge to the OPB bus.² The OPB bus in turn connects to the interrupt controller and to a variety of other components not shown in the figure.

Directly interfaced with the audio, video, mouse, and keyboard ports, are *sensor stubs* that monitor the ports for activity. These stubs are able to configure the necessary interface chips, and to generate interrupts when signals are acquired or lost. The interrupts are fed through the interrupt controller to the PowerPC processor, and are in turn dispatched to the appropriate software drivers.

Figure 5.1 also shows two dynamically instantiated cores in red: the AC97 FFT core and the video capture core. The system is in fact able to implement arbitrary logic, but these two cores were

¹64-bit 100 MHz IBM CoreConnect Processor Local Bus (PLB).

²32-bit 100 MHz IBM CoreConnect On-Chip Peripheral Bus (OPB).

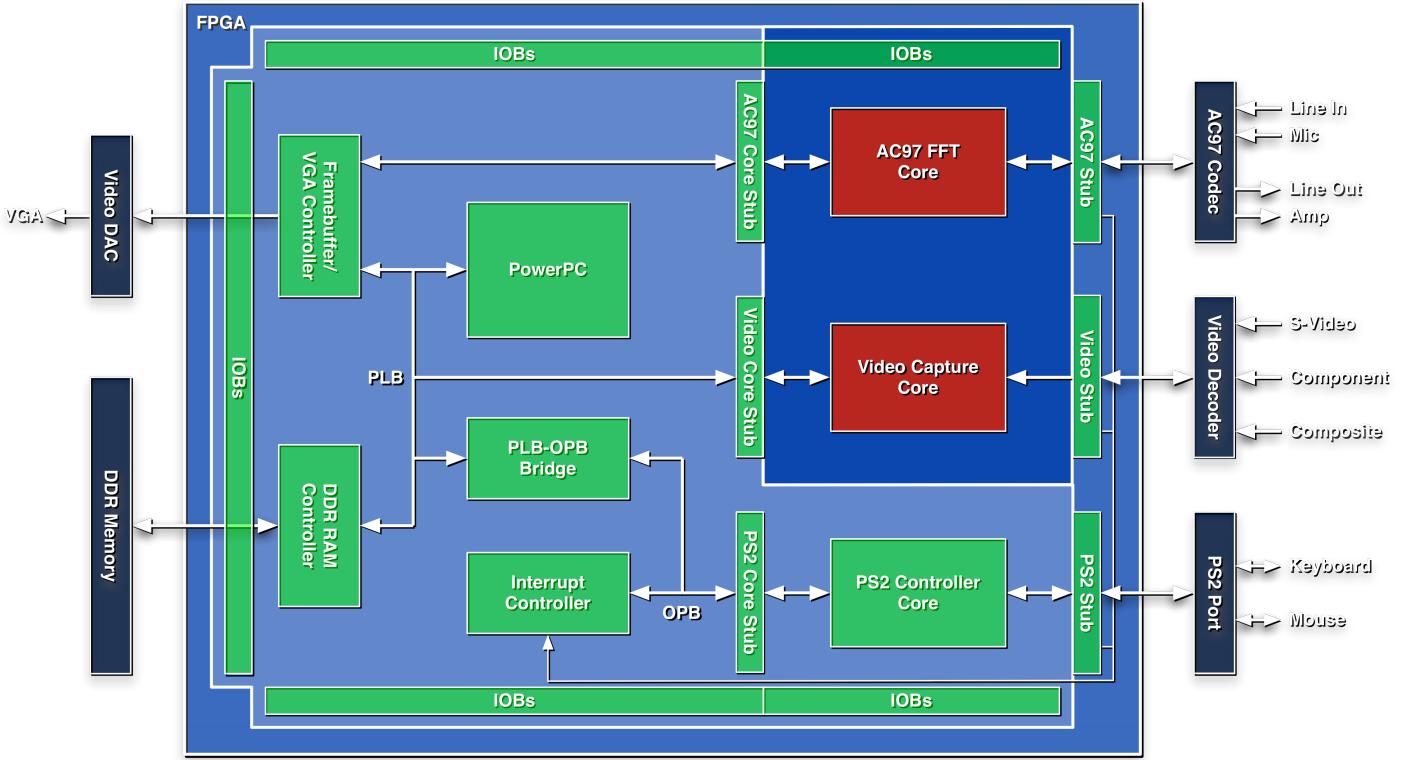


Figure 5.1: Demonstration system: high-level view.

specifically selected for demonstration purposes. The system also provides *interface stubs* that dynamically instantiated circuits can connect to. These stubs terminate in ports similar to half of a *bus macro* in the Xilinx Partial Reconfiguration (PR) flow. They form incomplete static circuits that would normally be trimmed by the synthesis or mapping tools, but are constrained to prevent trimming, and to expose named ports to the dynamic system. The objective is simply to constrain dynamic inputs and outputs to known locations, so that the runtime router is able to connect them as desired. Further details are provided in the remainder of this chapter.

Most of the base functionality was developed with Xilinx EDK 8.1 and ISE 8.1, for the sake of synthesis, mapping, placing, routing, and bitstream generation. Most of the custom cores were run through the same EDK/ISE process, except that synthesis was performed with Synplicity's Synplify Pro 8.4. Synthesizing the custom cores to EDIF³ from inside Synplify Pro ensured that they would remain as similar as possible, whether instantiated in the static design for development purposes, or instantiated dynamically to demonstrate system autonomy.

³EDIF is the Electronic Design Interchange Format—perhaps the most broadly accepted netlist format available.

5.2 Board Components

The requirements of the demonstration platform—with their implications and effects on design decisions—were thoroughly explored in the research proposal, and are reproduced in Appendix B. This section will consequently present the resulting system capabilities as a *fait accompli*, without revisiting the choices and dependencies.

The demonstration board is the XUP board described in Section 2.4, along with the VDEC1 video capture daughter card [59]. In broad strokes, the two boards include the following peripherals and capabilities:

Xilinx XC2VP30: Virtex2P FPGA with two PowerPC 405 hard cores⁴ and a large amount of configurable logic.

512 MB DDR SDRAM: Memory for Linux, drivers, user applications, and the SXGA frame-buffer.

IBM 1 GB MicroDrive: CompactFlash accessible through Xilinx SystemACE controller to support FPGA configuration, Linux bootstrap, and Linux filesystem.

Analog Devices ADV7183B: Video decoder chip capable of capturing NTSC, PAL, or SECAM video from component, composite, or S-Video inputs. Currently used to capture 24-bit NTSC 720×487 interlaced video at 60 Hz.

Fairchild FMS3818: Triple 180 MHz Video DAC. Currently configured for 1280×1024 resolution at 60 Hz, with a 108 MHz pixel clock.

National Semiconductor LM4550: AC ’97 audio codec, supporting 18-bit stereo recording and playback at up to 48 kHz, with a 12.288 MHz bit clock.

10/100 Ethernet MAC/PHY: Standard networking connection for storage or communications.

Dual PS/2 Port: Connection for standard mouse and keyboard.

RS-232 Serial Port: Console serial-link connection.

Other Connectors: SATA,⁵ USB,⁶ expansion connectors, micro-coax connectors, and more.

5.3 FPGA Components

Additional standard and custom logic is configured into the FPGA to support both internal and external peripherals. Some of the key components are:

⁴One of the two PowerPCs runs Linux, while the other one is available for dynamically instantiated circuits. At present, EDK automatically ties each of the second PowerPC’s inputs to one of the power rails, but those connections could easily be removed at runtime if necessary.

⁵Unfortunately, the connection does not comply with the Serial-ATA (SATA) standard, making true SATA compatibility impossible.

⁶The USB port is not accessible by the FPGA. It was provided by Xilinx as a practical and less expensive alternative to their Parallel Cable IV.

PLB Bus (plb_v34): 64-bit 100 MHz IBM CoreConnect Processor Local Bus (PLB). High-speed bus connected directly to the PowerPC processor.

OPB Bus (opb_v20): 32-bit 100 MHz IBM CoreConnect On-Chip Peripheral Bus (OPB). Peripheral bus bridged to the PLB bus.

DDR SDRAM Controller (plb_ddr): PLB controller for external DDR SDRAM.

RS-232 UART (opb_uart16550): Standard UART. Trial version of Xilinx IP.

10/100 Ethernet MAC (plb_ethernet): Standard Ethernet MAC. Trial version of Xilinx IP.

Interrupt Controller (opb_intc): Interrupt controller for PowerPC. Can be extended to accept interrupts generated by dynamic logic.

OPB ICAP (opb_hwic平): Controller for FPGA ICAP port, permitting partial internal dynamic reconfiguration.

Framebuffer (custom): PLB framebuffer for the SXGA output. Supports 24-bit fixed 1280×1024 resolution at 60 Hz (5 MB/s) from DDR SDRAM. Provides cursor functionality to interface with mouse driver. Provides FFT display and dual-ported BRAM to interface with dynamically instantiated audio FFT core.

SystemStub (custom): OPB IP Interface (IPIF) available for processor communication with dynamically instantiated circuits.

VideoCapture Stub (custom): PLB video capture stub. Provides dual-ported BRAM to support capture of 24-bit 720×487 interlaced video at 60 Hz into DDR SDRAM. *This is only a stub. The video capture core is dynamically instantiated.*

PS2 Core (custom): PS2 core connecting to PS2 stub, to support mouse and keyboard.⁷

Audio Stub (custom): AC'97 stub to monitor audio inputs and generate interrupts when audio is acquired or lost.⁸

PS2 Keyboard Stub (custom): PS2 stub to monitor keyboard port and generate interrupts when keyboard is connected or disconnected. Uses OpenCores.org PS2 interface [60].

PS2 Mouse Stub (custom): PS2 stub to monitor mouse port and generate interrupts when mouse is connected or disconnected. Uses OpenCores.org PS2 interface [60].

Video Stub (custom): Video stub to monitor video decoder and generate interrupts when video is acquired or lost. Uses OpenCores.org I²C controller [61].

GPIO Cores (opb_gpio): Interface to board LEDs, DIP switches, and pushbuttons.

A few comments are appropriate with respect to the Xilinx opb_hwic平. The Xilinx hwic平 core is an OPB interface to the device's internal ICAP, and a means to modify or read back state and configuration information. In the system developed here, the ability to perform partial dynamic reconfiguration through the ICAP is absolutely essential.

⁷Adapted from opb_ps2_dual_ref.

⁸Adapted from opb_ac97.

5.4 Internal Communication

Most of the static communication within the FPGA happens over the PLB and OPB busses, both of which support multiple masters and slaves. Communication with dynamic circuits is instead provided through stubs of various kinds, whose signals are exposed for convenience, much as if they were internal input/output ports.

The 64-bit PLB and 32-bit OPB busses both run at 100 MHz, resulting in theoretical maximum bandwidths of 800 MB/s and 400 MB/s respectively. In practice, typical transaction overheads run near 10 to 20 cycles per transaction, which makes individual reads and writes very costly. Bursting reads and writes are much better at amortizing the transaction overheads, and are therefore the method of choice for certain custom cores presented later in the chapter.

Stubs for communication with dynamic circuits can be implemented in a variety of ways. In the first case, nearly any kind of signal can be passed through a specially configured slice that has a well-defined location and an enable control. This is a lightweight way of interfacing dynamic circuits with the system, and is explained in more detail below. Only clock signals are poorly suited for this approach, because taking them off the clock tree usually introduces excessive skew. In the second case, custom or standard busses can be passed through collections of stub slices, allowing for any kind of dynamic core to be interfaced with the system.⁹ And in the third case, data can be pushed through dual-ported BRAMs, a method that is particularly well suited for data which crosses clock domains.

A Level 4 or higher autonomous system also needs some means of observing itself or its environment, and of detecting conditions that arise, as explained in Section 4.2.5. On the hardware end, that capability is provided by stubs that observe device inputs while passing those inputs along to static or dynamic circuits, a case more fully explained in subsequent sections.

5.5 Framebuffer

A significant number of IP cores were developed to assemble the full demonstration system. They include sensor stubs that observe system inputs, interface stubs that permit connections between static and dynamic hardware, and cores to support the visualization goals of the system.

A significant number of IP cores were developed to assemble the full demonstration system. The Framebuffer core reads pixel data from memory, generates video timing signals, and drives the 1280×1024 SXGA output. It can also display a square cursor that serves as a mouse pointer, and it can display frequency bands with data supplied by a dynamically instantiated FFT core.

⁹Although PLB and OPB stubs were developed, they are not presently being used within the system.

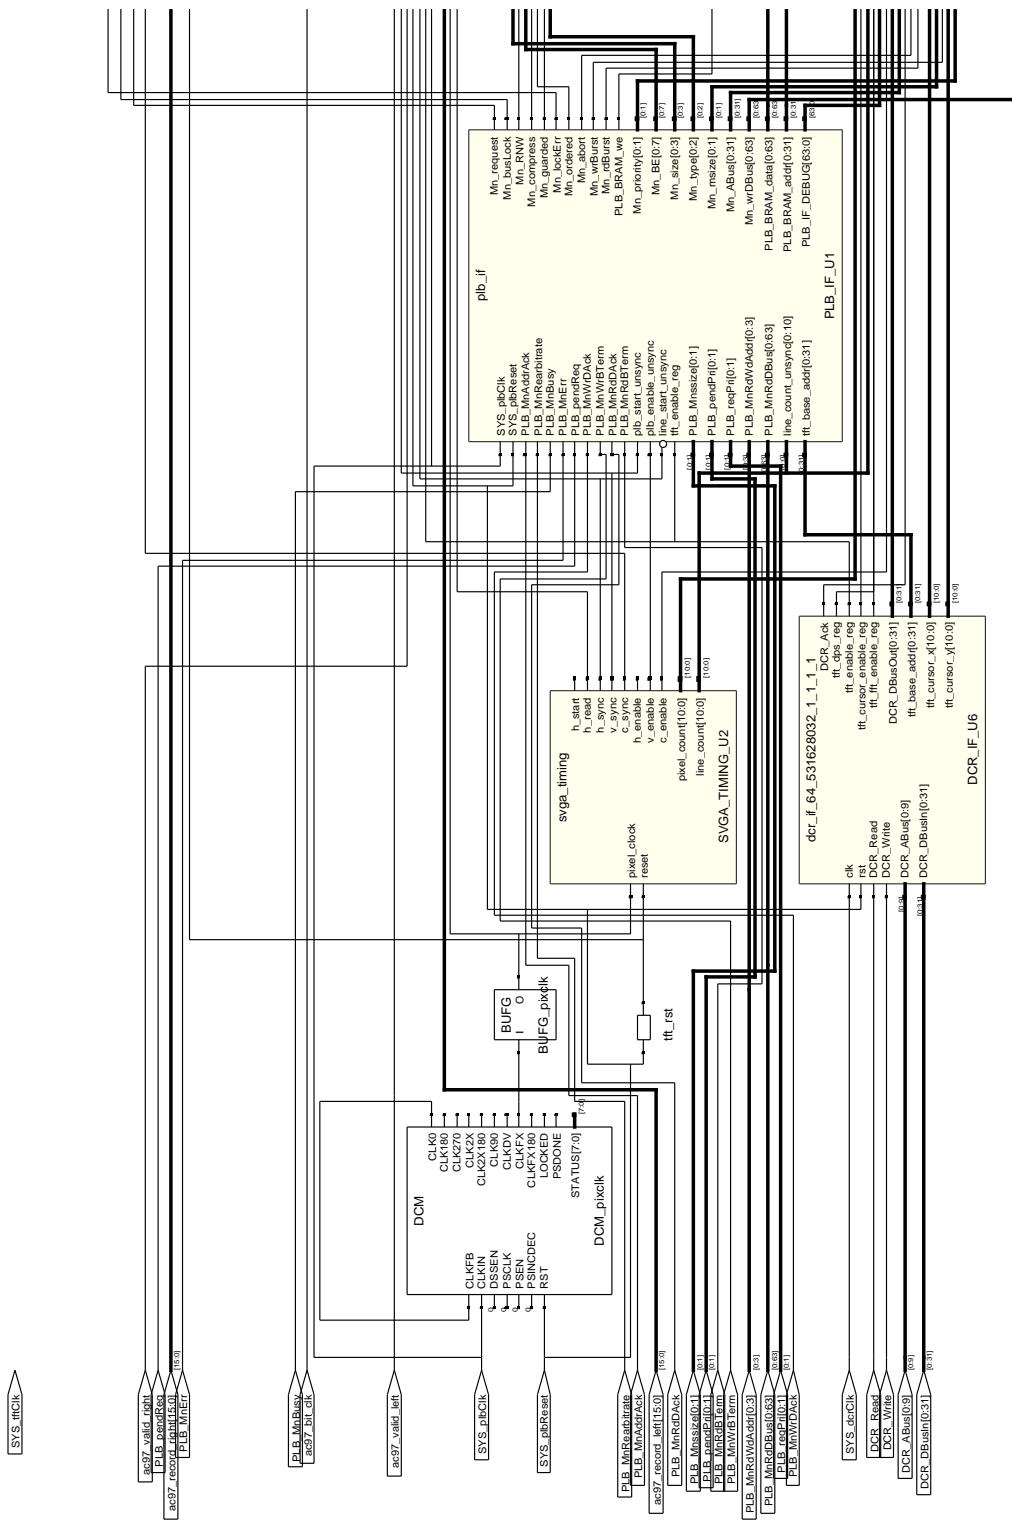


Figure 5.2: Framebuffer schematic (1 of 2).

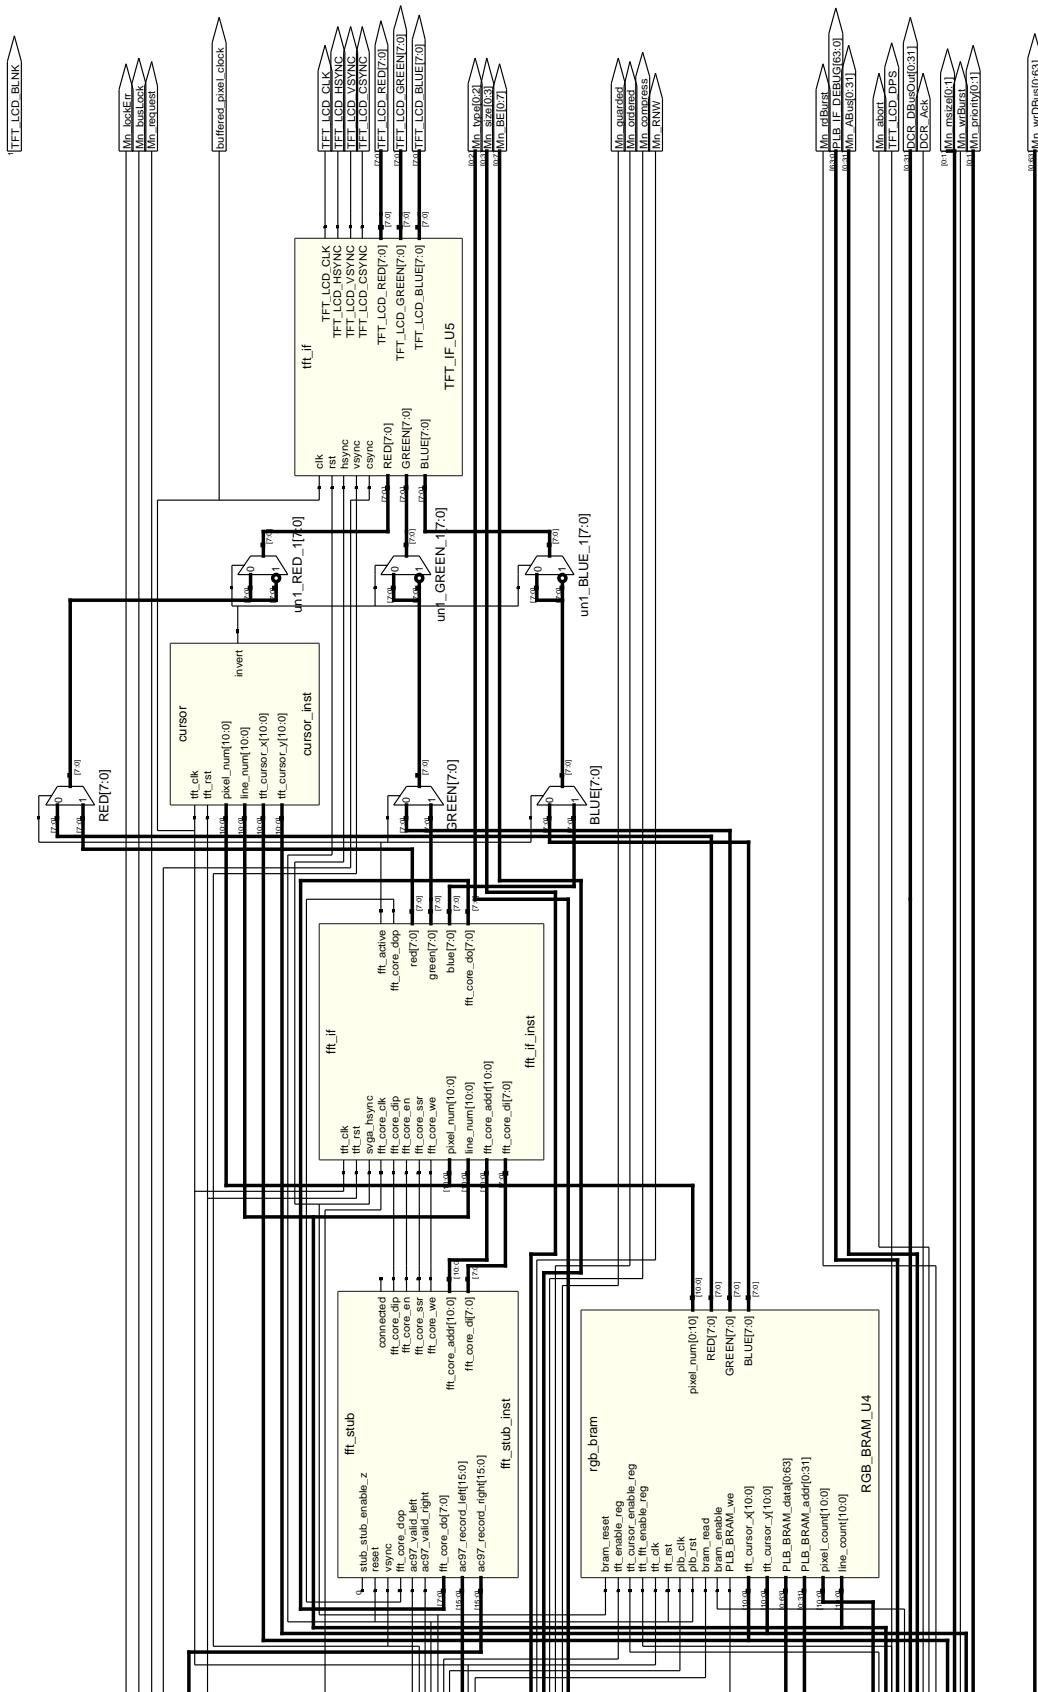


Figure 5.3: Framebuffer schematic (2 of 2).

The video bit clock for 1280×1024 at 60 Hz is 108 MHz, based on the timing data of Table 5.1. ($60 \text{ frames/s} \times 1066 \text{ lines/frame} \times 1688 \text{ pixels/line} = 107.96448 \text{ pixels/s.}$) This frequency is derived from the 100 MHz system clock with a Virtex2 Digital Clock Manager (DCM) configured for 27/25 multiplication.

Table 5.1: SXGA video timing parameters

Dimension	<i>Sync</i>	<i>Back Porch</i>	<i>Active</i>	<i>Front Porch</i>	<i>Total</i>
Horizontal (pixels)	112	248	1280	48	1688
Vertical (lines)	3	38	1024	1	1066

The Framebuffer acts as a PLB master to read bursts of pixel data from the DDR SDRAM memory. Each pixel is encoded in a 32-bit word, with 8 unused bits, followed by 8 bits each for red, green, and blue intensity. To better amortize the PLB transaction overheads, the read burst is 64 cycles long, with two 32-bit pixels per 64-bit PLB word, equating to 128 pixels read per burst. Ten read bursts thus fetch 1280 pixels, or one full line of SXGA output. The timing generation circuit also tells the PLB interface when to begin reading new line data after a horizontal sync, and when to reset its starting address after a vertical sync.

A number of settings for the Framebuffer can be configured through the PowerPC’s Device Control Register (DCR) bus. These settings include the base video address to use,¹⁰ a video display enable, a cursor display enable, and an FFT display enable.

AC’97 audio data is generated at 12.288 MHz, but must be displayed synchronously with the 108 MHz pixel clock. To sidestep cross-domain clocking issues, it seemed sensible to include a dual-ported BRAM, with one of the ports connected to the Framebuffer circuitry, and the other port terminating in a stub that a dynamic core could tie into. Although it would in some ways have been cleaner to write the FFT data to the appropriate location in DDR SDRAM, there was some concern about saturating the PLB bus’ capacity. The theoretical maximum transfer rate on the 64-bit PLB bus is 800 MB/s, and the Framebuffer reads 1280×1024 words, or 5 MB 60 times per second, for 300 MB/s. The video capture core pushes another 40 MB/s through the PLB bus, and the FFT display would have had similar requirements. Considering the various bus transaction overheads, and the fact that the PowerPC also accesses main memory through the same PLB bus, it seemed safer to generate the FFT pixel data from inside the Framebuffer, so that a much smaller amount of data could be transferred through a dual-ported BRAM.¹¹

A number of problems were encountered during the development of this core. Although the Framebuffer was partially inspired by the Xilinx plb_tft_cntlr_ref_v1_00_d reference core, the reference

¹⁰The base video address makes it possible to observe any part of the physical memory space. It’s quite fascinating to watch low memory as Linux loads, uncompresses, and boots itself.

¹¹One downside to generating the FFT display on the fly instead of reading it from DDR SDRAM, is that it doesn’t reside anywhere in memory. Consequently, screen captures taken from DDR SDRAM memory do not capture the FFT display, and must instead simulate it or leave it untouched.

design displays a lower resolution 640×480 at 60 Hz signal with a 27 MHz bit clock. When the resolution was increased to 1280×1024 at 60 Hz with its 108 MHz bit clock, the output signals from the FPGA to the video DAC exhibited excessive crosstalk and interference. The solution turned out to be configuring the IOBs driving the video DAC signals to *fast* slew rate, and a 12 mA drive strength.

Another problem was that the PLB specification allows bus masters to perform their transactions at one of four priority levels, but debugging with ChipScope Pro gave no indication that the priority level was being honored. In particular, when the video capture core is actively pushing data through the PLB bus and into DDR SDRAM, there are slight artifacts that show up in the SXGA output, and no combination of priority levels for the two cores seemed able to rectify the situation. This problem remains unresolved at present.

5.6 Sensor Stubs

Some of the stubs used in this system serve the purpose of detecting external conditions, and alerting the autonomous system to those conditions. As such, they play the part of the sensors required by the Level 4 autonomous system of Section 4.2.5. Alerts are provided to the system in the form of interrupt pairs—one to indicate activity startup, and another to indicate activity timeout—for each supported port: video, audio, mouse, and keyboard.

These sensor stubs engage in passive observation of their assigned ports, while passing data back and forth between the external devices and dynamic or static circuitry. Whenever data is received from an external device, activity is noted, and the timeout counter is reset. If no additional data is received before the timeout counter reaches its end, the system is notified that the peripheral is no longer available.

Each sensor stub performs its monitoring in the manner best suited to its peripheral, and that typically includes attempting to communicate with the peripheral. Details for each of the sensor stubs follows, but all of them share common infrastructure and conventions in how they notify the system of activity startup or timeout.

5.6.1 PS/2 Stub

The ps2 stub is a sensor stub that monitors a PS/2 port [62], and the system includes two such stubs to monitor both the keyboard and the mouse ports. Although PS/2 keyboards require little configuration, PS/2 mice have to be configured to send movement data. However, the keyboard and the mouse ports are electrically and functionally identical, and the distinction between them comes down to a matter of convention.

In the absence of activity, the stub periodically attempts to ping any device connected to its port, to determine when a device has been connected or disconnected. In the case of the keyboard, the stub sends an EE (Echo) command byte, and looks for an EE reply within a specific period of time. In the case of the mouse, the stub sends an E6 (Set 1:1 Scaling) byte, and looks for any acknowledgment. In both cases, if a device is connected to the port, the stub will flag activity and reset its timeout counter. This enables the ps2 stub to determine whenever a keyboard or mouse has been connected or disconnected. The ps2 stub probes for connected devices roughly ten times per second, and times out if no activity or response has been detected for one second.

The ps2 stub uses the OpenCores.org PS2 interface [60], to handle bus protocol and arbitration.

5.6.2 Video Stub

The video stub is a sensor stub that monitors the Analog Devices ADV7183B video decoder chip for video lock. Unlike the ps2 stub, the video stub does not look at the incoming video data, but instead requests interrupts from the chip whenever video lock is acquired or lost.

The VDEC1 board has component, composite, and S-Video inputs.¹² Instead of requiring that video be provided through only one of the three inputs, the stub sequentially tests each input when out of video lock. In each case, the stub sends configuration data to the ADV7183B to select the proper inputs and mode, and enables interrupts to signal acquisition or loss of video lock.

The video stub uses the OpenCores.org I²C controller [61], to handle bus protocol and arbitration. This is the only sensor stub in which all activity is detected through out-of-band signaling.

5.6.3 Audio Stub

The audio stub is a sensor stub that monitors the National Semiconductor LM4550 AC '97 codec chip for audio data. Unlike the video stub, audio data does not include any subcarrier to lock on, so the stub must look for input swings above some noise threshold to determine whether or not an audio signal is present.

Although the LM4550 supports multiple stereo inputs, the XUP board only provides a stereo line-level input, and a lower fidelity mono mic-level input. The stub periodically sends configuration data to the LM4550 if activity has timed out. This would normally search sequentially between the line and mic inputs, but two problems made this undesirable in practice. On one hand, the XUP board that was used here has a significant DC offset on its left input, and adjusting for that in a reasonably portable manner, significantly reduced the sensitivity of the audio stub. On the other

¹²Component video requires three inputs, typically labeled YCrCb, where Y is the luma, and Cr and Cb are the red-yellow and blue-yellow chroma, respectively. Composite video requires a single input that carries luma and chroma together. And S-video requires a dual input, with separate luma and chroma.

hand, audio data often includes periods of silence or near-silence, if only between audio tracks, and it was disconcerting to hear the stub scan the next input between each song. That problem could have been overcome by making the activity timeout period much longer, but a real system should be free to choose the solution best suited to its application, and this demonstration opted to only look for stereo line data.

The audio stub uses the Xilinx opb_ac97_v2_00.a to handle the AC97 protocol. This is the only sensor stub that is unable to conclusively establish the presence or absence of its intended input, because of the way that valid data can be obscured by the noise floor.

5.6.4 Unused Stubs

A number of additional *non*-sensor stubs were developed, but remain unused at present. These include PLB master and OPB slave stubs. Although these stubs were set aside for the time being, it might be desirable to recall them in the future, since they provide interfaces to flexible and well accepted buses.

5.7 Interface Stubs

The original stub developed for the demonstration system was an OPB IPIF peripheral. As the description suggests, the SystemStub was an interface between system OPB bus—conforming to the IPIF specification—and dynamically instantiated logic. Before the goals of this work were expanded, this stub was intended to be the only interface between dynamic circuits and the system.

The SystemStub provides a 32-bit address, data in, and data out busses, along with bus control signals and an interrupt input. Like most of the system’s peripherals, this core inhabits a 64 kB block within the processor’s address space. A portion of that provides access to IPIF memory-mapped registers, and the remainder of it is available for the connected dynamic circuitry to decode as it pleases.

The development of this particular stub initially followed a rather defensive track towards circuits that might connect to it, for fear that improper behavior would lock up the OPB bus and the remainder of the system—reminiscent of some of the security issues mentioned in Section 4.3.2. In practice, that fear turned out to be unfounded, and subsequent cores were developed with a more balanced view of dynamic circuits. Nevertheless, the software driver controlling this stub could disable it, enable it without interrupt support, or enable it with interrupt support.

Subsequent stubs for the audio, video, and PS/2 data were developed and incorporated into the system in place of more generic OPB and PLB stubs, for the sake of performance.

5.8 Custom Dynamic Cores

The dynamic cores developed for this work are intended to demonstrate the lower-level system capabilities—including the routing, placement, and reconfiguration that qualify this as a Level 2 system—and to provide functionality suitable for high-level autonomy—showcasing the ability to appropriately respond to detected conditions of a Level 6 system.

These cores must be instantiated by the system while it continues to run, and must be connected to the appropriate inputs, outputs, or stub ports. The two primary demonstration cores are a video capture core, and a Fast Fourier Transform (FFT) core. Each was first developed and implemented statically, before being permanently removed from the base system, to be implemented dynamically. The relevance of these two cores lies in their ability to process video and audio signals for real-time display on an SXGA monitor. A number of other dynamic cores were developed, but remain mostly unused at present.

5.8.1 Video Capture Core

The underlying video capture code was derived from the Xilinx `video_capture_rev_1.1` reference design provided for the XUP board. The original code extracts timing and sync information, performs chroma interpolation and colorspace conversion from YCrCb to RGB, and pushes the data into a line buffer for subsequent use by a VGA display driver. The static build of the core does away with the VGA interface, and instead provides an interface that bursts data over the PLB bus and into the Framebuffer’s memory space. The dynamic build connects to the system’s video core stub, through one port of a dual-ported BRAM. The use of the dual-ported BRAM reduces the difficulty of pushing 27 MHz video data onto a 100 MHz bus.

The original code derives a non-constant 13.5 MHz clock from the *line-locked* 27 MHz clock, and uses the derived clock to drive other logic. Whether intentional or not, that *gated clock* technique is undesirable for a number of reasons, including the fact that the resulting 13.5 MHz clock was badly skewed with respect to the 27 MHz clock.¹³ The final core instead pushes the 27 MHz clock through a DCM, while also capturing its 13.5 MHz divide-by-two output, using the latter more appropriately as a chip-enable. With this approach, both clocks are deskewed from each other, and the extracted pixel clock is still available for any logic that depends upon it.

¹³In practice, the skew depended very strongly on placement and routing, and could thus vary widely from one build to the next, sometimes causing highly unusual visual artifacts.

5.8.2 FFT Core

The FFT core is based on the OpenCores.org cfft [63] that performs Fast Fourier Transforms upon complex-valued data. The incoming audio data is purely real—with no imaginary component—so a complex transform is not strictly necessary, but the cfft core is freely available and fairly easy to use.

The cfft core is configurable in both transform size and input width. This implementation configures it as a 256-point 12-bit FFT with the imaginary input tied to zero, and pays little attention to its inner workings. Some extra logic is inserted in-line with the FFT input, to multiplex it between left and right channel audio data. In the general case, a complex FFT produces differing positive and negative frequency spectra, but for the real-valued audio data, the positive and negative frequency spectra are identical, and this FFT therefore effectively produces 128 significant points worth of frequency band data.

The left and right audio samples become available simultaneously, every 256 cycles of the 12.288 MHz AC '97 bit clock, which corresponds to a 48 kHz sampling frequency.¹⁴ To use a single FFT core for both audio channels, the data from one of the channels is delayed by five cycles of the 48 kHz clock, until the FFT is ready to accept its next input. Because the human ear senses sound logarithmically rather than linearly, the resulting frequency data is pushed through circuitry that approximates a logarithm function, before being pushed to the core's output ports.

5.8.3 Other cores

A number of other cores were explored during development, but most were never fully implemented. These include the OpenCores.org 3des¹⁵ core and the OpenFire soft-processor. A few of the cores that will be mentioned again in subsequent chapters are a 4-bit decimal counter, a 32-bit binary counter, a block of IOBs configured as inputs, and a block of IOBs configured as outputs.

5.9 Floorplanning

Floorplanning consists of constraining various parts of a design to specific physical regions in a device. This can be done for a variety of reasons, but is most commonly done to guarantee access to critical resources, to improve timing performance, or to reserve certain regions for future use.

Static designs often use constraints for performance reasons. The ISE place and route tools tend to work as hard as they are asked to, but no more. This means that by increasing the applied

¹⁴The 256-cycle frequency is a result of the AC '97 specification, and has nothing to do with the fact that the FFT is 256-points wide. This is simply a convenient coincidence.

¹⁵3des is a cryptographic core that performs Triple-DES encryption or decryption.

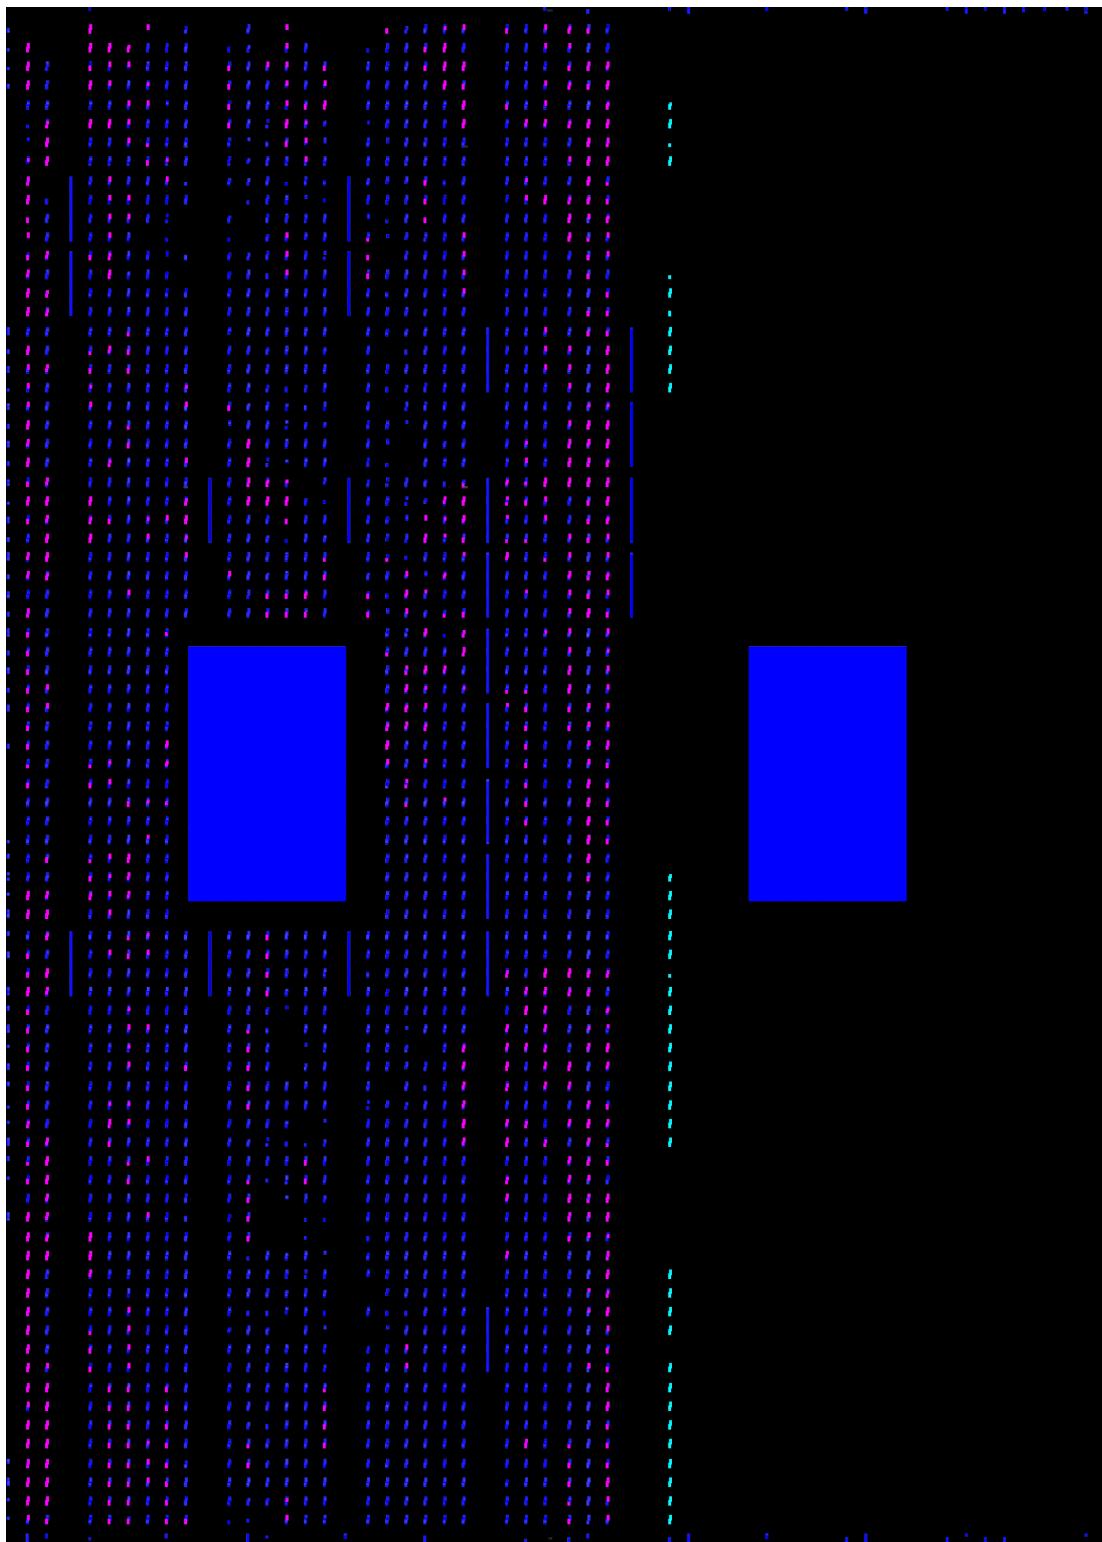


Figure 5.4: Demonstration system floorplan: static system resources are shown in blue and pink, while interface stubs are shown in teal.

constraints, it is possible to obtain smaller areas or better performance. However it is also easy to specify constraints that the tools are unable to meet.

Dynamic design may also need to manage available real estate since they will certainly be changing. In view of such changes, it is generally desirable to group unused or reserved resources into one or more clusters. Although this does not change the number of logic cells available, it does tend to reduce total routing length as well as routing congestion.

But dynamic designs may also require constraints that arise from underlying architectural issues. Section 2.3.3 discussed cases in which state corruption may occur during dynamic reconfiguration of a Virtex2P, and concluded that care must be taken with LUTs used in RAM or shift-register mode, and also with dual-ported BRAMs. The configuration architecture of the Virtex2P uses vertical configuration *frames*, spanning the entire height of the device, and that means that in order to protect corruptible resources, it is necessary to use horizontal separation. For the sake of simplicity, this work stipulates that any part of the base system circuitry may potentially contain corruptible resources, and therefore the static base system is entirely constrained to the left 55 % of the device, leaving the right 45 % of the device available for use by dynamic circuits. Figure 5.4 shows the system floorplan, with static resources in blue and pink, and interface stubs in teal.

5.10 Summary

The system hardware consists of the base design, that remains static over the life of the system, and of dynamic designs, that can be freely instantiated or removed from the running system. Interfacing between static and dynamic parts is facilitated by interface stubs, while detection of environment changes is facilitated by sensor stubs. And finally, the entire design is floorplaned to avoid state corruption that might occur during dynamic reconfiguration.

Chapter 6

Software

6.1 Overview

The demonstration system requires a considerable amount of support software, ranging from an operating system to special drivers, and including a Java Runtime Environment (JRE).

Figure 6.1 depicts the major software components that compose the system. This chapter discusses the Linux kernel, custom drivers, and Java Runtime Environment, all colored in green. Most of these components relate only indirectly to autonomy, with the exception of the driver that talks directly to the ICAP hardware. The autonomous server and controller, colored in mauve, are discussed in Chapter 7, as major system components that are directly involved in autonomy.

This chapter thus presents a necessary and important part of the overall system, though most of it builds on work done by others—the Linux kernel and the JRE are cases in point—without adding any groundbreaking novelty. However the custom drivers are original work, and include some special capabilities that directly support the autonomous server and controller, and the system’s custom hardware.

The discussion in this chapter begins with the Linux kernel build and the filesystem configuration, and then turns to each of the custom drivers, concluding with some brief points pertaining to Java. If the reader has sufficient understanding of Linux and Java, and is inclined to skim through this chapter, Section 6.3 on custom drivers is the part that most merits a little extra attention.

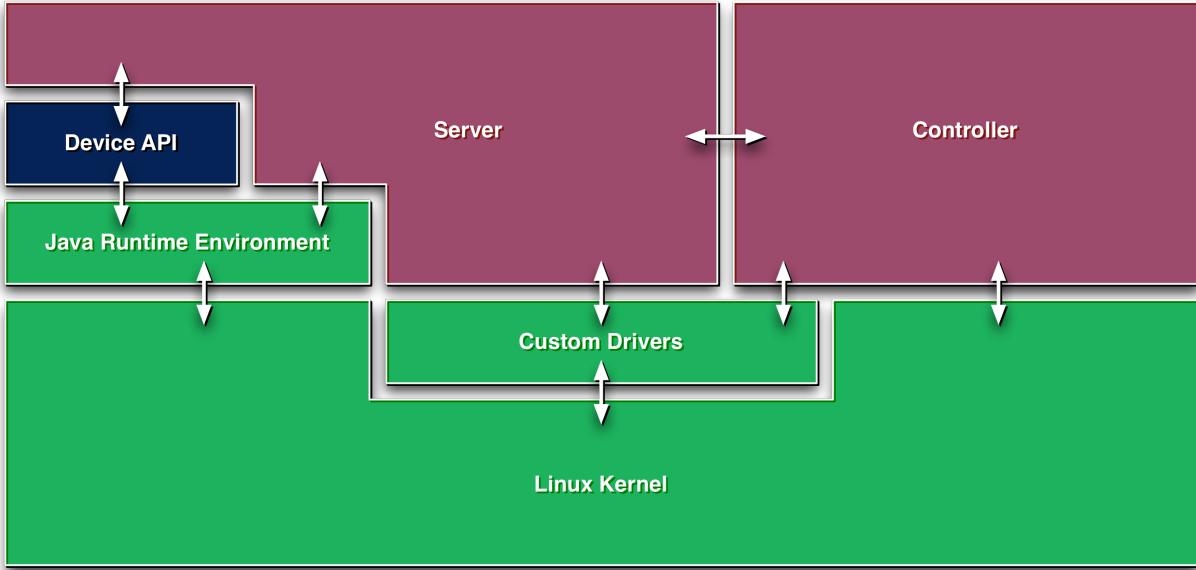


Figure 6.1: Software block diagram.

6.2 Linux

Linux as an operating system [26], and the various efforts to port it to the XUP board [29, 30, 31, 32], were previously introduced in Section 2.7. This section dives more deeply into the build and staging process involved.

6.2.1 crosstool

Most of the development took place on an Intel x86 Linux platform, but the target processor was an embedded PowerPC 405, so the kernel and all of the necessary software tools had to be cross-compiled for the target. Fortunately cross-compilation is fairly common in the Linux world, and utilities are available to facilitate the process.

The *crosstool* utility [35] is designed to build cross-compiling *gcc-glibc* [64] tool chains. This utility was used to build a PowerPC 405 Linux GNU tool chain for *gcc-3.4.4* and *glibc-2.3.3*.¹ The resulting tools—with the *powerpc-405-linux-gnu-* prefix omitted—are:

`addr2line, ar, as, c++, c++filt, cpp, g++, gcc, gcc-3.4.4, gccbug, gcov, gprof, ld, nm, objcopy, objdump, ranlib, readelf, size, strings, strip`

¹The PowerPC 405 does not support floating-point instructions, so the build was carried out with `GLIBC_EXTRA_CONFIG="--without-fp"` to prevent the generation of any floating-point code.

6.2.2 mkrootfs

Klingauf provides the mkrootfs script to automate part of the root filesystem generation [33]. The script generates many of the common filesystem directories—/bin, /boot, /dev, /etc, /opt, /home, /lib, /mnt, /proc, /root, /sbin, /tmp, /usr/local/, and /var—and populates them with suitable defaults.

Klingauf’s script was extended to install additional software, and to further configure /root and sshd, the secure shell daemon.² The additional software includes Java, OpenSSL [65], OpenSSH [66], and zlib [67].

6.2.3 Kernel 2.4.26

The UW EMPART project [29] was based on the 2.4.26 kernel because of the availability of Xilinx drivers for the various board components.³ In addition to the kernel configurations specified for EMPART, a few more changes were made for the present work:

- The kernel was configured to emulate floating-point operations in software, for the sake of any applications that might require that support. This may seem to run counter to the intent of the `GLIBC_EXTRA_CONFIG="--without-fp"` option for `gcc`, but in effect this tells the compiler not to generate floating-point code, while still providing kernel support for compiled software that already uses floating-point code, particularly the JRE.
- The kernel boot command line was modified to reserve the top 6 MB of memory for the video framebuffer. This constrains the stack and all kernel and user memory to the first 506 MB of memory
- A few other minor adjustments were made, relating to filesystem and driver support. One change in particular enabled FAT support, to facilitate rewriting the boot partition of the filesystem, and thus the base configuration bitstream that it contains.

It is worth noting that the 2.4 kernel uses `devfs` for driver registration, while the 2.6 kernel instead uses `udev`, as discussed in [68]. Although that change has only a minor impact on the drivers developed here, it does mean that those drivers will not directly work with the 2.6 kernel.

A few kernel patches are required for the build, and are detailed by BYU. The EMPART project provides these patches in the form of a *gzipped tarball*, along with a script that applies them to the kernel source.

²Secure shell access to the system through ssh is the best way to remotely interact with it. The regular Linux installation on this system only provides console access through an RS-232 port, but ssh access through ethernet is much faster, supports an arbitrary number of connections, and can be used from anywhere on a reachable network.

³Many of the drivers originate with MontaVista Linux [28] and were first developed for the Xilinx ML300 and ML310 boards.

6.2.4 Packaging the Kernel

In order for the system to boot Linux, it needs to obtain the kernel image from somewhere. This implementation simply places the compressed image on an 8 MB bootable partition on the CompactFlash MicroDrive, to be read by the SystemACE controller. The kernel image is actually packaged with the base configuration bitstream into the expected .ace file.

Klingauf and BYU also provide information on setting up a *bootp* server [33, 30], allowing the system to boot over a network connection, but given that the system was designed to demonstrate autonomy, it seemed more appropriate for it to boot on its own, without depending upon a *bootp* server.

6.2.5 BusyBox

For Linux shell utilities, Klingauf recommends BusyBox [34]: “BusyBox is a superb all-in-one utility, which gathers all important and/or useful utilities in one single binary. It comes with a kernel-like configuration interface and is [intended] to be cross compiled.”⁴ BusyBox can be configured in much the same way as the kernel, making it possible for the user to select the specific utilities that they want to include. The current configuration supports the following utilities:

```
[, [], addgroup, adduser, ash, basename, bunzip2, busybox, bzcat, cat, chgrp, chmod, chown, chroot, chvt, clear, cmp, cp, cut, date, dd, deallocvt, delgroup, deluser, df, dirname, dmesg, du, e2fsck, echo, egrep, env, expr, false, fdisk, fgrep, find, free, fsck.ext2, fsck.ext3, getty, grep, gunzip, gzip, halt, head, hexdump, hostname, id, ifconfig, inetd, init, insmod, install, ip, kill, killall, klogd, less, linuxrc, ln, logger, login, ls, lsmod, mkdir, mknod, mktemp, modprobe, more, mount, mv, netstat, nohup, nslookup, openvt, passwd, pidof, ping, pivot_root, poweroff, ps, pwd, rdate, readlink, reboot, reset, rm, rmdir, rmmmod, route, sed, sh, sleep, sort, strings, su, slogin, swapoff, swapon, sync, syslogd, tail, tar, tee, telnet, telnetd, test, time, top, touch, tr, true, tty, umount, uname, uniq, unzip, uptime, usleep, vi, wc, wget, which, whoami, xargs, yes, zcat
```

6.2.6 Other Utilities

In addition to the functionality provided by BusyBox, the system includes a few other necessary or simply convenient packages. Each of these is cross-compiled with the tools generated by crosstool:

OpenSSL 0.9.8d: An open source Secure Sockets Layer (SSL) and Transport Layer Security (TLS) implementation. OpenSSL is a fascinating and complex package that handles a wide range of cryptographic needs, but this system employs it primarily in support of OpenSSH.

⁴The reader is advised to use BusyBox 1.1 or later rather than BusyBox 1.0, which crashes as soon as the kernel invokes init.

OpenSSH 4.3p2: An open source Secure Shell (SSH) implementation. OpenSSH uses OpenSSL, and supports ssh and scp for secure connections and secure file transfer with the XUP board.

zlib 1.2.3: An open source compression library, referenced by OpenSSL and OpenSSH.

A number of other packages, including ntp, gcc, make, and kaffe, were compiled and built but not ultimately used in the system.

6.2.7 Filesystem

To simplify testing and debugging, the NFS volumes used for development of all the project software were mounted on the target system.⁵ The system also maintains a local copy of the project software, though the local copy is more prone to fall out of date with the data on the server. When the NFS volumes are mounted, they simply overlay the local data, and when they are not available, the local data is still visible and usable in the same place.

6.3 Custom Drivers

Custom Linux drivers were developed to communicate with the system's custom hardware. The custom hardware is mapped into the system's physical memory space, but on protected-memory systems like Linux, regular user code cannot access physical memory. It instead sees only the virtual memory space that the kernel and Memory-Management Unit (MMU) expose to it. Access to physical memory is only possible from kernel space, and drivers that live in kernel space are therefore used to access that memory. Some of the drivers also receive or generate interrupts, and that provides another reason to work from kernel space. Finally, driver functionality can be mapped into the filesystem, through simple and well-understood methods, making the functionality available to any code that can perform file I/O, from compiled code and interpreted Java code to shell scripts. Much of the infrastructure for these drivers is based on [69].

The following custom drivers are currently used by the system. Each of these will be discussed in greater detail in the following subsections.

ac97: AC '97 audio core driver.

fb: Framebuffer driver.

fpga: FPGA resource usage visualizer.

gpio: Interface to the XUP board LEDs, pushbuttons, and DIP switches.

icap: Interface to the FPGA's Internal Configuration Access Port.

ps2: PS/2 mouse and keyboard driver.

⁵The kernel as built was unable to share locks on exported volumes, so the volumes were mounted with *-o nolock*

stub: Interface for all sensor stubs and subscribers.

video: Video capture core driver.

6.3.1 AC '97 Driver: /dev/ac97

The ac97 driver interfaces with the low-level xac97 driver provided by Xilinx. It presently serves only to configure the LM4550's registers and audio paths, and optionally to loop sampled data from the inputs back to the outputs. Any settings applied to the LM4550 must be considered volatile, because the audio sensor stub overwrites those settings whenever it loses an audio signal and tries to reconfigure the LM4550. At present this driver supports neither read() nor write() operations, because the audio data is primarily intended for processing in hardware by the FFT core.

6.3.2 Framebuffer Driver: /dev/fb

The fb driver interfaces with the Framebuffer hardware core, and allows user space applications to control the framebuffer and draw on the SXGA display. The control functions are accessible through ioctl() calls:

- Set the framebuffer base address. This sets the physical memory address corresponding to the top left of the SXGA display. The caller can select one of two framebuffers, or can choose to display arbitrary memory contents.
- Enable or disable the display. When disabled, the framebuffer still generates video sync signals, but outputs all black pixel data.
- Erase the screen to an arbitrary color.
- Control the cursor position. This is intended to provide visual feedback for mouse tracking.

Drawing to the display happens by mapping the framebuffer's physical memory into the user's memory space with mmap() calls. All of the drivers that access physical memory must use ioremap() to remap the physical memory into the kernel space virtual memory, but the framebuffer must go a step further by mapping that memory into the user application's memory space. The mmap() function returns a pointer that allows the user to arbitrarily read from or write to framebuffer contents.

Implementing the driver-side portion of mmap() is not too difficult in general, but forcing immediate reads and writes with no caching was a little more complicated. The default behavior seems to be for the kernel to pool any changes to mapped memory, and to actually apply them to the underlying physical memory at its convenience. This behavior is much like that of a processor cache, where coherency is what matters most, and where flushing cached data to main memory is a lower priority. In this case, however, what is written to the physical memory directly determines what the screen displays, and any delay in flushing changes to the framebuffer memory result in visual artifacts. In

some cases, blocks or residual pixels took at long as ten or more seconds to update, resulting in unsightly behavior, irrespective of flush requests.

To force the kernel to truly remap virtual memory pages, and to not allow any caching, this driver borrowed code from pgprot_nocached() in drivers/char/mem.c, and from <http://www.scs.ch/~frey/linux/memorymap.html>. The resulting screen updates happen just as fast whether written from user space or kernel space, thus guaranteeing that the kernel does not perform any buffering in the background.

A few auxiliary utilities were developed to control or utilize the framebuffer: Command-line utilities can turn the display on or off, or perform a screendump by writing the current image to a .bmp bitmap file.

6.3.3 FPGA Driver: /dev/fpga

The name of the fpga driver may cause some measure of confusion. All reconfiguration of the FPGA happens through the /dev/icap driver, while the /dev/fpga driver instead displays current resource usage on screen, or provides the usage information to callers.

The fpga driver includes size and position information for every logic cell in the XC2VP30 device. Each cell is numbered following the same conventions as the Device API, and clients can set or query usage information with read() and write() calls and a simple protocol. The driver internally maps the framebuffer memory into kernel space, for the sake of drawing usage directly to the framebuffer.

There are two main applications in the demonstration system that update resource usage information, and converse with each other through the /dev/fpga driver: the server and the GUI. Both of these will be discussed at length in Chapter 7, but they are relevant here through their interaction with this driver. The typical scenario is as follows:

1. Upon startup, the server identifies the base resource usage and writes it to /dev/fpga.
2. The GUI reads the base usage and remembers it.
3. The user is free to interactively mask or unmask resources with the mouse.
4. Before any circuit placement, the server re-reads the resource usage, to synchronize its resource availability information with any user masks.
5. After every circuit placement or removal, the server updates the resource usage information.

All of this handshaking takes place behind the scenes. The user only sees the on screen updates anytime circuitry changes or anytime the user masks or unmasks resources with the mouse.

6.3.4 GPIO Driver: /dev/gpio/*

The gpio driver exposes the XUP board's LEDs, DIP switches, and push-buttons as filesystem devices: /dev/gpio/led4, /dev/gpio/dip4, and /dev/gpio/push5. General Purpose Input/Output (GPIO) has kernel level support through the xgpio driver, but for debugging purposes it was convenient and desirable to be able to read and write these resources through the filesystem.

6.3.5 SystemStub Driver: /dev/hwstub

The hwstub driver was originally designed to be the one driver necessary to communicate with dynamic circuits. But that role changed significantly as the thrust of this work expanded, and as the hwstub services were supplanted to support much more specialized and complex dynamic circuits. Nevertheless, the underlying SystemStub core does exist in the base system at present, and the hwstub driver still provides a way of controlling it. Data can be read from or written to connected dynamic circuits with read() and write() calls. The hwstub driver also allows the user to reset connected circuits, to enable write-only, read-write, or read-write with interrupts.

6.3.6 ICAP Driver: /dev/icap

The icap driver exposes the low-level Xilinx hwicap_v1_00_a driver as a filesystem device: /dev/icap. The low-level driver provides functions for communicating with the FPGA's configuration controller. This includes the ability to write full or partial bitstream to the device, and to read back device configuration or state information. The driver source has been made publicly available under the GPL license [70], and has been put to use by a small number of researchers [71].

The driver supports partial dynamic reconfiguration through file write() operations, to facilitate reconfiguration from anything from shell scripts to Java code. Internal bitstream headers are first stripped from the bitstream file, and the file contents can be written in arbitrary size and alignment. The driver does not inhibit full bitstreams or non-active partial bitstreams, but the first thing that such bitstreams do is shut down the device, so any attempt to write one of them is ill advised. Readback is not supported at present, primarily because it was not required by the demonstration system.

6.3.6.1 Errors in Low-Level hwicap_v1_00_a Driver

Development and debugging of the icap driver was hampered by an error in the low-level hwicap_v1_00_a driver that inserts an extra word in the bitstream. The effect of the extra word depends on where in the bitstream it occurs, and because the hwicap_v1_00_a driver does not split data along frame boundaries, it may just as likely affect frame data as configuration commands.

Unlike the CRC calculation error of Section 6.3.6.2, that did not actually corrupt any frame data, this error typically does corrupt frame data, and therefore interferes with the proper operation of the device.

To sidestep this error, the Linux ICAP driver does not directly use the offending `hwicap_v1_00_a` function, but instead uses a corrected version of that function. Two commonly held beliefs are that configuration data must always be written along frame boundaries, and that failing to provide frame data in uninterrupted fashion will stall the configuration controller. With the correction specified in the `comp.arch.fpga` thread [72], the present work has consistently been able to write data with arbitrary alignment and arbitrary interruptions, without so much as a glitch in device operation.

6.3.6.2 CRC Calculation Errors

To protect FPGAs against errors stemming from accidental bitstream corruption, it is common to verify the bitstreams with some sort of Cyclic Redundancy Check (CRC). The Virtex2 architecture allows the user to enable or disable the use of CRC checks. If CRC checks are disabled, the configuration controller expects to see a predefined constant where the CRC word would normally occur. If CRC checks are enabled, the configuration controller will calculate its own CRCs on the frame data that it receives, and compare the results with the CRC words included in the bitstreams.

The Virtex2 architecture uses the CRC16 polynomial $X^{16} + X^{15} + X^2 + 1$ in its bitstreams [19, *polynomial may not appear in revisions later than 1.3*]. Unfortunately the CRC calculation code in the Device API was ported from C to Java, without fully appreciating the subtle differences in sign-extension between those two languages.⁶ As a result of that difference, the CRC words generated by the Device API were occasionally incorrect, and therefore configurations that did not disable CRC checks would occasionally set the CRC Error flag, even though the configuration took place correctly.

CRC errors in an active configuration bitstream do not have quite the same effect as errors in a shutdown bitstream however. Shutdown bitstreams disable the entire device while the frame data is modified, and if CRC errors are detected, the configuration controller will prevent the device from starting up again. In active bitstreams, frame data may be written to the configuration frames before the CRC words are even encountered, meaning that erroneous data can be written to the frames, and the only indication of the error will be the `CRC_ERROR` flag in the configuration controller Status Register. Matt Shelburne's willingness to make his bitstream and CRC calculation code available for comparison significantly shortened the debugging effort [73]. When the CRC calculation error was corrected in the Device API, the `CRC_ERROR` flag issue disappeared.

⁶Java sign-extends implicit conversions from int to long, while C does not, even for identically-sized signed types.

6.3.7 PS/2 Driver: /dev/ps2/*

The ps2 driver provides mouse and keyboard support when these devices are connected to the XUP board. The underlying ps2 hardware core generates interrupts when data is received from either device, and the driver registers itself to process those interrupts. In the case of the keyboard, raw scan codes are first mapped to standard ASCII codes. And in the case of the mouse, the movement data is passed along unchanged. Commands can be written to the devices with write() calls, and data or responses can be read back with read() calls. Note that the ps2 sensor stub detection of activity is entirely transparent to this driver: it simply has no idea that the stub exists.

To communicate efficiently with clients, the ps2 driver supports the poll() call. Although the use of “polling” may sound horribly misguided to the reader, *polling* as the Linux kernel defines it is something quite different. In fact, it is simply the standard way for kernel drivers to support non-blocking I/O, in a manner that can notify listening processes as needed. While conventional polling is an inefficient way to perform communication or I/O, this kind of polling allows the driver and its clients to waste as little time as necessary when no input is available. The *non-blocking* aspect of this approach is also worth underscoring, because it allows for very responsive clients, that can listen to many inputs simultaneously. Despite the naming, the driver and clients actually interact as if interrupts were being sent from driver to clients.

6.3.8 Stub Driver: /dev/stub

The stub driver receives interrupts from sensor stubs, whenever activity begins or times out. Like the ps2 driver, the stub driver supports the poll() call, and simply generates codes to notify the client of the current activity status for the mouse, keyboard, video input, and audio input. This provides an easy way for the client software to stay abreast of what could be multiple changing input conditions.

6.3.9 Video Driver: /dev/video

The video driver used to play a fairly prominent role in communicating with the ADV7183B video decoder chip, before all of the configuration duties were transferred to a state machine in the video stub hardware. The driver used to communicate with the Xilinx i2c driver, to control the inputs, paths, and settings through an I²C protocol. At present the driver survives mainly as a software switch to enable or disable video capture.

6.4 Java

The Java language and environment are exhaustively documented by Sun Microsystems, Inc., and others [74]. This section does not dwell much on Java itself, but rather provides an overview of the terminology and environments, and the effort to use Java on the demonstration system.

In broad strokes, Java is an interpreted and highly portable language. Java source code is compiled into Java *bytecode*, and Java bytecode can run inside any conforming Java Virtual Machine (JVM). In some cases, parts of the interpreted bytecode are compiled on-the-fly into machine code suitable for the platform. This kind of compilation that occurs while the program is already running is called Just-In Time (JIT) compilation, and a *JIT* can denote a Just-In Time compiler.

Java also defines and provides a large number of standard APIs, that are expected to be available to Java programs at runtime. The combination of the JVM and the APIs largely constitute the Java Runtime Environment (JRE).

Terminology notwithstanding, the demonstration system requires a JRE, in order to execute key parts of its software. In practice, Sun’s “write-once, run anywhere” slogan is well applicable to desktops, workstations, and supercomputers, but less so to embedded computers, which come in a wide variety of processors and resource configurations. The point is simply that finding a suitable JRE for an embedded system can sometimes range from difficult to impossible.

6.4.1 Unusable Java Runtime Environments

Finding a JRE that was suitable for the target system proved to be surprisingly difficult. The first attempt used Sun J2SE 1.5 source code, but its dependence on a large number of desktop-grade packages severely complicated the build process. It then became apparent that the required HotSpot virtual machine could only target x86 or Sparc platforms, which disqualified it as an option. Similar attempts with Sun J2SE 1.4.2 yielded the same results.

Because the PowerPC 405 is an IBM processor, IBM Java for PPC 1.4.2 sounded like a possible alternative. Unfortunately the software invariably crashed shortly after startup if the JIT was enabled.⁷ The purpose of the JIT is to dynamically compile frequently used blocks of Java bytecode into native machine code, and the penalty for not using a JIT is significant. The startup time to launch the server software and initialize all of its databases with the IBM JRE and no JIT was 767 s—almost 13 minutes. The corresponding times with the actual JRE used by this project were 557 s with no JIT, compared to 175 s with a JIT, a more than 3 \times performance difference.

⁷Ben Hardill at IBM was very helpful in investigating the problem, but the final conclusion was that the 18 instruction sets supported by the JIT—most of them for desktop-grade processors—all included instructions that were incompatible with the embedded PowerPC 405.

The next option was to try Kaffe, a freeware JRE that can be compiled for just about any 32-bit processor supported by the GNU C compiler: x86/IA32, ARM, MIPS, PowerPC, Sparc, 68K, IA64, Alpha, PA RISC, and more [75, 76].

During the lengthy Kaffe build, the GNU Compiler for Java (GCJ) was also investigated [77]. GCJ compiles Java class files directly into native machine code, that can then be executed without requiring a JVM. Tests on the x86 *development* platform ran a representative application in 16.42 s with the regular Java JRE, versus 6.49 s with the GCJ executable. Testing GCJ on the target platform required rebuilding the toolchain with Java support and with a newer version of GCC. Unfortunately the corresponding GNU Classpath project [78] that provides the equivalent of the Java API, did not support all of the Java2 1.4.2 classes, so the resulting executables could not run on the target platform. Attempts to build a newer version of GNU Classpath failed because crosstool-0.38 was unable to support the necessary gcc-glibc versions for PowerPC 405 targets, and crosstool-0.43 unfortunately requires a 2.6 Linux kernel. Of all the solutions investigated, GCJ is the one that most clearly deserves further exploration.

6.4.2 Java SE for Embedded (PowerPC)

Interestingly enough, this project went full-circle in ultimately discovering a suitable Sun JRE: the *Java SE for Embedded (PowerPC)* [79]. Unlike most other parts of Java, Java SE for Embedded is not free but must instead be licensed,⁸ although use is permitted on a trial basis for a limited time. Sun was kind enough to allow this work to use Java SE for Embedded beyond the regular trial period, contingent upon non-commercial use.

Web comments suggested that invoking Java in -server mode would result in much better performance than invoking it in regular client mode. Unfortunately, passing the -server flag made no difference to Java SE for Embedded, and the the version information consistently stated that is was running “Java HotSpot(TM) Client VM”⁹ regardless of the -server flag. Apparently the HotSpot Server VM does not exist for all Java incarnations.

The most important thing was that Java was in fact able to run on the demonstration system, and that the significant amounts of Java code that this system uses were in fact able to run without modification. However, compatibility does not always guarantee identical runtime behavior, because of the way certain data structures are implemented. In particular, any code that uses hashes in Java should not make assumptions about the order in which hash entries exist—but the same is also true in other languages. In other words, Java SE for Embedded was fully correct, but the behavior between the code running on a development machine and on the target system was in fact different, and that did further complicate development, staging, and debugging.

⁸That was true as of time when this JRE was discovered.

⁹HotSpot is the name of Sun’s JIT compiler for java.

Chapter 7

System

7.1 Overview

Previous chapters have discussed the hardware and software underlying the system, as well as the roadmap guiding its development, and its theoretical motivation and background. This chapter considers the system from a top-level, focusing particularly on the autonomous server and on the autonomous controller. The reader is reminded that both the server and the controller run within the system, and in particular that the system does not require external oversight. Dividing responsibilities between the server and controller is a matter of good design practice on one hand, and of development convenience on the other hand, but the system still behaves as a unified whole.

Figure 7.1 presents a block diagram of the autonomous server. Software running on the PowerPC can drive the ICAP, thereby reconfiguring the configurable fabric. It should be noted that although the configurable fabric is shown adjacent to the PowerPC, it actually encompasses nearly everything in the FPGA, including the two PowerPCs and the ICAP.

Shown separately in the top right corner of the PowerPC/Software block are Linux and the JRE, both discussed in Chapter 6. These are of course not part of the server, and instead provide infrastructure for it. What remains in the PowerPC/Software block is either part of the server itself or of the custom drivers, though only the `/dev/icap` driver is shown here.

The autonomous controller is a lightweight application that monitors activity from the sensor stubs, and handles mouse and keyboard input, and generates directives for the server. The server in contrast handles the overwhelming majority of the system complexity, including the placing and routing, the model synchronization, and the reconfiguration.

This chapter begins with an examination of the server’s protocol, architecture, and library support. It then discusses the autonomous controller, before turning to the server itself and discussing its many components.

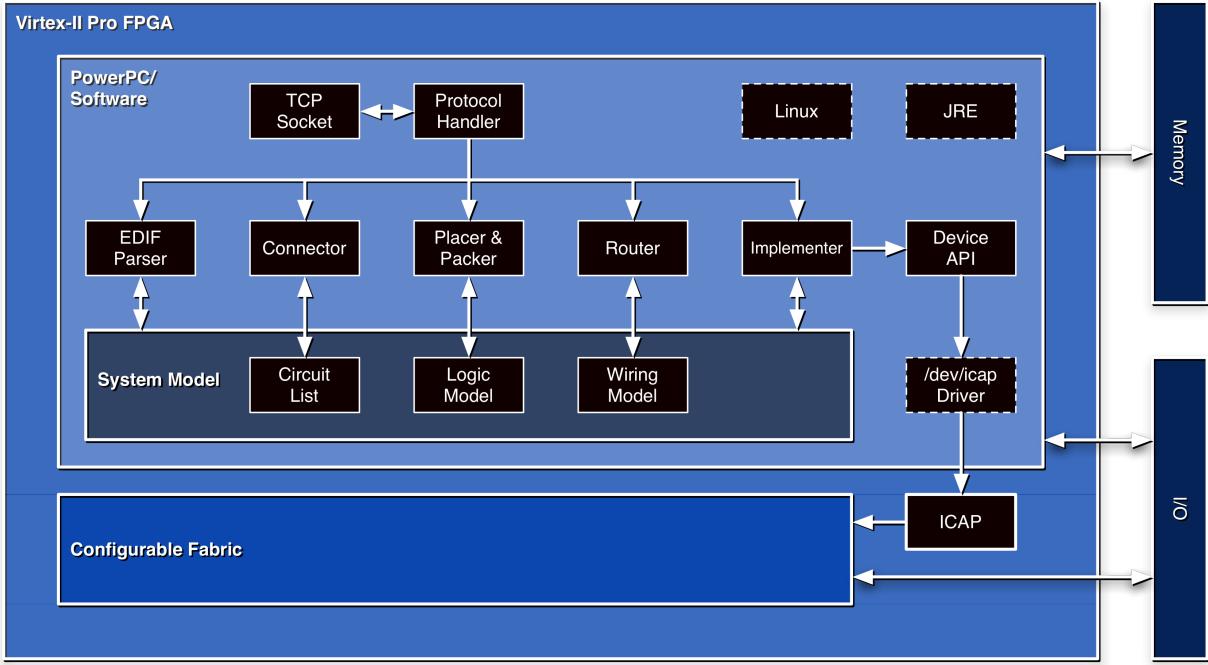


Figure 7.1: Server block diagram. Commands are received through a TCP port, and dispatched to the proper subsystems. Reconfiguration commands process pending changes, send the data to the Device API, and on to the ICAP through the /dev/icap driver. Although the configurable fabric is shown outside the PowerPC/Software block, the fabric actually encompasses the PowerPC.

7.2 Capabilities

The required capabilities of the system have a large influence on its design. Factoring into this are the Virtex2P architecture, the netlist primitives, the system model, and the separation between the server and the controller.

7.2.1 Protocol Specification

The fact that the server and controller are separate applications means that all of their interaction must adhere to some kind of protocol. That protocol is in turn influenced by the kinds of operations that the server and controller need to describe. In broad terms, the controller must be able to pass netlists to the server, request that those netlists be parsed, placed, routed, configured, connected, or removed, and request that the system be reconfigured. The actual protocol is text-based, and easy to invoke manually while developing or debugging. The protocol commands are as follows:

parse *circuit filename*: Parses a new circuit from EDIF file *filename*, extracting all primitives, cells, and nets, and associates it with designator *circuit* for future reference. *The server and controller handshake in the background to transfer the file contents.*

place *circuit*: Places the specified circuit. If top-level connections to the circuit are known, they may influence the placement. Can only be invoked on a circuit that was previously parsed.

route *circuit*: Routes the specified circuit. Can only be invoked after the circuit has been placed.

configure *circuit*: Configures all logic cells for the specified circuit. Can only be invoked after the circuit has been placed.

reconfigure: Completes any pending top-level routing, flushes all pending changes, generates a partial dynamic reconfiguration bitstream, and writes the bitstream to the icap driver. *This command is the only one that actually reconfigures the system. If no changes are pending, the command does nothing.*

rm *circuit*: Removes the specified circuit, and updates its top-level connections.

connect *source sink1 [sink2 ...]*: Connects the specified top-level ports, creating a new top-level net if necessary. Any number of sinks may be specified. Bus notation is supported.

disconnect *source|sink*: Disconnects the specified top-level port, removing the entire top-level net if necessary. If a source wire is specified, the entire top-level net is removed. If a sink wire is specified, only the branch driving that sink is pruned. Bus notation is supported.

A few additional commands are useful primarily in interactive or debugging contexts:

ls [*circuit*]: Lists the specified circuit, or lists all circuits—including the system circuit—if no circuit is specified.

-a: Lists top-level ports with each circuit.

-n: Lists top-level nets instead of circuits.

dump_xdl: Exports a top-level XDL description of the system, including all top-level connections. *Useful primarily in a debugging context.*

exit: Shuts down the controller, but leaves the server running. *Useful only in a debugging context, and fatal in an autonomous context.*

shutdown: Shuts down both the server and the controller. *Useful only in a debugging context, and fatal in an autonomous context.*

The controller also supports a *run* command that runs scripts. Run commands are not passed along to the server, but are instead recursively processed in the controller by sending each command in turn to the server.

7.2.2 Architecture Support

The Virtex2P architecture [18] includes 27 different types of logic cells. Most of them are configurable, but a few are not, and many are entirely unknown to most FPGA users. The reader is advised that these descriptions represent a best understanding, and should not be taken to be authoritative in cases that Xilinx does not document. Furthermore, not all of these cell types are available in every device in the family.

General Logic Cells:

SLICE: Standard logic cell. Highly configurable. Provides two 4-input Look-Up Tables (LUTs), two flip-flops, two multiplexers for higher-ordered functions, dedicated carry logic, dedicated sum-of-products (SOP) logic, and more.

TBUF: Tristate buffer. Not configurable, except for input signal polarity. Provides shared access to dedicated tristate wiring.

VCC: Logic 1 source. Not configurable. Provides constant logic 1, routable to other cell inputs.

Hard Core Cells:

RAMB16: 18,432 bit dual-ported memory (16,384 data bits plus 2,048 parity bits). Configurable as 512×36 , $1,024 \times 18$, $2,048 \times 9$, $4,096 \times 4$, $8,192 \times 2$, or $16,384 \times 1$, independently on each port.

MULT18X18: 18×18 unsigned multiplier. Configurable for synchronous or asynchronous operation.

PPC405: PowerPC 405 RISC processor. Not configurable, except for input signal polarity.

Clocking Cells:

BUFGMUX: Global clock tree buffer. Configurable. Allows clocks or special signals to use the high-fanout low-skew clock trees.

DCM: Digital clock manager. Highly configurable. Provides clock multiplication and division, deskew, and a variety of related functions.

Input/Output Cells:

IOB: Single-ended input/output block. Configurable. Provides access to corresponding package pad. Can be configured as input, output, or bidirectional, for asynchronous, synchronous, or Double Data Rate (DDR) operation. Supports a wide range of input/output standards.

DIFFM: Differential input/output master. Configurable. Can be used as single-ended, or the master of a differential pair. Supports special differential input/output standards.

DIFFS: Differential input/output slave. Configurable. Can be used as single-ended, or the slave of a differential pair. Supports special differential input/output standards.

DCI: Device-controlled impedance block. Configurable. Controls termination impedance on an input/output bank basis to support input/output standards requiring impedance matching.

PCILOGIC: Special PCI logic. Not configurable. Special support for high performance PCI output clock enable.

Gigabit Transceiver Cells:

GT: Gigabit transceiver. Highly configurable. Available in Virtex-II Pro but not Virtex-II Pro X. Supports a wide range of serial gigabit standards.

GT10: Multi-gigabit transceiver. Highly configurable. Available in Virtex-II Pro X but not Virtex-II Pro. Supports a wide range of serial gigabit standards.

GTIPAD: Dedicated gigabit/multi-gigabit input. Not configurable.

GTOPAD: Dedicated gigabit/multi-gigabit output. Not configurable.

CLK_N: Dedicated gigabit/multi-gigabit negative reference clock input. Not configurable.

CLK_P: Dedicated gigabit/multi-gigabit positive reference clock input. Not configurable.

Device Control Cells:

GLOBALSIG: Global signal control. Configurable. Controls the propagation of global signals GSAVE, GSR, GWE, and GHIGH, permitting these signals to be disabled in any desired clock domain.

CAPTURE: Global capture signal. Configurable. Triggers capture of state plane data for the sake of readback.

STARTUP: Global startup inputs. Not configurable, except for input signal polarity. Controls GTS, GSR, and the startup clock.

ICAP: Internal configuration access port. Not configurable, except for input signal polarity. Provides access to the configuration controller, and supports active configuration and readback from within the device.

Miscellaneous Cells:

BSCAN: Boundary scan interface (JTAG / IEEE 1149.1). Not configurable. Allows designs to interact with external boundary scan chains.

JTAGPPC: PowerPC JTAG interface. Not configurable. Provides access to the PowerPC JTAG port(s).

Testing Cells:

PMV: Internal ring-oscillator. Not configurable. Used for internal testing only, and not characterized by Xilinx.¹

RESERVED_LL: Internal test stub. Not configurable. Purpose unknown.

¹See comp.arch.fpga thread “Xilinx - one secret less, or how to use the PMV primitive”, posted 30 Aug 2006 02:52:49 -0700 by Antti Lukats. Note response “Re: Xilinx - no secret, you are not to use the PMV primitive” and following, posted Wed, 30 Aug 2006 07:33:26 -0700 by Austin Lesea.

Of these cell types, the most prevalent by far is the SLICE, accounting for 1,408 cells in the XC2VP2, and 44,096 cells in the XC2VP100. By convention, the term *SLICE* will usually be written *slice*, except in cases where emphasis is placed on the specific cell type.

7.2.3 Library Support

Xilinx provides access to over 250 functional design elements through the Xilinx Unified Libraries [20]. These elements are made available as *primitives* that can be referenced from HDL code. To retain a large measure of architecture independence, most primitives avoid directly referencing the cell types available in the supported architectures. Because these same library primitives appear in the generated EDIF netlists, the system must map them to cell types and corresponding cell configurations.

The mapping between library primitives and Virtex-II Pro cells may be one-to-one, one-to-many, or many-to-one. For example, a MULT18X18 library primitive corresponds exactly to a Virtex-II Pro MULT18X18 cell, so the mapping is one-to-one. By comparison, the ROM64X1 primitive requires four SLICE cells, in a one-to-many mapping.

The most common mapping by far is actually many-to-one, because of the way that Xilinx subdivides slices. The Unified Library primitives MUXCY, XORCY, ORCY, MUXF5, MUXF6, MULT_AND, must all be packed into slices, along with all flip-flop and latch primitives², and all LUT primitives and derivatives (AND, OR, XOR, NAND, NOR, XNOR).

The server internally models slices as a collection of subcomponents, with names following Virtex-II Pro conventions, rather than Xilinx Unified Library conventions:

LUT (FLUT or GLUT): Look-up table. Can implement an arbitrary one bit boolean function of four inputs. Can also be used as a ROM, or RAM, including the special case of a shift-register.

FF (FFX or FFY): Flip-flop. Can also function as a latch. Can include synchronous or asynchronous set and reset.

MUXCY (CYMUXF or CYMUXG): Carry chain multiplexer. Provided to support fast adders, wide functions, or sum-of-product constructs.

XORCY (XORF or XORG): Carry chain XOR gate. Provided to support fast adders.

ORCY: Carry chain OR gate. Provided to support sum-of-product constructs.

F5MUX: F5 multiplexer. Provided to support wide functions. Generates a function of the F LUT and G LUT outputs.

F6MUX: F6 multiplexer. Provided to support wide functions. Generates a function of two F5 outputs from up to four slices.

²When appropriate, flip-flops or latches may instead be packed into IOB cells.

AND (FAND or GAND): Carry chain AND gate. Provided to support multipliers.

These subcomponents are generally not independent, but instead share some wiring or settings, or are subject to mutual constraints. The net result is that both the placer and the mapper must be careful not to violate any applicable rules, as detailed more fully in Section 7.4.8. Some of these rules can be satisfied in advance by grouping components into standard forms, but all other rules must be checked by the placer for every candidate placement position, adding a significant penalty to the placer performance.

Another kind of one-to-many mapping exists in the case of device input/output. The differential master (DIFFM) and differential slave (DIFFS) cells in the Virtex-II Pro architecture can also function as regular single-ended IOB cells. In fact the Virtex-II Pro architecture contains almost no plain IOB cells, so the mapper and placer must be told that they can allocate DIFFM or DIFFS cells and transparently use them as IOB cells.

7.3 Autonomous Controller

The autonomous controller is responsible for monitoring changing inputs, controlling the server, and interfacing with the user. Unlike the server, that was developed in Java for compatibility with other CAD tools, the controller was developed in C to simplify interfacing with Linux device drivers.

The user display shows four distinct sections on a 1280×1024 screen, as seen in Figure 7.2. The top right section shows the real-time resource usage of the FPGA that the system is running on. A cursor controlled by a PS/2 mouse allows the user to mask or unmask resources, marking them unusable or usable for dynamically placed logic. The lower right section is a small console area through which the user can send commands to the server, and view responses from the server. The upper left section is reserved for 720×487 NTSC video when the video capture core is instantiated. And the lower left section is reserved for 128 bars of FFT display when the audio core is instantiated.

The controller begins by opening the following devices:

/dev/ps2/mouse: For mouse input. The controller pins the mouse position inside the FPGA section, allowing the user to click on resources to mask or unmask them. Mask data is retrieved by the server before each placement run.

/dev/ps2/keyboard: For keyboard input. The controller provides a miniature console through which the user can send commands to the server, and view its response.

/dev/fb: For SXGA monitor output. The controller maps the framebuffer memory into its address space to draw on the display. The fb device is also used to turn the cursor and the FFT display on or off.

/dev/fpga: For live resource usage visualization. The controller and the server use this to exchange resource mask information, and the server uses it to update placement information.

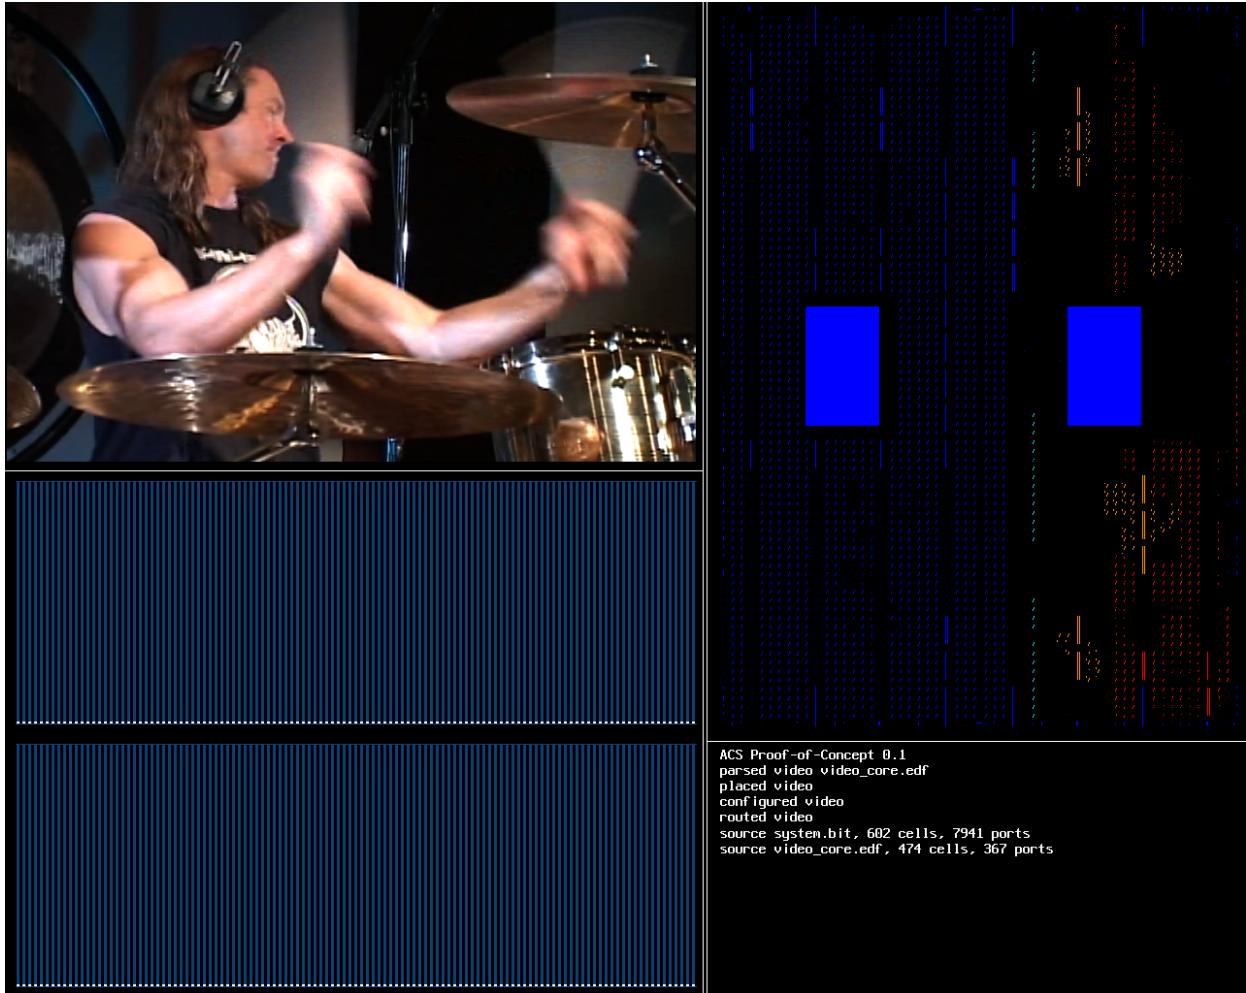


Figure 7.2: Autonomous controller screen capture. *Note that the FFTs cannot be captured and are instead simulated here. Their data goes directly from the audio core to the framebuffer core, without passing through main memory, and is therefore not available to the screen capture utility.* Video frame captured from [80].

/dev/stub: For notification of input changes. The controller receives notification through the stub driver whenever a signal is acquired or lost from the mouse, the keyboard, the video decoder, or the audio codec. These inputs are intended to trigger action on the part of the controller, and if appropriate, to cause the system to modify itself.

The controller initializes the various devices through their drivers, initializes the console area through which the user and the controller can interact, and initializes the framebuffer and draws the static portions of the display. It then creates a TCP/IP socket and initializes the client that

will interact with the server. The controller then registers the appropriate devices for polling,³ configures the mouse, and configures the keyboard.

When all of that initialization is complete, the controller tries to open a connection to the server, and pauses if need be while waiting for the server to start up. Once it has successfully connected to the server, it samples the FPGA resource utilization data that the server will have just finished updating, and it enters its event loop.

The event loop runs until the user explicitly sends an *exit* or *shutdown* command to the server. As long as it is running, it calls the Linux poll() function, which returns whenever any input becomes available from the mouse, keyboard, server, or any of the stubs. Each event is dispatched and processed by the appropriate handler.

If the user or the controller generate commands for the server, those commands are packaged and sent through the protocol handler to the TCP/IP socket and on to the server.

Figure 7.3 shows the resource utilization of the base system upon startup, colored in green. This gives an indication of the area overhead necessary to support autonomy on this particular device. Figure 7.4 shows an example of resources that the user has interactively masked out, colored in orange. And Figure 7.5 shows an example of the resource utilization after the video capture core and the audio FFT core have both been placed.

The reader has probably noted that there is much room for complexity and sophistication in the autonomous controller, but the present work has primarily aimed to lay the foundation for a Level 6 system. As such the controller does just enough to prove the concept, and everything beyond that is deferred for future work.

7.4 Autonomous Server

The server software handles everything pertaining to circuit instantiation and removal. Every request from the autonomous controller or from remote debugging clients, is parsed and processed by the server, and all reconfiguration takes place under its control, as depicted in Figure 7.1. Although the server makes no decisions about what circuits are instantiated, it does control how the circuits are placed and routed. In terms of the roadmap, the server functions as a Level 2 system, while the autonomous controller functions as a Level 6 system layered on top of the server's lower-level capabilities.

For compatibility with the Device API and with ADB, the server software is written entirely in Java. Its software is complex, and the discussion touches on some related issues, before turning squarely to the relevant subsystems that do the work.

³The reader is reminded that this “polling” is in fact extremely light on the controller, the drivers, and the operating system, and that it is the accepted manner of doing things under Linux.

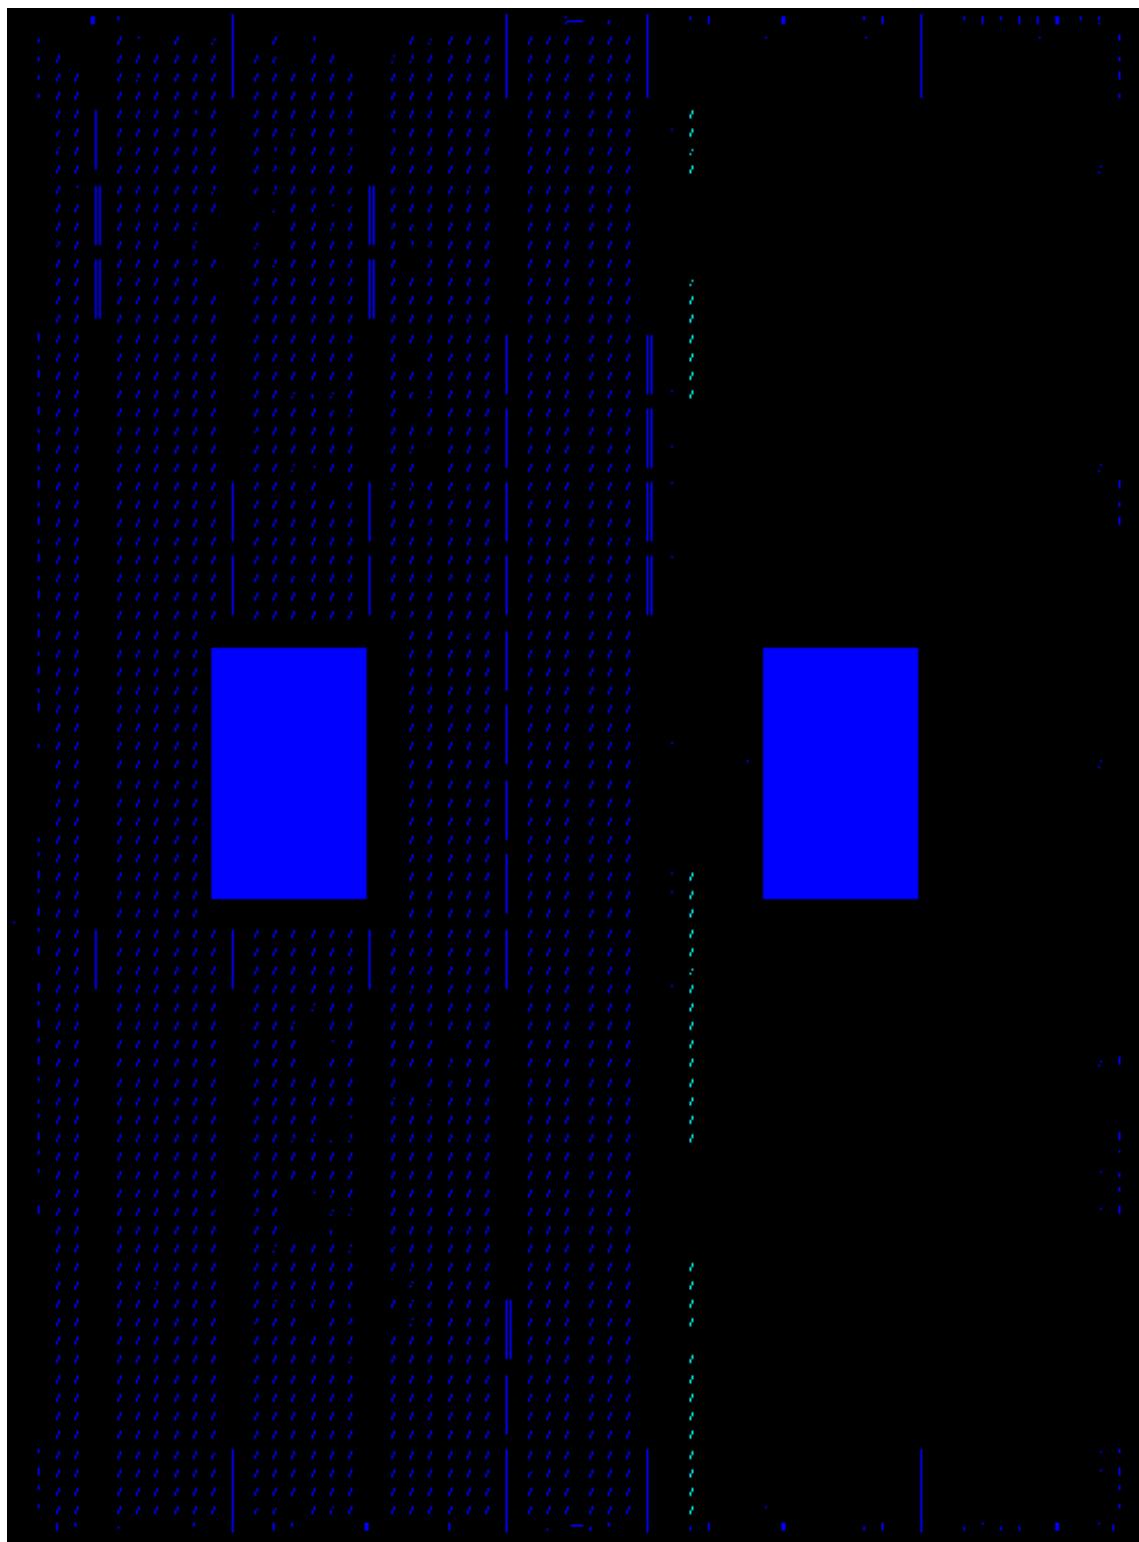


Figure 7.3: Resource usage at startup.

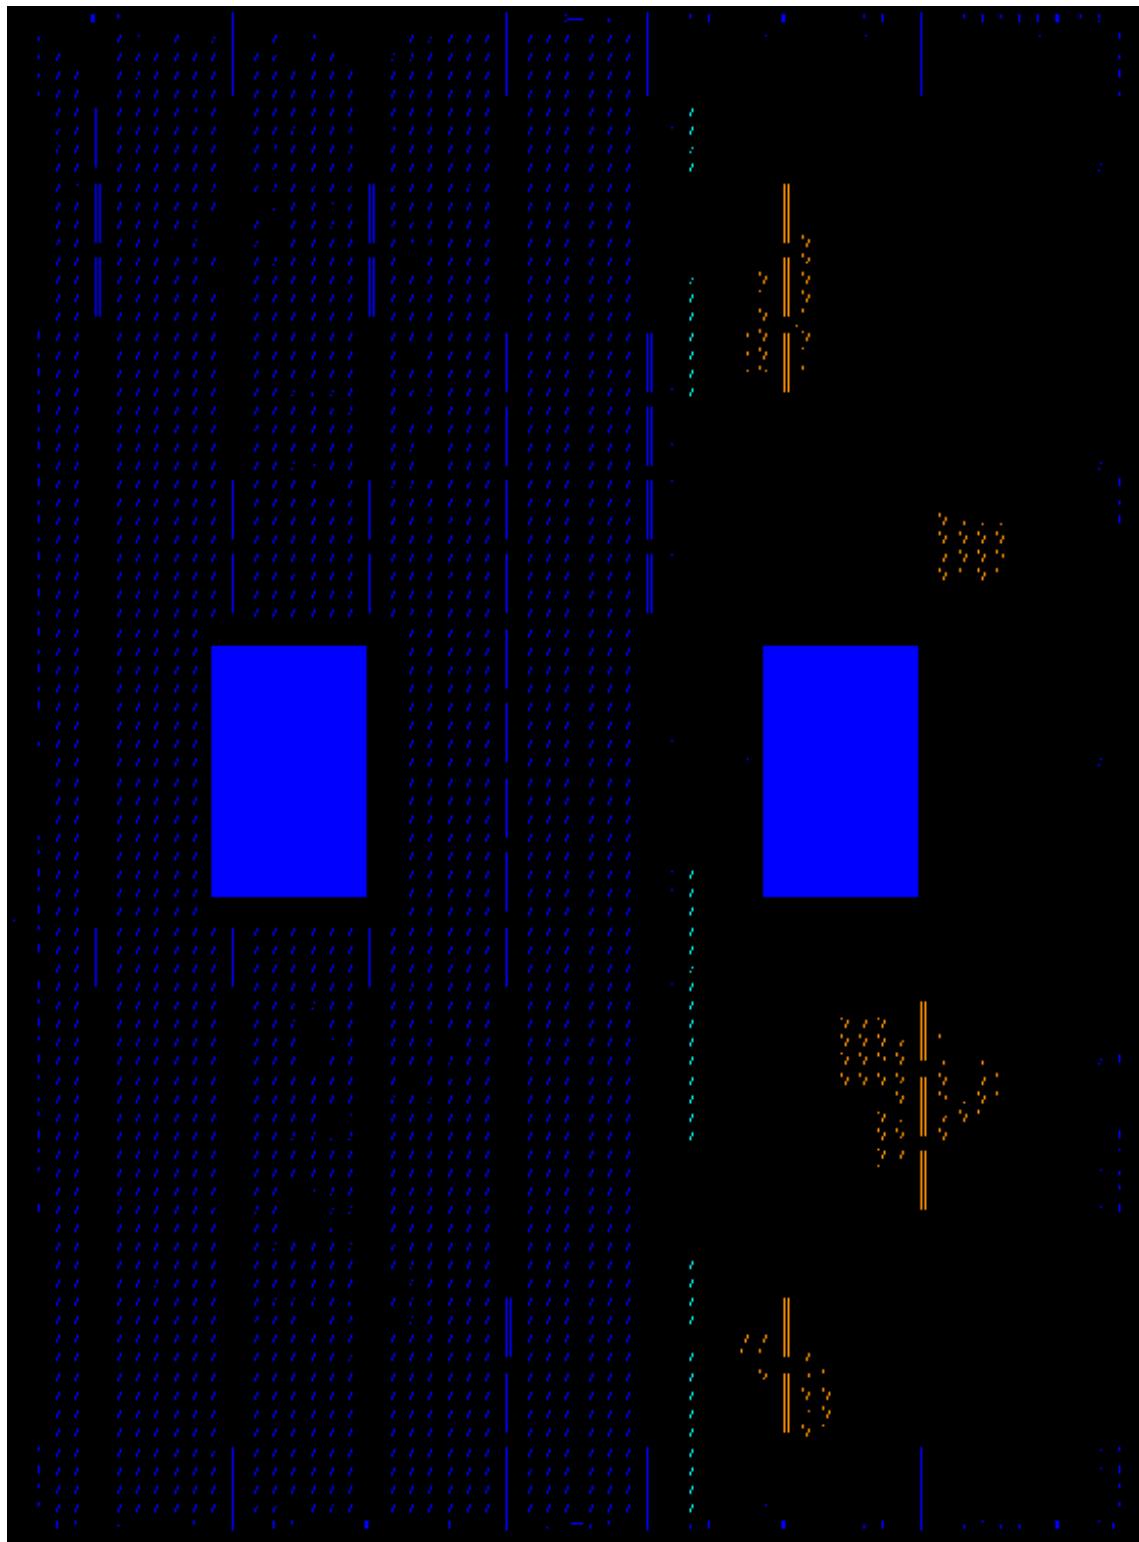


Figure 7.4: Resource usage after interactive resource masking.

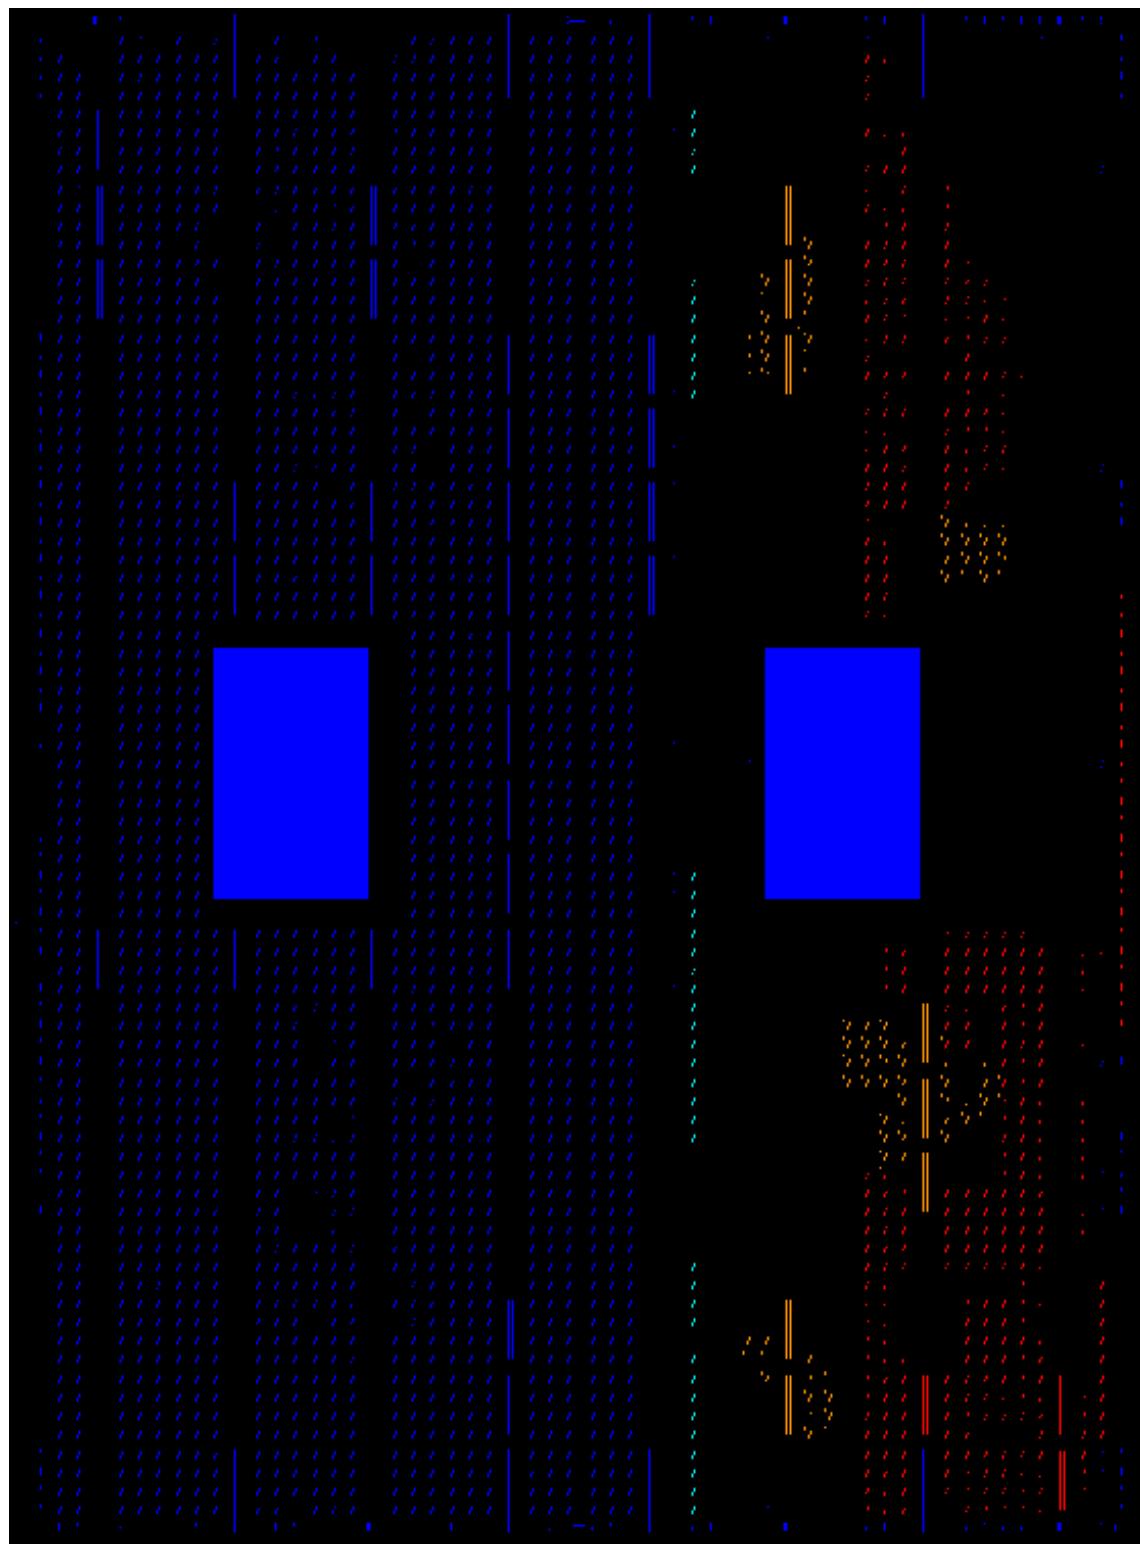


Figure 7.5: Resource usage after placement of video capture and FFT cores.

7.4.1 Circuits

The server’s model tracks every circuit that it knows about, including the system circuit itself. Every circuit is associated with a designator provided by the client, and all subsequent functions involving that circuit reference it by its designator.

The system circuit is a subclass of the regular circuit class. It is created by the server at startup, and it exists primarily to expose system ports to user circuits. These ports include 1) all IOB, DIFFM, and DIFFS cells, 2) all gigabit transceiver inputs, outputs, and clock references, 3) all global clock buffer outputs, and 4) all of the relevant wires and busses for sensor stubs and interface stubs. Access to these ports effectively allows circuits to interact with the system and its environment, as would be expected with a regular hardware circuit.

In keeping with the modelling discussion of sections 3.5 and 7.4.4, a system must model everything that it may change. Although all of the cells and wires used by the system are internally tracked to prevent the router and placer from interfering with them, it is not necessary to associate them with the system circuit, as that would consume extra memory while providing no clear benefit.

Once a circuit has been placed and routed, its net and cell lists are retained primarily to facilitate circuit removal. During removal, all device cells that belong to the circuit are released, and all net connections to the circuit are intelligently disconnected.

7.4.2 Protocol Handler

When the server starts, it opens a TCP/IP server socket on port 49,152,⁴ and listens for connection requests. Once a connection is established with a client—typically the autonomous controller—the server invokes the protocol handler.

The protocol handler receives text-based commands from the client, parses them, executes them, and notifies the client when the processing is complete. The only large data transfer between the two parties occurs when the client requests that a new circuit be parsed. At that point the server acknowledges the request, and the client begins to stream the circuit’s EDIF netlist to the server. Once the netlist has been fully parsed, the server again notifies the client, and both return to normal command processing mode.

⁴Port 49,152 falls within the range reserved for private or dynamic ports, and this default can be overridden when the server is invoked.

7.4.3 Cell Semantics

Three distinct kinds of logic cells are handled by the server. The first kind of cell is that of the EDIF netlist. EDIF cells are references to library components, and are in turn referenced by the nets that compose the netlist. EDIF cells are constructed during EDIF parsing.

The second kind of cell is the user design cell, related to an EDIF cell, though not necessarily in a one-to-one relationship. Because an EDIF cell may correspond to an exact device cell, or part of a device cell, or a collection of device cells, user cells are tied to a *primitive* that encapsulates the functionality of the EDIF cell, and participates in configuration and remapping operations. User cells are constructed by the applicable primitive after the entire EDIF netlist has been parsed.

The third kind of cell is the device cell that the system actually allocates and configures. User cells correspond to pseudo device cells. In some cases a single user cell will map to a device cell, while in the case of slices subcells, multiple user cells may map to a single device cell.

7.4.4 Model

The need for an autonomous system to model itself was discussed at length in Section 3.5. In simple terms, the model must support the same granularity and level of completeness as the changes that the system needs to make to itself. In the present work, the granularity is that of the smallest configurable element in the Virtex2P, and the level of completeness is just shy of the entire device. As previously discussed, neither LUT RAM resources nor tristate buffers⁵ are supported.

The model is composed of multiple facets. As discussed in Section 2.3.1, the Virtex2P is actually described by a *user* model as well as a *device* model, with the user model describing the CAD tool view of the Virtex2P, and the device model describing the configuration bitmap view of the Virtex2P.

In addition to the Virtex2P user and device facets of the model, the system also models itself in terms of dynamically instantiated circuits. From the time a circuit has been parsed, until the time that circuit is removed, the system knows about each of its nets and logic cells. This is primarily necessary to permit the system to remove circuits when necessary, but it could theoretically also allow the system to perform housekeeping chores on instantiated circuits, perhaps rearranging them or performing some form of online optimization.

The user and device facets of the model are fully provided by the Device API, while the circuits facet of the model is provided by the server itself. These three facets of the complete model are

⁵The lack of support for tristate buffers is largely due to a blind spot. On one hand, tristate lines implicitly involve multi-source nets, and the router currently doesn't know how to describe such nets. On the other hand, tristate buffers can add significantly more complex constraints to the placement problem, because tristate lines only connect horizontally within the same row. This omission is believed not to be too big of a limitation to the work, and it is worth noting that internal tristate buffers disappear entirely in Virtex4 and newer architectures.

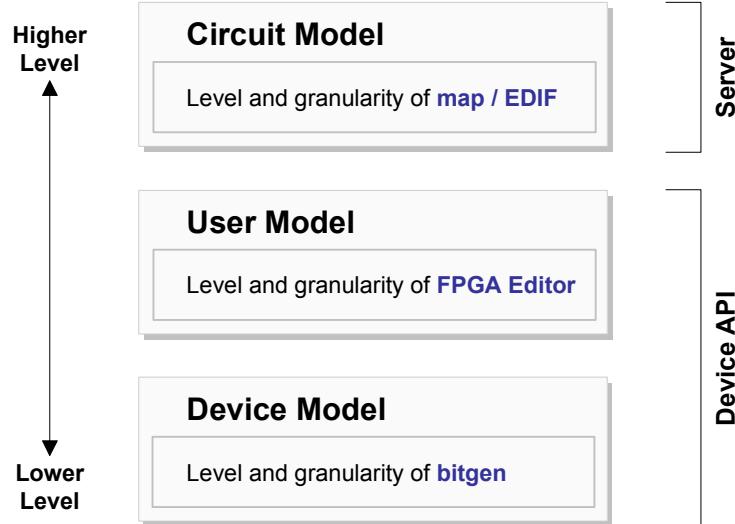


Figure 7.6: System model components.

fairly tightly coupled in spite of their differing features and granularities, and they together support all of the operations that the system performs upon itself.

Figure 7.6 shows the three facets of the system model, and their respective granularities. The device model is comparable in context and granularity with bitgen. The user model is comparable with FPGA Editor. And the circuit model is comparable with map and/or EDIF netlists. Although the model does not exist as a single distinct element inside the server, it is an integral part of the server's operation.

7.4.5 Libraries

Support for the Xilinx Unified Libraries was discussed in Section 7.2.3. Only the subset of library primitives necessary for the various dynamic cores has been implemented at present, but remaining primitives can be added as needed, and most are trivial to implement.

The underlying user model already supports all 27 logic cell types in the Virtex2P architecture, and most of the primitives in the Unified Libraries map to only one cell type. Because these primitives were designed for HDL coding simplicity, most of them are simple variations of each other, with only basic connection or configuration differences.⁶ That explains in large part how many hundreds of library primitives can map to roughly two dozen cell types.

⁶The reader may note that HDL instantiation of these components usually consists only of the library primitive name and the port mappings, with no parameter specifications. For cases that do require parameter declarations, those parameters usually correspond directly to cell configuration settings, and are thus passed along unchanged.

From a coding perspective, it is generally necessary to create a single class for each underlying cell type, and then to create subclasses for each supported library primitive. Most of the subclasses implement only minor configuration changes. Slice subcells are already transparently supported by a class that encodes the full complexity of that cell type,⁷ and most slice primitives merely need to define their names, because their usage is implicit in their port declarations.

The current list of supported library primitives may appear short, but it supports a wide range of dynamically instantiated circuits. Furthermore, deferring the coding of a new library primitive until a circuit actually requires it, streamlines the testing process by providing an explicit test case for it. The supported list, grouped by cell type, is as follows:

IOB/DIFFM/DIFFS (Inputs and Outputs): IBUF, IBUFG, OBUF

MULT18X18 (Multipliers): MULT18X18, MULT18X18S

RAMB16 (Block RAMs): RAMB16_S4_S4, RAMB16_S9_S9, RAMB16_S18_S18,
RAMB16_S36_S36

SLICE:

Carry Chain Elements: MULT_AND, MUXCY, MUXCY_L, XORCY

D Flip-Flops: FD, FD_1, FDC, FDC_1, FDCE, FDCPE, FDE, FDP, FDPE, FDR, FDRE,
FDRS, FDRSE, FDS, FDSE

Look-Up Tables: LUT1, LUT1_L, LUT2, LUT2_L, LUT3, LUT3_L, LUT4, LUT4_L

Utility: GND, VCC, INV

Wide Function Muxes: MUXF5, MUXF6, MUXF7

BUFGMUX (Global Clock Buffers): BUFG, BUFGP

The behavior and properties of these primitives are defined and documented in the Xilinx Libraries Guide [20].

7.4.6 EDIF Parser

The Configurable Computing Lab at Brigham Young University (BYU) provides two separate EDIF processing tools. A complete and robust tool set supports reliability studies for FPGAs exposed to radiation [81]. An alternate and lightweight tool set includes code and a grammar that is suitable for parser generation [82].

The edif.jj grammar available with the lightweight tool set can be preprocessed with UCLA's Java Tree Builder (JTB) [83].⁸ The resulting output can then be compiled by JavaCC, the Java Compiler Compiler [84], to generate a full parser for EDIF files.

⁷As previously discussed, this excludes shift-register functionality.

⁸The current version of JTB required the use of Java2 1.5, but because the target platform is limited to Java2 1.4.2, it was necessary to instead use JTB 1.2.2. This switch did not interfere with the use of JavaCC 4.0, and did not seem to adversely affect the performance or behavior of the resulting parser.

The BYU EdifParser class provides the functionality needed to parse EDIF streams, and to build syntax trees according to the edif.jj rules. The unaltered EdifVisitor class provides a way to visit every token in the syntax tree, but does only minimal processing on the tokens. That processing capability must instead be added to the class, to extract the library references and the actual cells and nets that are to be implemented.

Part of the difficulty in extracting EDIF information comes from the fact that many EDIF constructs can be infinitely nested, which makes the extraction much more complicated. In practice however, generated EDIF is not as unwieldy as one might fear, and a decision was made to only accept flattened netlists. Flattened netlists are easily generated, and their use sidesteps the majority of the nesting issue. Should it become necessary to support hierarchical netlists in the future, it may be desirable to use BYU's comprehensive EDIF tools, even though the resulting parsing would require significantly more time and memory.

As the EDIF reader visits the various EDIF tokens in depth-first mode, it builds structures that describe the library references, as well as the cells and nets that are encountered. After all tokens have been visited, the final cell and net data is added to a circuit object, where it resides as long as the circuit remains in use.

Before the parsing begins, the parser maps each of the supported library primitives to classes that provide the necessary functionality. This permits the EDIF reader to associate each cell in the circuit with the appropriate primitive.

A final step permits the circuit design object to fix up internal cells or nets. This was necessary to remove certain buffers that Synplify Pro had inserted but had confusingly flagged for internal routing, and that were consequently failing to route.⁹ Another useful change would be to absorb inverters on cell inputs into the cells themselves, but this is not yet supported. It is believed that these and many other adjustments are routinely handled by the *map* tool in the regular ISE flow.

7.4.7 Placer

Although much research has been published on placement inside FPGAs, almost all of it assumes a uniform $2D$ array of identical logic cells, an assumption that is far too simplistic for modern FPGAs. Modern FPGA architectures instead include dozens of cell types, of differing sizes, and occurring in widely varying numbers and patterns, frequently poking large holes in the assumptions. Two placement tools that were available might have helped in the current work. Unfortunately the Versatile Place and Route (VPR) tool developed by Betz and Rose [48] was unusable because of its unfamiliarity with true device layouts. The incremental placement tool developed by Ma [49] is similar in that it targets the Virtex architecture. These tools were very good for the architectures

⁹ *Intra*-slice routing is handled by configuration settings rather than by the router, so there is no good way to connect internal ports through *inter*-slice wiring. Since the buffers in question serve no useful purpose, they are simply removed.

and abstractions that they targeted, but they are unable to solve the more general problem posed by modern architectures.

The present work is designed to support *all* of the logic cells in a device, and that means that the placement problem must be solved for the more complex general case. However although placement is a necessary part of this work, it does not constitute its main focus, and the quality and performance of this placer are presumably less than they would be if developed by research groups with expertise in placers.

Placement algorithms fall into three broad categories, as defined by Sherwani [46]: 1) constructive, 2) analytic, and 3) random. The constructive approaches typically begin by searching for mincuts, to partition the circuit graph, and then proceed by constructing a placement. The analytic approaches set up systems of equations to express requirements and constraints, which they solve to derive a placement. And the random approaches begin with a random placement to which they apply perturbations while seeking to minimize some cost function.

When the problem domain is well understood, and particularly when the device geometry is regular and homogeneous, the constructive or analytic approaches are good choices. But the complexity of real FPGA architectures demands that considerably more expertise be applied for those approaches to work. An alternative is to use a random placement algorithm, that understands very little about its problem domain, and simply applies perturbations that are checked for legality. It is worth noting that many commercial and research placers appear to be hybrids, using at least some measure of annealing, but first structuring the problem into a more manageable form.

7.4.7.1 Abandoned Algorithm: Simulated Annealing

For the sake of simplicity, this work initially implemented simulated annealing [46, 85] as a means of generating a placement solution. Simulated annealing is inspired by the annealing process in metals, beginning at a very high temperature where molecules move around widely, and gradually cooling to lower temperatures where molecules settle into a global energy minimum. Pseudo-code for the algorithm is presented in Figure 7.7. From a placement perspective, simulated annealing can produce good results, but takes a long time to run.

The simulated annealing algorithm uses exponential functions, which in turn require the use of floating point operations in software implementations. But floating point operations are costly on processors such as the PowerPC 405, that do not have dedicated Floating Point Units (FPUs), and that must instead emulate the functionality much more slowly in software. If data values fall within certain ranges, it is sometimes possible to substitute scaled integers for floating point values, and to precompute exponentials. It is also possible to make use of table look-ups in certain cases, as done by Sechen et al in TimberWolf [85]. The code developed for this implementation attempts to reduce the costly floating point overheads, but does not try to optimize the algorithm: What is really needed is an algorithm that can perform many orders of magnitude faster, and that kind of speedup cannot be achieved with mere coding optimizations.

```

Algorithm SIMULATED-ANNEALING
begin
    temp = INIT-TEMP;
    place = INIT-PLACEMENT;
    while (temp > FINAL-TEMP) do
        while (inner_loop_criterion = FALSE) do
            new_place = PERTURB(place);
             $\Delta C$  = COST(new_place) - COST(place);
            if ( $\Delta C < 0$ ) then
                place = new_place;
            else if (RANDOM(0,1) >  $e^{\frac{\Delta C}{T}}$ ) then
                place = new_place;
            temp = SCHEDULE(temp);
    end.

```

Figure 7.7: Simulated annealing algorithm. Reproduced by permission from Figure 5.4 in [86].

The perturbations that are applied to the placement are coded so as to be reversible. Because the cost function that evaluates the suitability of a perturbation is quite complex, it is easier to first apply and evaluate the perturbation, and then simply reverse it if the algorithm rejects it. Part of the cost function includes the net length, but this implementation follows the TimberWolf approach in estimating the length of a net as half of its bounding box circumference.

Although the regular simulated annealing algorithm calls for the number of perturbations to be some large multiple of the number of cells, it was modified to instead be a multiple of the number of placeable objects, such that a supercell¹⁰ would count as a single placeable element. This modification was made after observing that large adders that consisted mostly of a carry-chain supercell, were being moved around unnecessarily even though the entire supercell always retained its fixed geometry.

Although the performance of the simulated annealing placer was suitable for a proof-of-concept, it wasn't sufficient for demonstration purposes. Placing a 45 cell circuit with 500 perturbations per cell per temperature step, required 90 seconds of actual processing time on a 2.4 GHz Pentium 4 with a 512 kB cache and 757 MB of main memory. Those numbers should be considered in light of the fact that a 32-bit counter requires 155 cells,¹¹ and that most circuits of interest require 1,000 cells or more. Furthermore, the demonstration system uses a 300 MHz PowerPC 405 with a 16 kB cache and 506 MB of main memory, and the PowerPC 405 does not include the high end pipelining and scheduling of the Pentium 4. Given performance more than an order of magnitude slower,

¹⁰Supercells, described in Section 7.4.7.4, are groups of cells that function together, frequently with fixed geometries. It therefore makes sense to place supercells as discrete blocks.

¹¹155 cells may seem surprisingly large for a 32-bit counter, but this includes all of the carry-chain elements necessary to achieve good timing performance.

the demonstration system would take over 15 minutes to place this relatively small circuit. That performance was unacceptable, and simulated annealing was therefore abandoned as a placement algorithm.

Simulated annealing begins with no understanding of the circuit that it is placing, and proceeds forward guided only by increases or decreases in the overall cost that it sees, much like a guessing game where one is only told whether they are getting “warmer” or “colder.” Given enough running time, the algorithm gives good quality results, but that same running time makes it better suited to traditional off-line ASIC or FPGA design, than to a dynamic and responsive system.

Observation of the annealing algorithm in action makes it clear that a lot of time is spent discovering a basic circuit geometry. In retrospect, it seems that a better approach might have been to start with a constructive placement algorithm, and then perhaps run a brief annealing phase to optimize the results.

7.4.7.2 Resource Map and Filtering

Any placement algorithm must have at its disposal the list of available cells. Some cells are used or reserved by the base system, while other cells are used by user circuits. Other cells can be filtered out as needed to support floorplanning or special requirements. Dynamic floorplanning is not currently implemented in this system, but it could be used to simplify the placement problem and to impose broader structure upon user circuits.

The cell availability is described by a resource map. The map tracks every logic cell in the device, and all subcells described in earlier sections. It also tracks the absolute and relative positions of cells, as well as their types and fully-qualified names.

Because of the manner in which the placer and mapper function, they must be able to efficiently look up cells by cell type, and must also be able to efficiently locate cells at relative offsets from arbitrary positions. To support these requirements, the cell map maintains three structures. The first one maps cell indexes to cell information. The second one maps cell types to arrays of cell information. And the third one maps cell positions by type to cell information.

For the XC2VP30 target device, the total number of cells plus subcells is 210,041. The resulting map structures consume 9 MB of memory.¹² This size penalty is undesirable, but the map structures permit much faster lookups than would otherwise be possible. When circuits depend upon relative placement constraints, the number of lookups can become very large. For example, a 32-bit counter which positions all of its cells relative to a central carry chain, makes 7,377,611 map lookups during placement, and it is believed that trying to search through simple lists would require a great deal more time.

¹²This is due in part to Java’s costly memory overhead for objects, as previously discussed in Section 4.6 of [22]. This memory usage should be seen as overhead associated with the device, and once initialized, the map remains invariant, no matter what circuits are implemented.

7.4.7.3 Constraints

A variety of constraints can affect placement, and while these may put more demand on the placer infrastructure, they also tend to simplify the rest of the placement.

Some constraints may never be violated. Logic cells that are already in use or reserved by the system must never be reused. Any user filtering of cells must be respected, to support floorplanning or special requirements. And logic constrained to internal or external ports, including device input/output pads or explicitly specified cells, must also be strictly respected. Each of these is treated as a *hard* constraint.

There are also certain *logical* constraints that stem from the architecture of the device. In the Virtex-II Pro architecture, dedicated carry chain resources like MUXCY and XORCY *must* be placed in exact ascending sequence within a slice column. Other structures, such as tristate TBUFs or sum-of-product blocks using ORCYs must be vertically aligned to use dedicated tristate or sum-of-product wiring. Various SLICE subcells may also need to be packed together to avoid large delays.

All of the logical constraints can be avoided by substituting regular LUT functions for special resource types, as was done in the JHDLBits project [50], but the resulting performance and area penalty would be severe. This system instead absorbs related logic into *supercells* that are placed as a group, relative to a common anchor point, in the manner of Relationally Placed Macros (RPMs) used by the Xilinx ISE tools.

7.4.7.4 Inferring Cell Groups

A handful of common and performance-critical functions are accessible in Xilinx FPGAs through dedicated but highly constrained routing resources. Although the placer must first infer relevant cell groups, a group that has been inferred can be treated as a single *supercell* for placement purposes. The inference step is essential to ensure the validity of the resulting placement, but it also simplifies the remainder of the placement.

Certain kinds of cell groups do not have fully fixed geometries. For example, it is possible to connect multiple drivers to the same tristate line, but that generally only requires that the TBUF driver cells reside in the same row, with few column restrictions. Grouping related TBUF cells would therefore result in a structure that was not fully rigid, but instead retained internal degrees of freedom. These kinds of structures are presently too complex for the placer to handle, and TBUF placement is therefore unsupported. Tristate lines implicitly involve nets with multiple drivers—and are in fact the only construct in the Virtex-II Pro architecture where such nets occur—but the reader is reminded that the router does not yet support multi-driver nets.

The following group types are supported:

Carry Chain: Combines connected MUXCY cells in ascending sequence on a carry chain. This group will also absorb any connected XORCY gates, with attached flip-flops if applicable, as well as LUTs driving the MUXCY select input, and AND (MULT_AND) gates driving the MUXCY 0 input. The carry chain can start and end at various places, but in general the carry propagates from SLICE_XnYm to SLICE_XnY(m+1), with two stages in each slice. The last MUXCY will include a final XORCY and LUT if present.

F5 Mux: Combines two LUT outputs, using predefined dedicated wiring, with one MUXF5 and two LUT subcells.

F6 Mux: Combines four LUT outputs, using predefined dedicated wiring, with one MUXF6, two MUXF5, and four LUT subcells. The slice positions are constrained to SLICE_XnYm and SLICE_XnY(m+1), where m is even.

F7 Mux: Combines eight LUT outputs, using predefined dedicated wiring, with three MUXF6, four MUXF5, and eight LUT subcells. The slice positions are constrained from SLICE_XnYm to SLICE_X(n+1)Y(m+1), where n and m are even.

FF-LUT Pair: Combines connected LUT/flip-flop pairs. This is a commonly inferred construct that both simplifies placement and improves timing. Connected FF-LUT pairs are not inferred in cases where the LUT drives more than just its local flip-flop.

Groups involving TBUF drivers and sum-of-product (SOP) blocks¹³ are not yet implemented, because both retain internal degrees of freedom that complicate the placement. Another reason for not yet implementing these group types, or F8 muxes, is that they have not yet cropped up in any instantiated circuits.

7.4.7.5 Placer Algorithm

The poor runtime behavior of the simulated annealing placer forced the decision to find a more suitable alternative. Although simulated annealing is appealing when no domain knowledge is available, it conversely holds that a little bit of domain knowledge can significantly reduce dependence upon questionable random perturbations. It is common for constructive algorithms to include low-temperature annealing as an additional opportunistic step, glancing around the vicinity of the current solution—much like shaking a pile of pebbles to flatten it. There is objective reason to believe that some kind of hybrid approach would also have worked well here, but it remains that placement is not the prime focus of this work, and the search for really good solutions is respectfully left to the true placement experts.

Many possible approaches were considered, and even coded to varying degrees, but were ultimately discarded. In most cases—including the *mincut* approach for example—the algorithms did well for small circuits, but scaled poorly as the circuit size increased beyond some nominal point.

¹³SOP blocks internally use carry chains, and connect to horizontally top-aligned adjacent blocks to implement wide functions.

The current algorithm is presented in Figure 7.8. It begins by building clusters of cells that are neither *special*, nor *synchronous*. In practice, that means that the placer grows clusters of *combinational* logic, but does not pull in cells like RAMB16s or IOBs. In clustering asynchronous cells, the placer tackles the islands of logic between registers—since that logic will likely form the *critical paths*¹⁴—and momentarily leaves special cells to their own coordinate systems.¹⁵

```

1 | Prepare clusters
2 | Grow clusters of asynchronous slice cells
3 | Absorb synchronous cells and special blocks
4 | Determine cluster resource and geometry constraints
5 | Build 2D weighted cluster connectivity matrix
6 | Order clusters by connectivity weight
7 | Place clusters
8 | Iterate starting at most tightly connected cluster
9 |   Determine acceptable cluster aspect ratios
10 |   Iterate through candidate positions
11 |     Allocate position that minimizes wire length
12 |   Prepare cells
13 |   Build 2D weighted cell connectivity matrix
14 |   Place all special supercells in cluster
15 |   Order cells by connectivity weight
16 |   Place cells
17 |   Iterate through candidate positions
18 |     Enforce cells packing restrictions
19 |     Place cells by connectivity weight
20 |   Proceed to next most tightly connected cluster

```

Figure 7.8: Server placer algorithm.

While growing clusters, and looking at connectedness between cells, the placer ignores connections from power or ground rails, or from nets with 50 or more sinks. Power and ground nets are easy to squeeze in after the rest of the placement is satisfactory, and large nets tend to be clock or reset lines, neither of which requires special consideration.¹⁶ A key to growing the clusters is to understand how connected they are with respect to each other, and the placer works through that question with a two-dimensional connection matrix. The connection weight between any two clusters is the sum of all connections between them.

¹⁴The critical path of a circuit is the slowest path—the one responsible for limiting the maximum operating frequency of the entire circuit.

¹⁵Special cell types are much less plentiful than slices. This not only makes it easier for the placer to pick locations for them, but also allows them to passively influence the geometries of the clusters.

¹⁶To be fair, the importance of clock and reset signals is critical, but clocks have their own dedicated wiring available, and reset signals tend to connect to *everything*, making it too easy to end up with a single enormous cluster, rather than many smaller ones.

Once the clusters are formed, the placer looks at each one to determine its resource requirements, and any geometric constraints that may apply. The constraints might include minimum height or width dimensions, or might require specific relative positioning, as in the case of F7 mux structures. The placer identifies candidate aspect ratios for each cluster, and selects one based on estimated total wire length to its peers.

The algorithm makes a number of simplifying assumptions, to keep the problem size within bounds. One idea that worked nicely on paper, but not nearly as well in practice, was to look at a set of connected elements, begin with the two most tightly connected ones, look for the next most tightly connected element, and set it to the left or the right of the first two, continuing until all elements had been laid out in linear fashion. The placer would then take the linear arrangement, and fit it into a two-dimensional block, allowing it to fold back upon itself at will. The premise was that the best possible linear arrangement could only get better if it were allowed to fold, so as to bring many of the elements closer together. This approach was used for both the placement of clusters within the device, and for the placement of cells within each cluster. In practice however, although the elements may have ended up in close proximity to each other, the routes that connected them varied with congestion, and sometimes ended up being much longer than expected.

The floorplanning of each circuit depends in part upon any anchor points that it may have: constrained resources, or top-level connections to other circuits. Exceptions are made in the case of BUFGMUX global clock buffers, since their wires will always follow the clock tree, and since these buffers are equidistant from the clocked elements that they drive.

This placer is at least somewhat greedy in its behavior, and while that makes it good at finding compact placements, that also increases the likelihood that the resulting placement will not be routable. That problem is addressed by forcing extra padding into the clusters in proportion to their sizes. Furthermore, the largest and least flexible blocks—such as carry-chains—are placed first, leaving it to the more flexible components to fit into the remaining spaces. What matters most is that the algorithm works, produces good quality results, and is much faster than the annealing implementation that it replaces.

7.4.8 Mapper

Components in the Xilinx Unified Libraries that pertain to slices actually describe functionality in terms of slice subcells rather than full slices. The placer is able to work with these subcells, but the Device API and the remainder of the system are not, and the slice subcells must therefore be packed back into regular slice cells before the routing and configuration can be changed. *This packing stage applies only to slices, and not to any other device cell type.*

The packing is complicated by the fact that slice subcells are not entirely independent, and furthermore that there is a considerable amount of overlap in the wires that they use. To guarantee proper operation, the packing is constrained by a series of rules. A subset of these rules is enforced

during the placement stage, to keep the placement from generating invalid results. The full set of rules is then applied to determine the final configuration of each affected slice. The rules checked during placement are as follows:

1. The slice may not contain a mix of synchronous and asynchronous flip-flops/latches.
2. The slice may not contain a mix of flip-flops and latches.
3. The slice may not contain a mix of positive-edge and negative-edge triggered flip-flops/latches.
4. The flip-flops/latches may not be driven by different clocks.
5. The flip-flops/latches may not be driven by different resets.
6. The flip-flops/latches may not be driven by different chip-enables.
7. The X output may be driven by only one of XORF, FLUT, or F5MUX.
8. The BX input may drive only one of F5MUX, CYMUXF or the carry chain input. *Conservative*.
9. The Y output may be driven by only one of XORG, GLUT, F6MUX, or ORCY.
10. The BY input may drive only one of F6MUX, CYMUXG or FFY.
11. The X flip-flop/latch cannot be driven by BX if BX is already in use by another net.
12. The Y flip-flop/latch cannot be driven by BY if BY is already in use by another net.

Note that rules marked as *conservative* are slightly more restrictive than necessary, to avoid performing complex checks during placement. These restrictive rules ensure that the post-placement packing will be able to generate a valid configuration. The only penalty for these conservative rules is that the resulting placement may require a slightly larger area because it prevents the placer from allocating competing subcells in the same slice.

When the placement has been finalized, the full configuration of each slice is determined, and the full packing takes place. It is worth noting that the rules presented here are simpler than they would have to be if the LUT RAM and shift-register functionality were supported.

One of the things that strongly affects configuration decisions is whether subcells connect to local wires or general wires or both. General wires tie in to the general routing matrix, and can reach most any point in the device. Local wires are dedicated wires that only connect one slice to another, and include the carry chain, the shift chain, the sum-of-products chain, and the F5/F6 mux lines. Local wires cannot connect to general wires without passing through slice logic.

Another kind of distinction exists between internal and external slice ports. External ports are those that show up on the periphery of the cell and are able to connect to other cells, whether through local or general wiring. Internal ports are not accessible from outside the slice, and originate when the external wiring is not fully defined. For example, if a top-level circuit has a port driven by a LUT output, but with no other connections specified, the mapper will not know where to direct that output, so the LUT output will be treated as an internal port. That internal port will not

be resolvable until the top-level circuit's port is connected to another port through a top-level connection.

The proper usage of local versus general wiring is suggested by `_L` and `_D` suffixes on Xilinx library primitive names, although the EDIF generated by Synplify Pro shows that the suggestion is occasionally misleading. When the suffix can be used on a primitive, absence of the suffix denotes general wiring, presence of the `_L` suffix denotes local wiring, and presence of the `_D` suffix denotes *dual* outputs with both local and general wiring. In some cases, support for dual outputs would require splitting entire nets into local and general subnets. That capability is not currently supported, and cases that would require it are disallowed.

The checks and decisions made during full packing are presented in Appendix A.

The packing stage relies on circuit net and cell information to determine how to configure slices, and it is therefore not possible to replace any of the nets or cells until the complete packing has been determined. Failure to defer the replacement leads to situations where some cells have been replaced while others have not, with the new ones no longer bearing any resemblance to the Xilinx Unified Library primitives from which they came.

7.4.9 Router

The router is thoroughly discussed in [22], but a number of changes and additions were made to it in support of the present work.

- The router can accept a timeout delay. If a timeout is set, any route that takes more than the specified amount of time fails. This sidesteps situations that would cause the A* algorithm to visit every remaining wire in the device, while trying to find a path to a destination that was no longer reachable. In cases where the net has multiple sinks, the timeout duration applies to each sink individually, and not to the entire net.
- The router can provide listeners with information about each arc traversed by a route. This is included in support of the server's XDL export capability.
- The router can route multi-sink nets in the order given, or reorder them from farthest to nearest or nearest to farthest. This capability tends to reduce overall congestion for large circuits.
- The heuristic can take advantage of local routing information for each tile type, to better evaluate distance and cost for the cost function.
- The heuristic can accept a list of reserved wires that it is not allowed to use. Although this could be used for fault tolerance purposes, it is primarily used to reserve circuit ports that may be needed in the future. In the unusual case of dynamically implemented circuitry, certain

circuit ports may be unused when the original routing is performed, but may be needed at a later time for other connections. There is no way to foresee this possibility, so the only safe option is to reserve all circuit inputs and outputs.¹⁷

- The heuristic follows target sink wires back to their first connecting point. This addresses situations in which a specified target wire is never driven from inside the tile that it resides in. Much of the routing proceeds by looking at the distance from the target tile, and if no connections to the target wire can actually be made from inside that tile, it makes little sense to keep heading towards it. An example would be almost any of the RAMB16 inputs, that reside in BRAM site tiles, but that can only be driven from BRAM interface tiles.
- The heuristic can accept a wire filter object that accepts or rejects any wire visited. This could be effectively used to constrain the routing inside or outside of some arbitrary region, geometric or logical. It could also prevent the router from using certain kinds of wires, or from using defective wires.

The reader is reminded that the router does not have access to or knowledge of timing information, and that it does not support passthroughs.¹⁸

7.4.10 Configuration

Every logic cell that is used in a circuit must be configured, unless it happens to require only default settings. In each case, the primitive associated with the user cell is responsible for remembering and applying the configuration settings.

Some configuration settings are generated or passed along in the EDIF by the synthesizer—look-up table INIT settings, for example—while others can be directly inferred from the primitive name—RAMB16_S36_S36 configures both of its ports in 512×36 mode. Still others must be determined in more complicated fashion by the primitive, as is the case for most slice subcells.

Regardless of where the settings are initially determined, they are added to each cell’s property hash map by name. Once a user cell has been placed in one or more device cells, the primitive can iterate over the named properties and set the device cell’s resource settings accordingly. The inner workings depend on whether the resource type is a named value, a bit vector, or a string.

The Device API tracks the resource settings for all instantiated cells. Whenever a resource is changed, the Device API flags that resource and any dependents for evaluation before the configuration bitstream is generated. This process includes a number of complexities, but the Device API handles everything transparently from the perspective of the server.

¹⁷One specific problem that this change solves involves the Virtex2 BX and BY slice inputs. These wires are in fact bidirectional, and can effectively be used as jump off points to reach other inputs. If a different net later needs to use one such BX or BY input, the route will fail immediately because the input wire is no longer available.

¹⁸A passthrough is a portion of a net that passes through logic resources, using them as impromptu wiring resources.

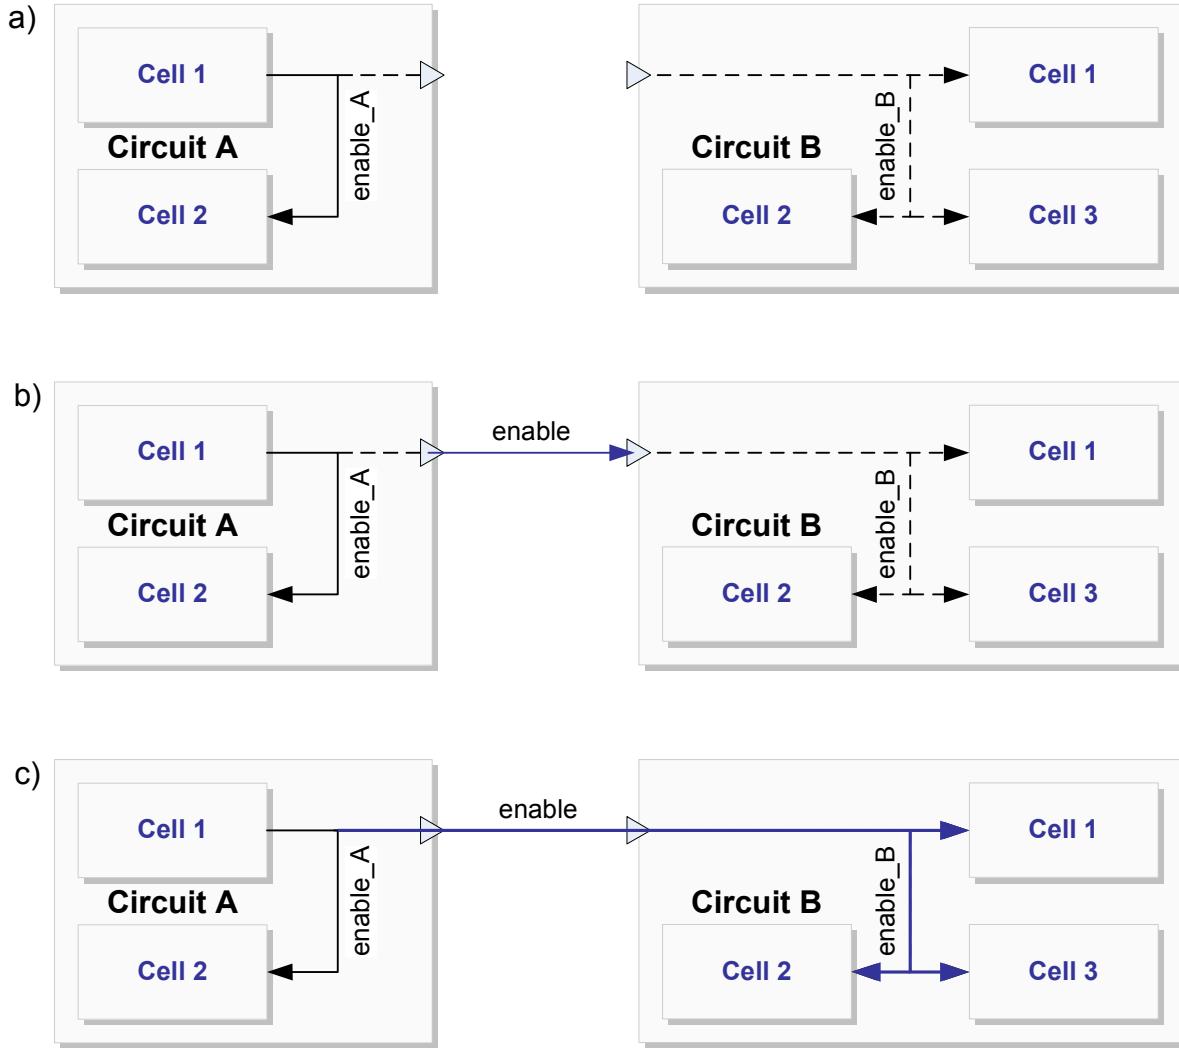


Figure 7.9: Top-level net: a) Two circuits, A and B, both contain enable nets. The *enable_A* net is partly routed, while the *enable_B* net is fully unrouted. b) Top-level net *enable* is defined, connecting from one top-level source to one top-level sink. c) Top-level net *enable* is routed, extending *enable_A* to three sinks on *enable_B*.

7.4.11 Connections

Although the routing stage makes all necessary connections *inside* each circuit, it is still necessary to establish top-level connections *between* circuits. This capability allows the system to dynamically and arbitrarily add, remove, or change top-level connections.

In a circuit that has been placed, routing a regular net involves looking up the device cell pins that correspond to each net port, and passing those pins to the router. It is not possible to do this before the circuit has been placed. Conversely, unrouting a net requires removing all connections, from the sink pins all the way back to the source pin.

Top-level connections are quite different. Firstly, top-level connections can be specified even before a circuit has been placed—though not before it has been parsed—and top-level connections may in fact influence the placement. Secondly, the ports in top-level nets often describe multiple pins, rather than a single pin, even when the net is a not a bus. And thirdly, top-level nets sometimes have to connect to user cell pins that may have remained *internal* ports inside a slice, because the mapper could find no connection to them. In such cases, the server will ask the owning primitive to try to pull the internal port out to an external device cell pin.

Figure 7.9 shows a top-level net named *enable*, connecting Circuit A net *enable_A*, and Circuit B net *enable_B*. In Step a) we discover *enable_A* partly routed, while *enable_B* is fully unrouted. In Step b) we create a top-level connection from a single port in Circuit A to a single port in Circuit B. But it is apparent in Step c) that to actually route the new top-level net, it is necessary to extend *enable_A* to three cell pins in *enable_B*.

Top-level busses are more complex only to the extent that the system has to create top-level connections for each wire in the bus. But it is worth noting that there are no shortcuts—the system must intelligently create each connection.

Referring back to the example in Figure 7.9, if the system needed to remove top-level net *enable*, it would intelligently trim each of the branches to Cell 1, Cell 2, and Cell 3 in Circuit B, resulting in a circuit that once again looked like Step a). Again, there are no shortcuts when disconnecting busses—each wire in the bus may have very dissimilar connections.

7.4.12 XDL Export

An XDL export function was added to the server, to facilitate timing analysis, validation, and debugging. XDL is a Xilinx file format that can be converted to and from NCD, which means that it can fully describe the configuration of the device. XDL for each instantiated circuit is exported after the circuit is routed and configured. And XDL for the top-level connections is exported on-demand through one of the protocol commands.

This export capability required some extensions previously discussed for the router, because the *xdl* conversion tool requires that each arc in every net be fully specified.¹⁹

To obtain meaningful results, the instantiated circuits, along with the base system, and all of the top-level connections, had to be merged into a single *valid* XDL file—a painfully complicated

¹⁹In the very common case where an arc between two wire segments can be expressed in a number of ways, it was also necessary to infer the conventions that the *xdl* tool expected to see.

process, performed by a Perl script, that reprised the net merging performed for top-level nets. However, that effort made it possible to generate an NCD file that captured the entire configuration of the dynamic system. In the XDL to NCD conversion process, the *xdl* tool performs full Design Rule Checking (DRC)—a very useful debugging aid. Once the NCD is created, *timingan* can perform full timing analysis on the design, and return true f_{\max} data. And FPGA Editor can open the NCD and display all of the nets, cells, and configuration for the combined static and dynamic system. These external verification capabilities inspired a large measure of confidence in the validity of the system and its operation.

It is important to note, however, that the XDL exporting truly is an export only. There is no way to import XDL data back into the system. And although it is possible to perform *offline* debugging and timing analysis, none of the timing data can be fed back into the system—the router still has no knowledge of timing.

Chapter 8

Operation

8.1 Overview

This chapter considers the operation of the system. It looks at system and circuit preparation, at server, controller, and operating system startup, and at normal operation and behavior. The chapter concludes by discussing performance, external observability and analysis, and the overheads involved in supporting autonomy.

Figure 8.1 depicts the system startup and operation. When the system is powered on, it begins by configuring the FPGA, booting the Linux kernel, installing custom drivers, and invoking *init*, the parent of all other processes. The autonomous server and the autonomous controller can then be launched, with each one attending to its own initialization.

When the server completes its initialization, it opens a server socket, and waits for connection requests. In most cases, the controller will have been waiting for the server to open a socket, and when that socket becomes available, a connection is established right away. Although the server and controller both run on the demonstration system, and communicate through a TCP socket, it is also possible for a remote interactive client to connect instead from anywhere on the network.

Depicted in dark blue in Figure 8.1, are events that may arise during the course of operation. At present, these originate outside the system, on inputs that it observes, but it is entirely possible that other events should arise internally. Dynamic responses are shown in red: If a suitable response is identified, the controller communicates with the server to implement the response and reconfigure itself.

The actual demonstration allows the user to interactively mask out resources, and allows the system to respond to detected conditions by dynamically instantiating or removing circuitry. The implications of this will be evaluated in terms of the Chapter 4 roadmap.

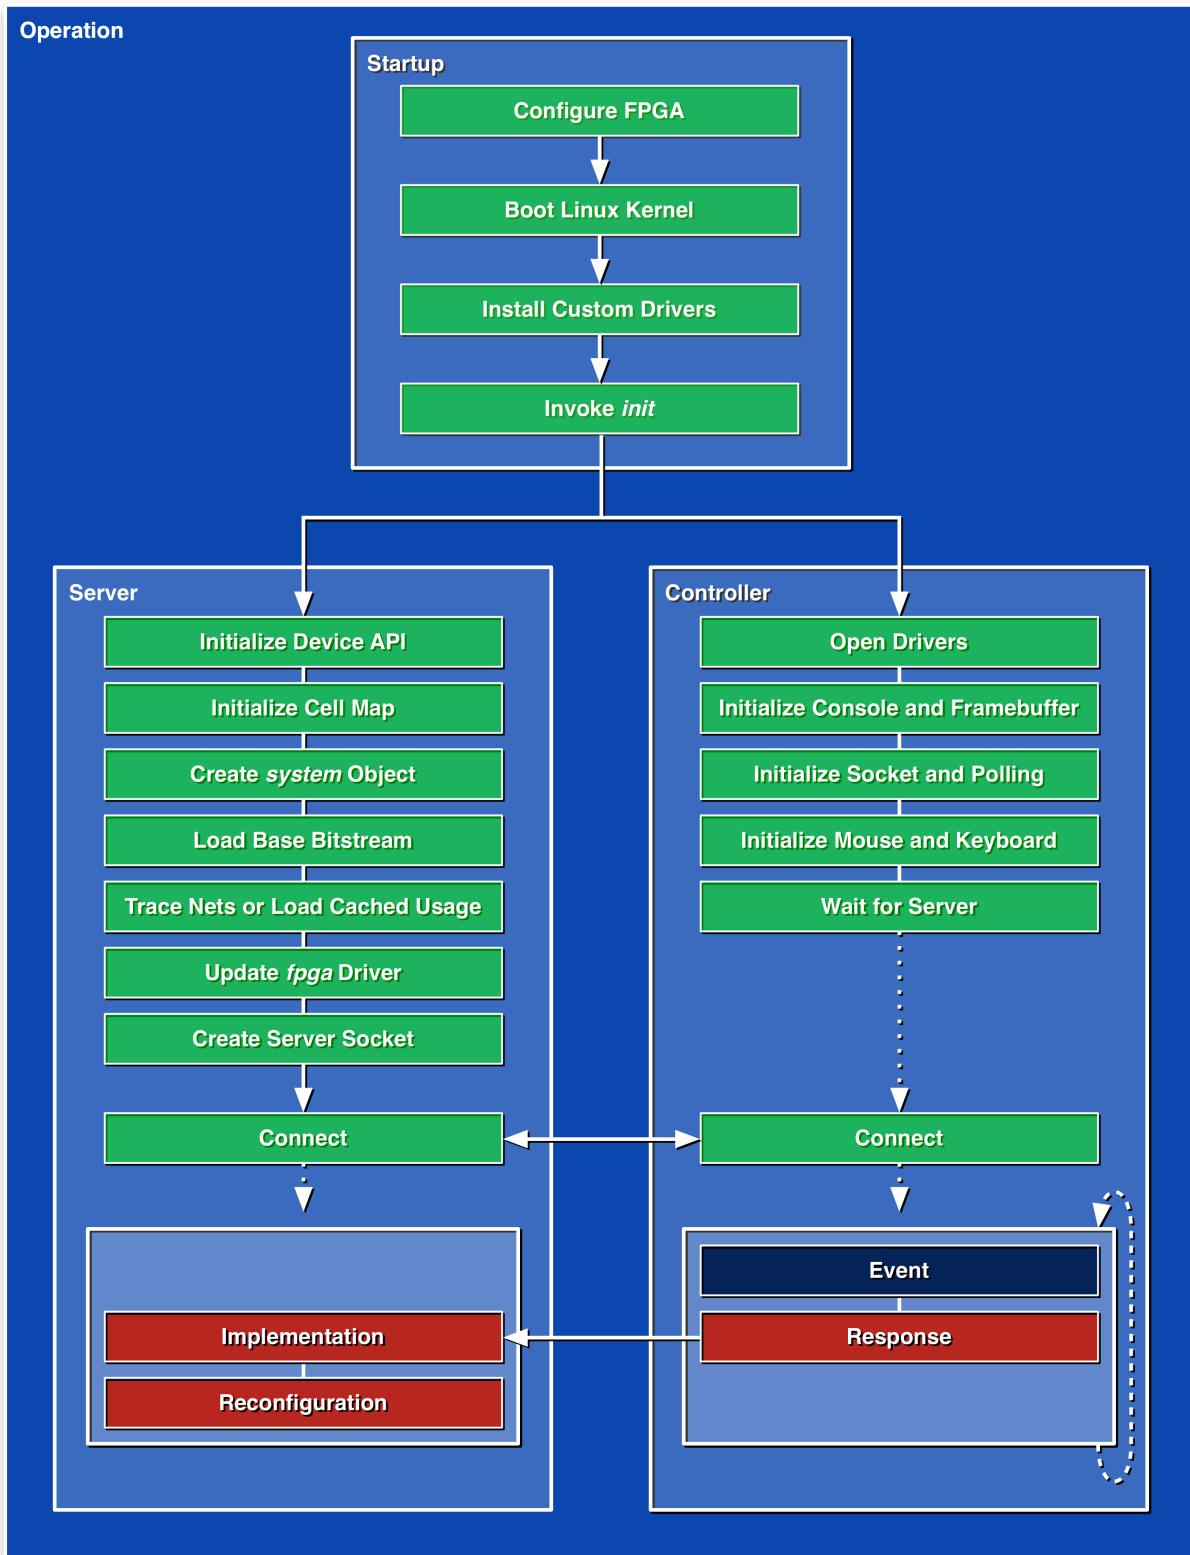


Figure 8.1: Startup and operation.

The second half of the chapter is devoted to analyzing the system: Its dynamically instantiated circuits can be exported as XDL, to be observed and quantified with external tools. And its static base system can be measured in terms of the overhead that it imposes upon an embedded system. Results are presented and conclusions are drawn, but one particular metric is especially important, and makes this the first known demonstration of hardware autonomy: The system actually works.

8.2 Preparation

The server and controller software for the demonstration system are available through its filesystem. If the XUP board is connected to its normal network, it mounts an *nfs* share at the appropriate mount point, and otherwise, it simply uses its local copy of that same data from the MicroDrive. For a system that purports to be autonomous, it is useful to be able to show it disconnected from other computers and thus from external control. And yet the best known autonomous systems—*living systems*—are often highly connected, without feeling insecure in their autonomy: external connectedness does not imply external control.

An autonomous system must either begin its existence with a fully populated response library, or must be able to augment that library along the way. The current demonstration system has no ability to expand its library—though nothing precludes that in the future—so it must be populated ahead of time.

Preparation of the dynamically implemented cores is fairly straightforward at present, though a small number of limitations and conventions apply. On the limitation side, it has been made clear that TBUFs, LUT RAMs, and shift registers are unsupported, and that dual-ported RAMB16s are subject to possible state corruption. None of these are inherent limitations, but are instead the result of Virtex2P architecture properties, coupled with some design trade-offs. If it seemed acceptable to suspend circuits while overlapping frames were reconfigured, the corruption issues could be sidestepped.

As for conventions that apply to the dynamic cores, the cores are presently synthesized without input/output pad insertion,¹ and are output as flat² EDIF netlists, because the absence of hierarchy simplifies things slightly for the EDIF parser. All of dynamic cores were synthesized with Synplify Pro [87].

¹Synthesizers routinely tie top-level circuit ports to input/output pads, which is appropriate for circuits that will be implemented as top-level designs in the target FPGA. Note that in a small number of cases, the designer may actually *want* pad insertions with this system.

²Flattening a netlist removes all hierarchical information from it.

8.3 Startup

Starting up the demonstration system is a surprisingly complex process, even though it can happen transparently. Many pieces of the following narrative have been mentioned along the way, but they are repeated here in context for the sake of coherence.

8.3.1 Linux

The MicroDrive that provides the root filesystem also includes a dedicated FAT16 partition, suitable for storing a SystemACE file to be read by the SystemACE controller upon startup. The SystemACE file contains the compressed Linux image, bundled with the base configuration bitstream.

When the system is turned on, the SystemACE controller looks at the first partition on the MicroDrive, finds the SystemACE file, configures the FPGA with the embedded configuration bitstream, loads the compressed Linux image into memory, and invokes the bootloader.

The Linux bootloader inspects its memory space, uncompresses the Linux image, and begins booting the kernel. Once the kernel and its built-in drivers have finished loading, the kernel invokes *init*, which serves as the parent of all other processes.

8.3.2 Drivers

Once Linux is fully loaded, the custom drivers can be loaded as well. All but a few of the drivers described in Section 6.3 are loaded into the kernel:

gpio (/dev/gpio/*): Interface to the XUP board LEDs, pushbuttons, and DIP switches. Used for debugging purposes as a system heartbeat.

fb (/dev/fb): Framebuffer driver. Allows the autonomous controller and the fpga driver to draw on the screen and control hardware options.

fpga (/dev/fpga): FPGA resource usage visualizer. Allows the autonomous controller and the server to exchange usage data, and to display it on the screen for user feedback.

icap (/dev/icap): Interface to the FPGA’s Internal Configuration Access Port. Allows the server to reconfigure the FPGA as directed by the autonomous controller.

ps2 (/dev/ps2/*): PS/2 mouse and keyboard driver. Allows the user to interact with the autonomous controller.

stub (/dev/stub): Interface for all sensor stubs and subscribers. Provides notification to the autonomous controller when mouse, keyboard, audio, or video signals are acquired or lost.

8.3.3 Server

When the server is launched, it begins by initializing the Device API databases for the Virtex2P XC2VP30-FF896-6. This loads all device logic, wiring, and packaging information into memory, and makes the information accessible to the server through an API.

Once the Device API is initialized, the server proceeds to initialize its internal cell map. The cell map actually consists of three different data structures described earlier, aimed at improving placer performance by giving it very fast indexed access to cells.

The system circuit object is created next. It exposes all of the device's IOBs and global clock buffer ports, as well as the interface stub ports. This step makes it possible for dynamic circuits to connect to all system IOBs,³ all global clock buffers, and all logical interface points into the remainder of the system.

The server then loads the current configuration bitstream into memory. If the resource usage information was previously cached, the server reads the contents of the cache file to populate its internal resource usage lists. Otherwise, it traces every net in the bitstream to determine which wires, connections, and cells are being used by the base system. The reader may note that reading cached usage information is sensible enough, since the base configuration bitstream changes only infrequently, and always under direct human control. Whenever the base configuration bitstream is modified, the resource usage cache has to be deleted to force a full trace of the new bitstream. In cases where the Device API does have to trace through the full bitstream, the startup time increases accordingly.

In a small number of cases, there is some disagreement between the server and the ISE tools over what device cells are in use. The tracer conservatively estimates that if there is a connection from a cell output to a cell input, then both of the cells involved must be in use. However, the ISE router sometimes allows nets to bounce off cell input wires, without actually affecting the cell logic. In such cases, the server would incorrectly flag the cells as used, even though the ISE tools would show them as unused. This results in a reduced number of usable cells, but causes no other harm, particularly since the placer confines its slice allocation to the large region reserved for it.

The one other case where disagreements may arise between ISE and the server, involves output pads that are driven by a default VCC connection. Almost all cell inputs that are undriven have an implicit connection to VCC, and if nothing else about a cell differentiates it, the server will be unable to distinguish it from an unused cell. The specific instance in which this occurs involves two IOBs used by the Ethernet MAC, that drive VCC through two output pads. These two cases could result in an error if the server tried to reuse the IOB cells in question, so the appropriate correction is hardcoded into the server startup code.

When the server has updated its internal usage information, it sends that information to the fpga driver, which draws the base resource usage on the screen, and makes it available for the

³This includes the ability to tie into existing connections—the IOB need not be unused.

autonomous controller to query. The server then creates a TCP server socket and listens for incoming connections.

8.3.4 Controller

The autonomous controller can be launched in parallel with the server, once the custom drivers have been loaded. Because its initialization is far simpler than the server's, it simply waits for the server to finish initializing if need be.

As previously described in Section 7.3, the controller begins by opening a number of devices:

- /dev/ps2/mouse:** Mouse input. Allows the user to interactively reserve or release logic resources.
- /dev/ps2/keyboard:** Keyboard input. Allows the user to interact with the server, and obtain system information.
- /dev/fb:** SXGA monitor output. Allows the controller to draw its Graphical User Interface (GUI), to display mouse activity, and to display keyboard input and server responses.
- /dev/fpga:** Resource usage visualization. Allows the controller to display reserved resources, and to obtain usage information staged by the server.
- /dev/stub:** Sensor stub notifications. Provides notification to the controller of any acquisition or loss of signal activity on the monitored ports: mouse, keyboard, video decoder input, or audio codec input.

The controller then initializes the console—through which the server and the user interact—as well as the framebuffer, the TCP socket, the client interface that implements the server protocol, and the device polling. It initializes the mouse and keyboard if applicable, and connects to the server, or waits until the server is ready to establish a connection. The controller then reads the base resource usage that the server has provided, and enters its main event loop.

8.4 Operation

Once the system is up and running, most of the complexities are either already taken care of, or entirely transparent to the user and to the autonomous controller.

8.4.1 Interaction

The user can interact with the system by connecting or disconnecting devices from the PS/2 mouse port, the PS/2 keyboard port, the stereo audio line input, the composite video input, the component

video input, or the S-video input.⁴ When a mouse is connected, the user can click on resources in the FPGA view, to mask or unmask any resources of their choosing. When a keyboard is connected, the user can interactively send commands to the server.

8.4.2 Events

Events that do not change the system's inputs are handled statically. Network data is sent to the console and the client interface for display and processing. And mouse and keyboard events are likewise processed for masking or typing purposes. It is certainly plausible that such events—or even software conditions—might require dynamic changes to the system, but implementing that in the present situation would have been somewhat contrived. The reader is reminded of this work's philosophy that change should happen for a reason, rather than happening randomly or arbitrarily.

Events that *do* change the system's inputs are handled differently, however. Any device that becomes connected may require simple configuration of the device. That is true for both the mouse and the keyboard, so any activity detection on those ports begins by configuring their rates and settings. The video decoder and the audio codec will already have been configured by the sensor stubs, before the controller is informed of the new activity.

Aside from any device configuration, the controller presently invokes scripts when devices are connected or disconnected. The scripts are simply named after the device, with an `_on` or `_off` suffix appended. For example, when the mouse is connected, the controller loads and runs the script named `mouse_on`. The *response library* of Chapter 4 is therefore simplistic in this implementation, because each condition that arises maps to exactly one response. It is important to note that more sophisticated response libraries can be added in the future, and Level 7 or Level 8 systems can be layered on top of the infrastructure that the current work provides.

8.4.3 Robustness

The quickest way to disrupt the system is to give it an outdated copy of the configuration bitstream. Firstly, if the resource usage of the base system is not cached, the server and Device API will infer it by tracing through the bitstream's wiring, as noted earlier. That will give the server an incorrect understanding of which wires and logic cells are in use, and as it goes about modifying itself, it runs the risk of coopting resources that should not be available to it.

Secondly, when the system reconfigures itself, it will be writing incorrect configuration frame data. Normally one might expect to perform a read-modify-write operation on the device: reading the current configuration data, applying new settings to it, and writing the modified configuration

⁴The video stub only monitors one of its inputs at a time, so if a particular input is active, connecting or disconnecting the other inputs has no effect whatsoever.

data back to the device’s configuration controller. But the Device API instead loads the base configuration bitstream when it starts up, and subsequently applies changes to both its internal model, and to the device itself. As long as both the model and the device remain synchronized with each other, the system can do little to harm itself. But if the model and the device differ, the system will corrupt itself as soon as it attempts to reconfigure. In almost all cases, that corruption is fatal.

Fortunately, this situation can only arise if the configuration bitstream with which the SystemACE controller configures the device, differs from the configuration bitstream that the server reads in when it starts up. As long as the two are synchronized when the system is first fielded, this problem will not arise. This problem could be sidestepped entirely if the server were able to read the base configuration bitstream from the SystemACE file on the reserved FAT16 partition, or if the icap driver supported readback and the server read the base configuration bitstream directly from the device.

There are many other things that can go wrong during operation. If a circuit includes embedded location constraints, but those locations are already in use, the system can do little more than report the condition and fail. The same is true if the system runs out of logic or wiring resources. These cases and others would have to be addressed before any system could be deemed robust. But the present work is exploratory, and is thus not obligated to perform at the level of a mature commercial system.

8.4.4 Time-Limited IP

Under normal conditions, an autonomous system should never have to shut down. However the base hardware in this system uses a few Xilinx Intellectual Property (IP) cores running in trial mode. The opb_uart16550_v1_00_c and the plb_ethernet_v1_00_a cores are available for purchase from Xilinx, but it is also possible to register to use these cores in trial mode at no charge.

It is also worth noting that OpenCores.org has Ethernet MAC 10/100 Mbps and UART 16550 cores that could have served as replacements for the plb_ethernet_v1_00_a and the opb_uart16550_v1_00_c. Experience shows that the OpenCores code is generally of good quality, although Xilinx code frequently provides additional implementation options, and is thought to be better mapped and optimized.

8.5 Validation

This section describes the capabilities that qualify the system for each of the claimed levels of autonomy on the roadmap, and describes the actual demonstrated system capabilities.

8.5.1 Level 0: Instantiating a System

The FPGA is configured through a configuration port by the SystemACE controller at startup. It is also possible to log on to the demonstration platform, and write partial configuration bitstreams directly to the ICAP through the /dev/icap driver. *It is imperative that the partial configuration bitstreams be compatible with the base configuration.* The availability of a configuration port establishes Level 0 capability.

8.5.2 Level 1: Instantiating Blocks

Level 1 capability is superseded by Level 2 capability. The allocation map and primitive router of Level 1 are encompassed by the detailed model and full router of a Level 2 system.

8.5.3 Level 2: Implementing Circuits

A client application can run locally or remotely to interact with the server, using the protocol described in Section 7.2.1. Through the client application, the user can tell the server to parse EDIF circuits, to place, route, configure, and connect them, and to reconfigure itself. The implicit internal model, along with the placer and router, establish Level 2 capability.

8.5.4 Level 3: Implementing Behavior

Level 3 capability is not implemented. The system is unable to synthesize and map behavioral functionality inside itself. This is deferred for future work.

8.5.5 Level 4: Observing Conditions

The autonomous controller receives notification of events from the sensor stub driver. Without requiring participation from the server, the controller displays input change notices in its console area, whenever audio, video, mouse, or keyboard signals are detected or lost. The sensors and the monitor establish Level 4 capability.

8.5.6 Level 5: Considering Responses

For a given notification from the sensor stub driver, the controller references a corresponding script. The mapping between the type of condition and the script is currently hardcoded, so the response library is fixed and the Level 5 capability is shallow at best.

8.5.7 Level 6: Implementing Responses

The autonomous controller responds to detected conditions with the help of the autonomous server. Server scripts selected by the controller are used to dynamically implement or remove circuitry from the system. The implementer establishes Level 6 capability.

A Level 6 system is also supposed to include some means to request assistance or provide notification if a response could not be identified or cannot be implemented. The present system has a response for each supported condition that may occur—a corresponding script—and therefore neither requests assistance nor provides failure notification, although those capabilities would be interesting to add in the future.

8.5.8 Actual Demonstration

The actual demonstration proceeds through the startup process described earlier, and includes launching the autonomous server and autonomous controller. Of the possible inputs that the system observes, only the mouse is initially connected. When the controller and server have established a connection, and have performed their handshaking with the fpga driver to display and exchange the base system's resource usage, the user can use the mouse to mask out resources in some area that the placer is likely to use. The keyboard can also be plugged in—being immediately detected and configured by the controller—to interactively query the autonomous server or implement functionality.

A video signal is then connected to one of the video decoder's inputs, and the controller receives notification of the event through the sensor stub driver. The controller reads the `video_on` script, and sends its commands to the server, providing underlying EDIF files where appropriate. The server parses the various EDIF files, places, routes, configures, and connects the resulting circuits, and reconfigures itself as directed by the script. The placement naturally avoids any resources that the user masked out, and dynamically updates the displayed resource usage as it goes. When the implementation is complete, the full-resolution live video shows up on the display, next to the resource usage, as shown in Figure 7.2. The user can disconnect the video, or the mouse or keyboard, and the controller will continue to process the events as appropriate.

8.6 Analysis

The demonstration system is still subject to tuning and continued work, so the quantitative analysis presented in this section is more of a representative snapshot than a final pronouncement. Even under dynamic conditions, the system tools may need to adjust their parameters according to resource availability and congestion. Placements that are tightly packed may conversely be more

difficult to route, and without the benefit of timing-driven and sophisticated routing, the maximum frequency of the resulting circuits sometimes improves when the distances are not minimal.

8.6.1 XDL Export

The server's XDL export capability, described in Section 7.4.12, affords this work the benefit of full Design Rule Checks (DRC) and timing analysis, along with observability of the *complete* system—a snapshot of the base hardware and dynamic circuits combined.⁵

As previously stated, the XDL export capability is an export only: there is no way to import XDL into the system. Furthermore, although the analysis information is useful for offline metrics and debugging, it cannot be fed back into the system. In particular, the fact that timing analysis can be performed on the full system, does not give the router access to timing data. The router remains blind in that respect.

8.6.2 Performance

Performance numbers for dynamically instantiated circuits are generally most useful when they consider the circuit in the context of the full system. The actual connections naturally add to the routing delays, as do the stubs on both the system and circuit ends, and these are not fully characterized by themselves. Conversely, it isn't possible to directly compare the performance of dynamic circuits between the Xilinx ISE tools and the system presented here, without statically embedding the dynamic circuits within the base system. This section therefore presents some head-to-head results that show the relative performance of the tools, as well as some circuit-specific results inside the base system. Subsequent results make use of the circuits presented in Table 8.1.⁶

Head-to-head comparisons between the ISE tools and the Autonomous Computing System (ACS) demonstration system are presented in Table 8.2. These numbers are taken directly from the Linux *time* utility, in the case of implementation times, or from *timingan*, in the case of f_{\max} .⁷

The f_{\max} result for the ISE tools exceeds that for the ACS system in each case. This is due mostly to the high-quality timing-driven routing performed by the ISE tools: The ISE router is believed to be based on the *PathFinder* algorithm [88], and therefore has a very good ability to resolve contending resource demands from circuit nets. Some additional insights can be gleaned from

⁵The XDL output must first be translated to NCD, the Xilinx Native Circuit Description format, with the *xdl* tool in -xdl2ncd mode. The NCD format is suitable for use by FPGA Editor and *timingan*, and DRC is automatically performed during the translation from XDL.

⁶The difference between the *video_core1* and *video_core2* circuits stems from the decision to interface the cores to dual-ported BRAMs instead of crossing clock domains and interfacing to the OPB or PLB busses.

⁷The reader may note that f_{\max} for the *dec* circuit looks improbably high on a -6 speed grade Virtex2P. Those numbers are simply being reported as provided by *timingan*, with no pronouncement concerning their validity.

Table 8.1: Dynamic test circuits

<i>Circuit</i>	<i>EDIF Nets</i>	<i>EDIF Cells</i>	<i>Description</i>
dec	10	9	4-bit decimal counter
counter32	157	155	32-bit binary counter
ps2_core	865	788	PS/2 keyboard/mouse OPB interface
video_core1	1832	1567	Video-capture (ITU.R BT 656 to PLB interface)
video_core2	1014	1000	Video-capture (ITU.R BT 656 to BRAM interface)
rfft	3981	3835	Stereo 256-point audio FFT core

Table 8.2: Dynamic circuit performance comparison

<i>Circuit</i>	Xilinx ISE			ACS System			
	<i>f_{max} (MHz)</i>	<i>Slices</i>	<i>ccm4 (s)</i>	<i>f_{max} (MHz)</i>	<i>Slices</i>	<i>ccm4 (s)</i>	<i>xup0 (s)</i>
dec	928	5	0:44.50	724	3	0:05.68	2:45.56
counter32	312	29	0:44.21	256	25	0:06.03	7:33.10
ps2_core	265	349	1:00.64	77	230	0:20.29	16:41.43
video_core1	213	466+7	1:08.14	102	470+7	0:13.56	14:18.91

inspecting corresponding routes from the two tools in FPGA Editor. Although detailed routing analysis is beyond the scope of this work, observation shows “cleaner” routes coming from the ISE tools than from the ACS system. Unfortunately, these routes are difficult to properly visualize without the benefit of a suitable CAD tool or a large plot, and are not reproduced here.

Table 8.2 also shows implementation times for both the ISE tools and the ACS system, on *ccm4*, a 2.4 GHz Quad Core2 Linux machine, with 4 MB of cache for each of the four processors, and 2 GB of main memory. These times show the ACS flow running 4–8 times faster than the ISE flow, albeit with lower quality results. But it remains that the ACS tools must run on *xup0*, a 300 MHz PowerPC 405, with 16+16 kB of cache, and 512 MB of main memory. The implementation times on *xup0* are 29–63 times slower than on *ccm4*, a telling measure of the disparity between workstations and embedded systems. To really bring the performance into focus, one should consider that an event that required instantiating the *video_core1* circuit would take nearly 15 minutes to complete the instantiation.

The discussion now shifts from ISE and ACS head-to-head comparisons to full-system context results. These results are obtained by taking the circuit XDL output and the base system XDL output, and letting a script carefully merge them with the XDL output for the top-level connections. This approach provides a usable XDL file that corresponds to the system’s configuration at the time the output was generated. After the merged XDL is generated, it is translated to NCD format,

at which time the full DRC is performed.⁸ The resulting NCD is then loaded into timingan, along with the base system's Physical Constraints File (PCF), to tell timingan which constraints it should verify. Table 8.3 shows the required and possible performance for two circuits of interest for the demonstration.

Table 8.3: Embedded dynamic circuit performance

<i>Circuit</i>	<i>f_{req} (MHz)</i>	<i>P_{req} (ns)</i>	<i>P_{act} (ns)</i>	<i>Slack (ns)</i>	<i>f_{act} (MHz)</i>
video_core2	27.000	37.037	9.200	27.808	109
rfft	12.288	81.380	20.603	60.777	48

The video_core2 circuit connects to the VDEC1 decoder's 27 MHz Line-Locked Clock, meaning that it must achieve a maximum period of 37.037 ns. The timingan output shows the worst delay to be 9.200 ns, corresponding to 27.808 ns of slack. In other words, including clock skew, and setup and hold time requirements, the implemented circuit could actually run at 109 MHz. The rfft circuit connects to the AC'97 codec's 12.288 MHz bit clock, meaning that it must achieve a maximum period of 81.380 ns. Its actual worst delay is 20.603 ns, which means that it could run at 48 MHz.

Although the video_core2 circuit can run as fast as 109 MHz, the rfft falls well short of that, which explains why earlier attempts to interface complex circuits directly to the 100 MHz OPB or PLB busses were scaled back. In some cases the circuits might actually run fast enough, but since there are no timing guarantees, metastability would be a serious concern, particularly since it would affect major system busses.

Additional data was collected to help determine whether the f_{\max} results could be improved most effectively with changes to the placer or with changes to the router. This was done by exporting circuit XDL data, and re-routing with the ISE *par* tool. To further understand the behavior, the tests were run again, this time with *par* re-placing and then re-routing. Finally, both of these steps were repeated, with a timing constraint that forced *par* to work very hard. The results are presented in Table 8.4.⁹

A row showing ACS and ACS in the Placer and Router columns, indicates that both placement and routing were performed by the ACS system. ACS and ISE indicates placement by the ACS system, and re-routing by ISE. And ISE and ISE indicates re-placement and re-routing by ISE. For the video_core2 circuit, *par* is able to increase circuit performance by 59.19 % without even really trying, and by 76.87 % when trying very hard. When *par* is directed to re-place the circuit, those numbers change to 160.92 % and 165.35 %, respectively. Clearly *par* is able to do a good job with the routing, even if the placement is less than ideal. But for this circuit, *par* doesn't gain much additional performance when also directed to re-place the circuit.

⁸A small number of warnings appear, but these have been analyzed and dismissed as materially insignificant.

⁹These results differ somewhat from those of Table 8.3, because of different tuning parameters in effect at the time of testing, but they still refer to the same circuits.

Table 8.4: Embedded dynamic circuit comparative analysis

Placer	Router	$P_{\text{req}} (\text{ns})$	$P_{\text{act}} (\text{ns})$	$f_{\text{act}} (\text{MHz})$	Speedup (%)
video_core2					
ACS	ACS	n/a	11.097	90.11	
ACS	ISE	37.037	6.971	143.45	59.19
ISE	ISE	37.037	4.253	90.11	160.92
ACS	ISE	4.000	6.274	159.39	76.87
ISE	ISE	4.000	4.182	239.12	165.35
rfft					
ACS	ACS	n/a	47.842	20.90	
ACS	ISE	81.380	43.296	23.10	10.50
ISE	ISE	81.380	42.996	23.26	11.27
ACS	ISE	10.000	16.525	60.51	189.51
ISE	ISE	10.000	12.006	83.29	298.48

The results are a little different for the FFT circuit, where par increases circuit performance by only 10.50 % when it isn't trying, and by 189.51 % when trying very hard. When par also re-places the circuit, the numbers change to 11.27 % and 298.48 %, respectively. In this case, par improves little over ACS when it doesn't need to, but improves spectacularly when forced to work hard. Even without re-placing the circuit, par can improve performance by nearly 200 % just by re-routing it. It seems then that the results are inconclusive as to whether ACS performance would improve most in response to router changes or to placer changes. It would probably be necessary to run similar tests on a broad range of circuits to gain a better understanding of the underlying behaviors.

8.6.3 Inspection

In addition to facilitating full timing analysis, the XDL export capability also simplified the debugging process, and provided insight into problems that the router encountered. For example, during experimentations with placer tuning, the router would sometimes run into too much congestion, and fail to route certain nets that appeared easily routable. The truth of course was that the circuits being routed were only responsible for part of the congestion, and the remainder of it could be attributed to the base system. The fact that FPGA Editor was not able to route these nets either—when taken in the context of the full system—lent a measure of external credibility to the router's troubles. The placement was further tuned to try to reduce congestion problems.

Visualization of placement and routing, with full net and cell names, can also be invaluable for development and debugging purposes. Figure 8.2 shows the XDL view of a slice configured by the system, while Figure 8.3 shows the FPGA Editor view of that same slice. And Figure 8.4 shows a sample of logic placed and routed by the system.

```

1 inst "video/SLICE_X72Y111" "SLICE", placed R25C37 SLICE_X72Y111, cfg "
2   CYOF::#OFF CYOG::#OFF
3   CYINIT::CIN
4   FFMUX::FXOR GYMUX::G
5   XBMUX::#OFF YBMUX::#OFF
6   CYSELF::#OFF CYSELG::#OFF
7   DXMUX::1 DYMUX::0
8   SOPEXTSEL::#OFF SYNC_ATTR::ASYNC
9   CLKINV::CLK SRINV::SR CEINV::CE
10  BXINV::#OFF BYINV::BY YBUSED::#OFF
11  FFX:YCrCb2RGB_inst_B_int_20:FF
12  FFY:YCrCb2RGB_inst_X_int_5:FF
13  SRFMUX::0 REVUSED::#OFF
14  F:YCrCb2RGB_inst_un4_B_int_axb_18:LUT:D=(A1*~A2)+(~A1*A2)
15  G:YCrCb2RGB_inst_N_175_i:LUT:D=(A1*~A4)+(A2*~A4)+(A3*~A4)
16  FFX_INIT_ATTR::INIT0 FFY_INIT_ATTR::INIT0
17  FFX_SR_ATTR::SRLOW FFY_SR_ATTR::SRLOW
18  FXUSED::#OFF F5USED::#OFF
19  XORF:YCrCb2RGB_inst_un4_B_int_s_18:
20  XUSED::0 YUSED::0
21  F_ATTR::#OFF G_ATTR::#OFF
22 ";

```

Figure 8.2: Configured slice sample (XDL view). Setting names and values correspond exactly to FPGA Editor conventions. EDIF cell names embedded where appropriate.

8.6.4 Overhead

It seems appropriate to revisit the drawbacks attributed to autonomous computing systems in Chapter 1, and to quantify them at least in part. The items mentioned were area and performance penalties, the need for commercial-grade incremental and distributed versions of the tools, the added complexity of testing, and the paradigm shift for engineers.

Concerning the area penalty associated with an autonomous system, the base system that is necessary to support autonomy uses 7,426 or 54 % of 13,696 total slices, along with 613 other logic cells. But the potential state corruption issues restricted the floorplan a little further, such that no dynamically instantiated slices were permitted in the same columns as static logic. This requirement is more relaxed with newer architectures, in which configuration frames no longer span the entire height of the device. The penalty is further amortized in the case of larger devices, or in cases where a single controller in a larger system may assume responsibility for multiple configurable devices.

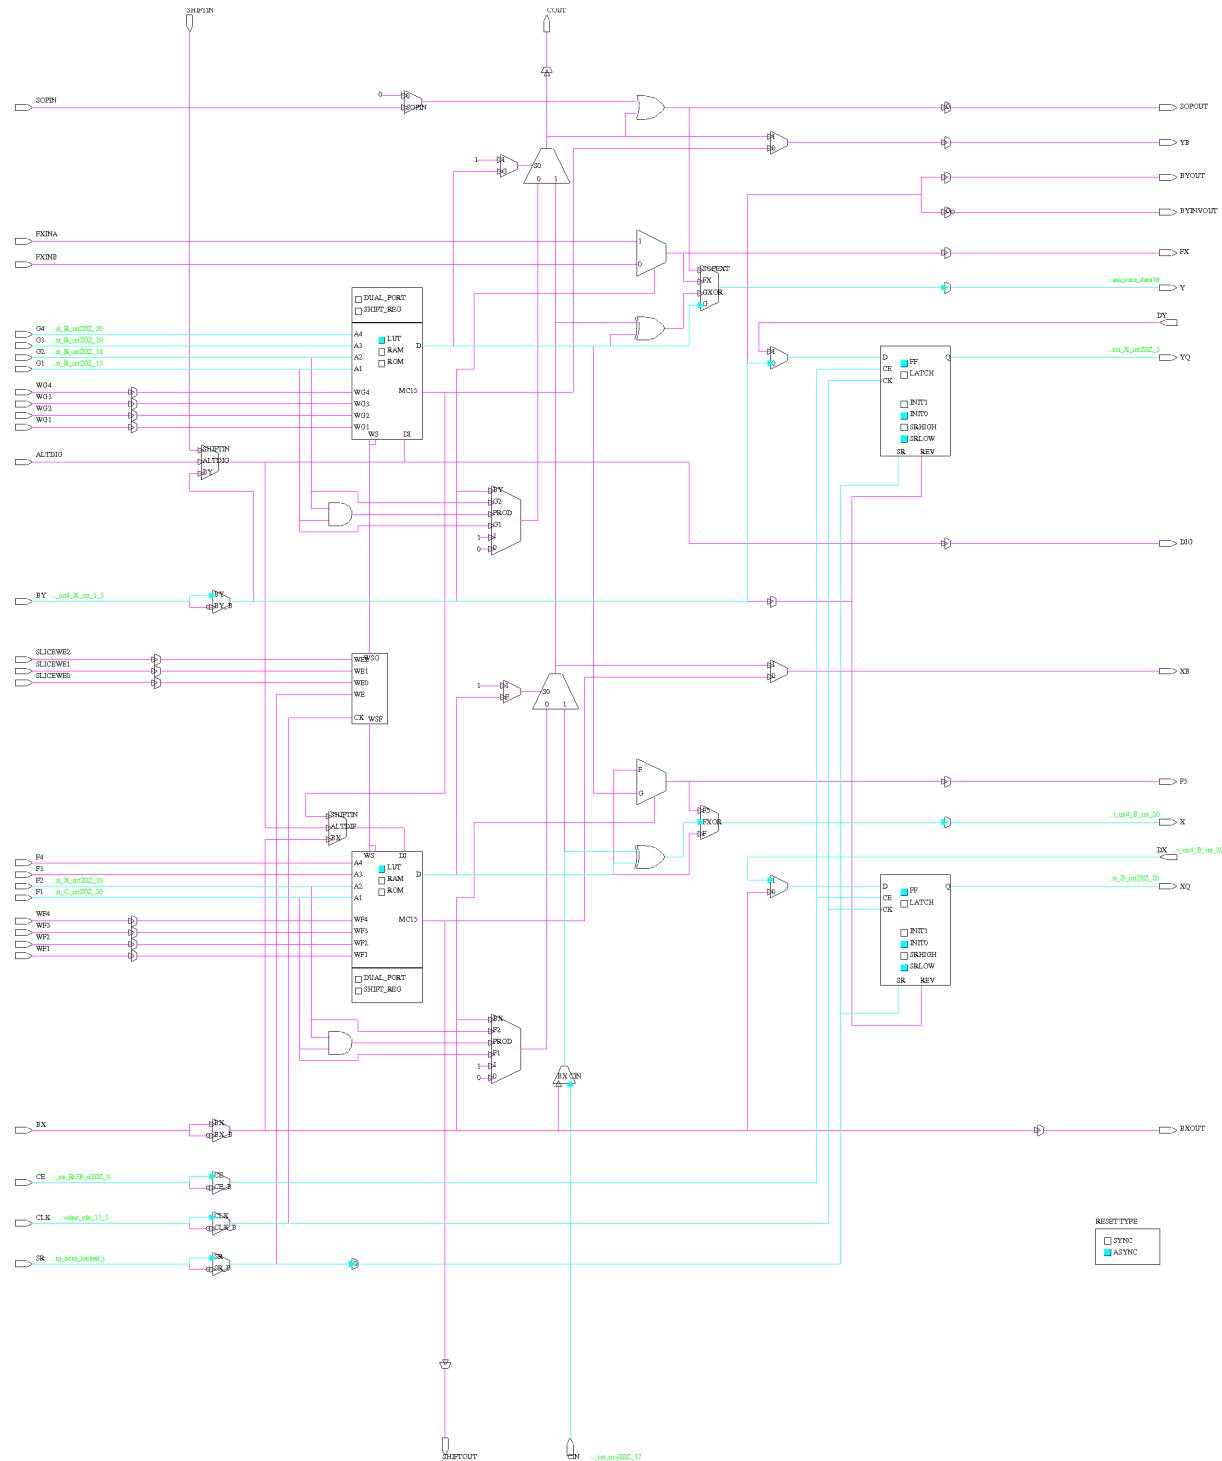


Figure 8.3: Configured slice sample. Exported as XDL, translated to NCD, and opened with FPGA Editor. Some EDIF cell names are only visible when the relevant resource is clicked.

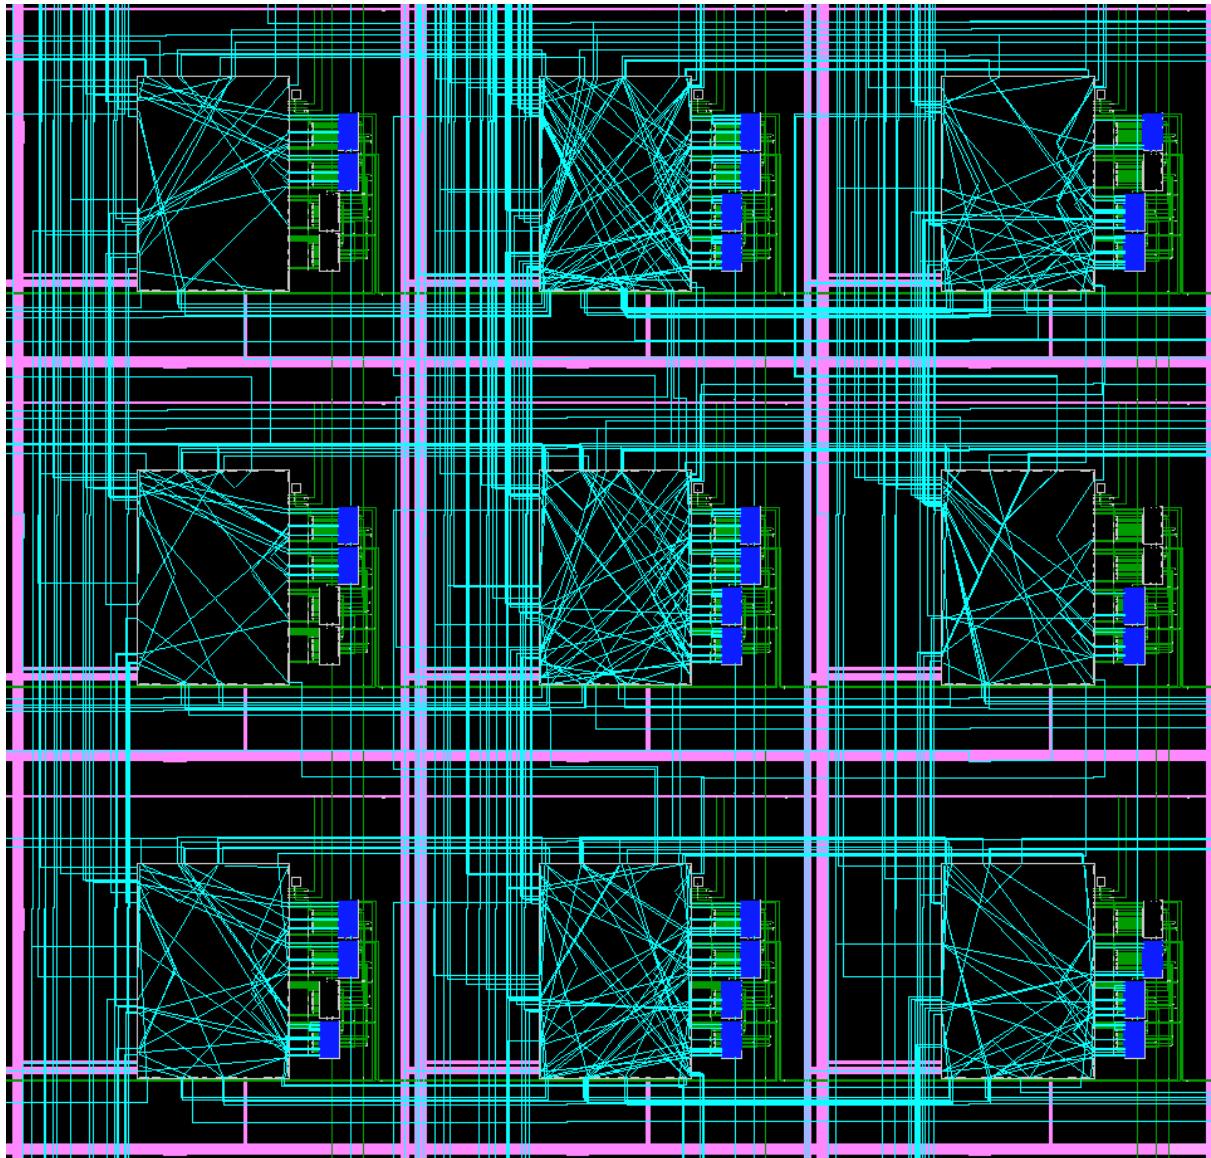


Figure 8.4: Placed and routed sample. Exported as XDL, translated to NCD, and opened with FPGA Editor. Blue lines are routes generated by the system, while green and purple lines are only shown for extra context.

One may also consider that although the base system required a large percentage of the XC2VP30 device, it also provided a full Linux kernel, capable of performing a variety of tasks. The opposite way of looking at the matter is to consider a system that already contains a processor and memory, and to ask how much additional circuitry is required to support autonomy. The likely answer is that not too much would be required beyond the ICAP interface.

There is of course also a performance penalty for implementing designs in FPGAs rather than ASICs, because of the extra layer of configurable logic, but those trade-offs are best discussed by their respective proponents. However, although many designs use FPGAs entirely for static logic, without ever performing runtime reconfiguration, the present work is one case that has no ASIC alternative. ASICs are inherently static—though some are beginning to include blocks of configurable fabric—while this autonomous hardware system is inherently dynamic.

The need to develop CAD tools embedded within the systems that they support has been demonstrated by this work. It is possible to fold the necessary tools into the system, and with some extra robustness and timing data that tool vendors already know how to provide, embedded versions of these tools could be developed, as long as vendors could find a financial justification.

The testing issue remains, but experience with this work suggests that the difficulties may not be quite as large as initially feared. It *is* true that multiple unrelated circuits crowded together into a figurative tight space have a larger chance of interfering with each other. But it also remains true that first-order behavior of circuits depends mostly on their connections, and circuits that don't connect to each other may be able to coexist quite happily. At least some measure of that is borne out by experience with this system, and with Xilinx Embedded Development Kit (EDK) projects: Complex systems can be designed with versatile busses that support multiple masters and slaves, and that arbitrate properly between components. In other words, there are likely to be some semantic rules that can be layered on top of the structure of the system, to help keep things orderly.

And as for the suggestion that hardware and software designers would have to adapt their ways of working and thinking, that too remains a strong likelihood. The author's past experience in switching from procedural programming to object-oriented programming, is testimony to the difficulty of learning entirely new ways of doing things. But the long term benefits of that switch are compelling, and many new programmers these days begin directly with Java or C++, two prominent object-oriented languages. Perhaps then, the corresponding transition for hardware developers may also happen smoothly.

Chapter 9

Conclusion

9.1 Claims

This work has contributed a roadmap for the development of autonomous computing systems—systems that are able to modify their own hardware while running—and has implemented a system that proves the concept and extends the state-of-the-art.

9.2 Narrative

The larger promise of hardware autonomy extends well beyond this initial effort, and yet little of what was accomplished here had been developed or even proposed elsewhere. It is true that some researchers have sought to implement individual parts of this—notably runtime routing and placement work by Vahid, Lysecki, Ma, and of course Keller and the JBits team—but frequently in custom architectures, and almost never with the same intent.

Brebner’s insightful writings are probably those closest in spirit to this work, and Kubisch et al. are the only ones known to have attempted anything comparable. Many other researchers work at the much larger granularity of Networks-on-Chip, or heterogeneous multi-processors. In fairness, many of those approaches will be the first to garner commercial acceptance. But this work aims for systems that control themselves at the finest available granularity, and apologizes only for the fact that such systems cannot yet make *actual* physical changes to themselves.

In keeping with the philosophy of this work, change should happen because something motivates that change: it should be a response with intent, not an arbitrary perturbation. Furthermore, the more complex the system, the more appropriate that it shoulder its own weight, and assume responsibility for its resources and operation. When those responsibilities are transferred to the

system, and the system is able to model and understand itself, it also becomes able to offer a simpler interface to its environment and users.

The larger motivation of this *macroscopic* work is paradoxically to build a bridge to understanding the *microscopic* or fundamental interaction of hardware and information, as explained in the Preface. Sometimes extrema have much in common, as is true in high-energy physics, and the fundamental nature of hardware-information interaction remains a mystery that will hopefully be plied to uncover some of its secrets.

9.3 Roadmap

The proposed roadmap for autonomy has been validated to the extent that it was tested by the implementation phase of this work. The roadmap consists of a hierarchy of levels, each one building upon its predecessors, and collectively providing all of the capabilities necessary for a system to function autonomously.

- Level 0 has no autonomy.
- Level 1 performs primitive slot allocation.
- Level 2 performs full routing, placement, and reconfiguration.
- Level 3 performs synthesis and architecture mapping.
- Level 4 observes itself and its environment.
- Level 5 searches for possible responses to conditions observed.
- Level 6 autonomously implements responses.
- Level 7 grows by combining known information.
- Level 8 learns by evaluating its actions.

The lower levels of the roadmap are grounded in implementation tools, while the upper levels reach as high as artificial intelligence may take them. The middle levels meld sensors, searches, and cost functions. There are no theoretical obstacles to any of these levels, though many of them present significant engineering challenges, particularly by requiring workstation-class tools to run on lean embedded systems.

There are many issues that accompany these levels of autonomy. Policies or guidelines must be put in place to ensure that the system performs useful work, instead of interminably reinventing itself. Internal security practices must be implemented to ensure that the system does not open its gates to malicious circuitry. And proper testing and separation between circuits must be enforced to protect the system's behavior. Each of these is a complex issue in its own right, coming together in even more complex ways within an autonomous system.

Insights gained from the implementation phase and from discussions with other researchers [89], make it clear that the proposed roadmap does not have to be followed exactly. The implementation described in this work is best classified as a Level 6 system, even though it skips Level 3 synthesis and mapping capabilities. Furthermore, others have implemented higher-level sensing and response functions, without extending their systems down to fine-grained hardware levels.

On the basis of the insights gained, a refinement of this roadmap would likely do away with its single linear axis, and replace it with two orthogonal axes: one describing the system's ability to change itself, in granularity as well as high-level abstractness, and the other describing the system's ability to observe and decide how to respond to conditions. A more complete development of these ideas is deferred for future work: The present roadmap has served its purpose well, and can be appropriately refined in due time.

9.4 Implementation

The implementation phase has confirmed the validity of the roadmap, and has demonstrated the first known instance of a system reconfiguring itself—with bitstreams that it generated—while continuing to operate.

The initial proposal promised to demonstrate a Level 2 system, with the intent of providing foundational capability for subsequent work. Level 2 capability was achieved, and feedback received from certain parties drove the work still further.

The underlying Level 2 system provides core functionality: the ability to parse circuits, place them, configure them, route them, connect them, generate bitstreams, and reconfigure itself with those bitstreams. Layered on top of the Level 2 system is a collection of sensors, with a fixed response library, and the ability to implement responses. This constitutes a simple Level 6 system, that is not yet particularly impressive, but that should suffice to convince anybody that hardware autonomy truly is achievable.

Any system of this complexity opens broad doors for further work—foundations are only useful when structures are built on top of them.

The ability of a system to model its own resources is a key component of its ability to assume responsibility for those resources. Knowing which resources are available and unavailable, and knowing the degrees of freedom that each one possesses, is part of what makes autonomy possible. Without this ability, a system could do little more than arbitrarily perturb itself—hardly the intent of this work.

Beneath the modeling and the server functionality, is a custom hardware system built inside the XUP Board's Virtex2P FPGA. A PowerPC immersed inside the FPGA fabric is tied to PLB and OPB busses, memory and ICAP controllers, and a variety of other standard and custom

peripherals. Custom sensor stubs, capable of communicating with the mouse, keyboard, audio, and video devices, serve as the system's eyes. And a sizeable reconfigurable area is available to implement custom hardware cores.

Some restrictions were placed on the dynamically instantiated cores, to avoid state corruption issues with LUT RAMs, including the special case of shift-registers. It would be nice not to have to impose any restrictions, but the restrictions are at least easily specified in Synplify and other synthesis tools, in exchange for some performance and area penalty. Apart from the stated exceptions, this system supports *all* of the logic and wiring inside the FPGA, and not just a uniform array of slices, as is often the case.

Instantiated circuits can also be exported to other tools in the context of the full system. This gives the project offline access to full Design Rule Checks, full timing analysis against the base system's constraints, and full visualization in FPGA Editor.

Accompanying the standard and custom hardware is a Linux kernel and filesystem, along with custom drivers, a Java Runtime Environment, a set of CAD tools targeting the system itself, and a GUI that permits user interaction. The user can draw with a mouse on the FPGA usage display to mask out resources that the system is not permitted to use, and can query the system and directly send commands to it.

9.5 Evaluation

The single most important metric for this system is whether it actually works. The fact that it does is a validation of the proposed roadmap, at least up to or approaching Level 6, and a validation of a complex set of CAD tools, embedded inside the very system that they target.

If the system did not work, it might do nothing at all, but in much greater likelihood, it would spontaneously kill itself as soon as it attempted self-reconfiguration. That results from the fact that the base system touches nearly every configuration frame in the device, not with logic resources, because those were constrained to a specific region, but with routing resources that have to pass through the reserved area. The system did in fact spontaneously kill itself on a number of occasions during development, whenever the base system and the server's configuration bitstream were accidentally left out of sync. Such is the importance of the system maintaining an accurate model of itself.

Beyond the conclusion that the system really does work, there are a some additional ways of evaluating its performance. In a head-to-head comparison with identical circuits, the custom tool set was shown to run at least four times faster than the corresponding Xilinx ISE tools, though the ISE tools of course provide much higher quality results. But the very same tests on the XUP board ran up to 60 times slower, because of the disparity in processing power between the workstation and the embedded system. In absolute terms, it can take 15 minutes to implement a circuit that

contains a few thousand EDIF cells, though the actual duration depends heavily on the placeability of the circuit.

Another metric is the frequency at which the dynamically implemented circuits can run. Although the system has no way of verifying or enforcing timing requirements, it is possible to analyze timing performance offline with the help of the XDL export and the Xilinx ISE tools. Two key circuits, the video capture core and the audio FFT core, need to run at 27 MHz and 12.288 MHz respectively, and were determined by the ISE timing analyzer to run at up to 109 MHz and 48 MHz respectively. But those numbers belie the true situation, explained below:

On one hand, the margin of error could quickly dwindle if placement or routing congestion became big enough problems. On the other hand, the 27 MHz and 12.288 MHz frequencies were adopted because it became clear that the cores might not achieve the required 100 MHz necessary to interface directly with the PLB and OPB busses. The placer and router are simply not yet mature enough to yield much better results, and even a sophisticated router would be severely limited without access to the proper timing data.

Finally, the system can be evaluated in terms of its area overhead. The autonomous infrastructure described here requires an operating system to support a Java Runtime Environment, and a minimum of 75 MB of memory. The operating system in turn requires a processor to run on, along with supporting busses or peripherals. When everything is added up, the base system includes a 512 MB SDRAM memory stick, and uses 7,426 slices. For many smaller embedded systems, that would be a prohibitively expensive, but those smaller systems tend not to be the ones that require autonomous capabilities. In greater likelihood, systems that require autonomy would already include a significant amount of the processing and memory capabilities that were required, and the costs would be amortized in a much more palatable way. Still further amortization would be realizable in the case of systems that contained multiple FPGAs, if it was acceptable to delegate supervisory duties to one of the FPGAs.

9.6 Reflections

A number of statements and arguments about autonomous systems were made in earlier chapters, and it seems fitting to revisit them at this time.

It was argued that autonomy makes systems more flexible and adaptable, able to communicate at higher levels of abstraction, and able to utilize imperfect devices. Designers can thus focus less on system infrastructure and more on component functionality, knowing that their designs will also be more portable.

Flexibility and adaptability: The demonstration system presently has only a limited degree of functional flexibility, because levels 5 and 6 were implemented fairly superficially. Although the sensor stubs can detect mouse, keyboard, audio, or video activity, the system's response

to conditions is still quite simplistic. However, it remains possible to provide completely new circuitry to the system, at the lowest possible granularity, and connect it in arbitrary fashion—a claim that no other existing system can make.

Portability and abstraction: Because the system includes full Level 2 capabilities, Virtex2P circuits can be provided to it in EDIF format, and are therefore portable to any Virtex2P device. Although the system has only been implemented on the XUP board and its XC2VP30, there is no reason why it could not be implemented on a larger part, and remain fully compatible with the EDIF netlists. Had a Level 3 synthesizer and mapper also been developed, it would have been possible to describe circuits directly in HDL, which would have provided compatibility with *any* configurable architecture.

Defect and fault tolerance: The fact that the user can interactively mask or unmask logic resources, demonstrates at a very fine granularity the system’s ability to avoid faulty or defective portions of the device.¹ Although it was impractical to demonstrate the same kind of masking for wiring resources, the router can accept masks in the very same way, *at the granularity of a single wire*. This returns to the economic argument, that if a system can tolerate imperfect devices, the manufacturing yield need no longer be tied to 100% correctness. It would be interesting to know how the yield and pricing might change if large devices could tolerate 10% defects, for example. This is likely to become a pressing issue at nano-scales, where the defect-free yield is identically zero.

Component-based design: The availability of a system infrastructure allows designers to develop circuit cores, with relatively arbitrary interfaces. As long as the necessary signals are available at runtime, there is no need to pre-define standard busses or other interfaces.² Furthermore, the designer may not know or care what kind of system the core will run on, and can instead focus entirely on the circuit they are developing. This is analogous to a software programmer writing a utility and taking the underlying operating system services for granted.

9.7 Future Work

In some ways this work has merely laid a foundation and raised new questions. But *meaningful* new questions—which allow us to peer a little further than before—are an important component of research, and are therefore welcomed.

Because this work was performed independently and without sponsorship, it does not benefit from established ties to other efforts. Some have suggested building ties with the International Conference on Autonomic Computing (ICAC) [57], as a means to gain additional traction in the research community. But ICAC is focused almost entirely on software—using the term *hardware* only to

¹The system is able to avoid resources that have been masked out, but defect detection and fault detection for masking purposes is a matter left to the appropriate experts.

²Though a designer may not need to use standard busses and interfaces, it is possible that they would choose to do so anyway, to simplify handshaking. Or if a designer had two existing but incompatible circuits, he or she could simply develop some glue circuitry and instantiate it dynamically.

denote commodity desktop computers—and is driven more by shorter-term commercial goals than by longer-term research, so the fit may not be ideal.

A much more interesting and fruitful connection might perhaps be established with the *cybernetics* community. Wikipedia describes cybernetics—springing from the work of Norbert Wiener—as “the interdisciplinary study of the structure of complex systems, especially communication processes, control mechanisms and feedback principles . . . closely related to control theory and systems theory”:

Contemporary cybernetics began as an interdisciplinary study connecting the fields of control systems, electrical network theory, mechanical engineering, logic modeling, evolutionary biology and neuroscience in the 1940s. Other fields of study which have influenced or been influenced by cybernetics include game theory, system theory (a mathematical counterpart to cybernetics), psychology (especially neuropsychology, behavioral psychology, and cognitive psychology), and also philosophy, and even architecture.” [90]

Direct continuations of this work depends upon the availability of the Device API, which has already been described as proprietary and unreleased software. Without the Device API, this system would be unable to model its logic and wiring resources, and would be unable to generate configuration bitstreams: Any system that cannot update its own configuration cannot exhibit hardware autonomy.

Unfortunately, there are no alternatives available for the Device API. Various researchers have attempted to reverse-engineer Xilinx configuration bitstreams, and have met with varying degrees of success, but none of the published work does an adequate job of supporting the *complete* device. Restrictions on the Device API are also the reason why the remainder of this work cannot be released to the community, and why it would in any case do them little good.

The Device API targets the Virtex, Virtex-E, Spartan-IIIE, Virtex-II, Virtex-II Pro, and Spartan-3 architectures, but even the most advanced of these is already two generations old, trailing behind the Virtex-4 and Virtex-5 architectures. It is conceivable in time that the restrictions on the Device API may be lifted, but the opportunities for using it before it becomes obsolete are quickly diminishing. For this reason, although the implementation phase of this work might be extended a little bit, it is unlikely that it would grow into a larger effort.

If continued work was possible, it would make most sense to add timing data to the router, and replace its semi-greedy algorithm with a PathFinder derivative [88]. This would allow the system to guarantee timing on any circuits that it implemented—or decline to implement circuits if they failed to meet timing—and would probably provide the single most effective and compelling argument for its adoption in real systems.

One remaining capability that would have been nice to demonstrate is dynamic *re*-placing, to handle cases of resources failing while the system is running. For example, if the user masked out resources

that were already in use by a dynamically instantiated circuit, the autonomous controller would remove and re-implement the “damaged” portion of the circuit. Clever algorithms could probably be identified or developed for this purpose, but at present the system would have to completely remove the circuit, and re-instantiate it as if it were an entirely new circuit.³

9.8 Conclusion

This work has surpassed the goals of the initial research proposal. It has presented and tested a roadmap to guide the development of autonomous computing systems, and it has implemented a demonstration system that proves the viability of hardware autonomy. The resulting system is able to observe its environment, and respond to changing inputs, by implementing and connecting circuitry inside itself, and reconfiguring itself while continuing to operate.

Special thanks are due to Mr. John Rocovich and the Bradley Foundation, to Xilinx Research Labs, to Sun Microsystems, Inc., and to the Virginia Tech Configurable Computing Lab.

³In such cases, there is no expectation that the *re*-placing step will be glitch-free. On the contrary, if a design requires glitch-free operation, it should use Triple-Modular Redundancy (TMR) or other standard approaches. The system’s ability to work around faults should be seen as complementing TMR, rather than replacing it.

Appendix A

Slice Packing

A.1 Introduction

The discrepancy between the slice subcells referenced by the Xilinx Unified Libraries and the actual slice cells that are available in the Virtex-II Pro architecture, is resolved by the mapper of Section 7.4.8. The mapper provides some basic rules to guide the placer, and once the placement has been generated, the full packing is applied to replace all slice subcell references with slice cell references.

A.2 Rules

The checks and decisions made during full packing are as follows:

1. If ORCY is used:
 - (a) If the ORCY input is tied to ground, the SOPEXTSEL mux is set to 0, otherwise it is set to SOPIN.
 - (b) The ORCY output is not allowed to drive both local and general wires.
 - (c) If the ORCY output requires general wiring, the GYMUX mux is set to SOPEXT, and connects it to external slice port Y.
 - (d) If the ORCY output requires local wiring, it is connected to external slice port SOPOUT.
2. If F6MUX is used:
 - (a) Input 0 is tied to slice port FXINB.
 - (b) Input 1 is tied to slice port FXINA.
 - (c) Input S is tied to slice port BY.

- (d) The F6MUX output is not allowed to drive both local and general wires.
 - (e) If the F6MUX output requires general wiring, the GYMUX mux is set to FX, and connects it to external slice port Y. If port Y is already in use by a different net, an error is generated.
 - (f) If the F6MUX output requires local wiring, it is connected to external slice port FX.
3. If F5MUX is used:
- (a) If the 0 input is not driven by the GLUT output, an error is generated.
 - (b) If the 1 input is not driven by the FLUT output, an error is generated.
 - (c) The S input is connected to external slice port BX. If port BX is already in use by a different net, an error is generated.
 - (d) The F5MUX output is not allowed to drive both local and general wires.
 - (e) If the F5MUX output requires general wiring, the FXMUX is set to F5, and connects it to external slice port X. If port X is already in use by a different net, an error is generated.
 - (f) If the F5MUX output requires local wiring, it is connected to external slice port F5.
4. If XORF is used:
- (a) If the LI input is not driven by the FLUT output, an error is generated.
 - (b) The FXMUX mux is set to FXOR, and connects it to external slice port X. If port X is already in use by a different net, an error is generated.
 - (c) If the carry input originates in the previous stage of the carry chain, the CYINIT mux is set to CIN, and connects it to external slice port CIN. Otherwise the CYINIT mux is set to BX, and connects it to external slice port BX.
5. If XORG is used:
- (a) If the LI input is not driven by the GLUT output, an error is generated.
 - (b) The GYMUX mux is set to GXOR, and connects it to external slice port Y. If port Y is already in use by a different net, an error is generated.
6. If CYMUXF is used:
- (a) If the S input is driven by VCC, the CYSELF mux is set to 1.
 - (b) If the S input is driven by the FLUT output, the CYSELF mux is set to F.
 - (c) If the S input is driven by anything else, an error is generated.
 - (d) If the carry input originates in the previous stage of the carry chain, the CYINIT mux is set to CIN, and connects it to external slice port CIN. Otherwise the CYINIT mux is set to BX, and connects it to external slice port BX.
 - (e) The CYMUXF output is not allowed to drive both local and general wires.
 - (f) If the CYMUXF output requires general wiring, the XBMUX mux is set to 1, and connects it to external slice port XB.

- (g) If the DI input is driven by the FAND output, the CY0F mux is set to PROD.
- (h) If the DI input is driven by the FLUT I0 input, the CY0F mux is set to F1.
- (i) If the DI input is driven by the FLUT I1 input, the CY0F mux is set to F2.
- (j) If the DI input is driven by VCC, the CY0F mux is set to 1.
- (k) If the DI input is driven by GND, the CY0F mux is set to 0.
- (l) If the DI input is driven by external slice port BX, the CY0F mux is set to BX.
- (m) If the DI input is driven by anything else, an error is generated.

7. If CYMUXG is used:

- (a) If the S input is driven by VCC, the CYSELG mux is set to 1.
- (b) If the S input is driven by the GLUT output, the CYSELG mux is set to G.
- (c) If the S input is driven by anything else, an error is generated.
- (d) The CYMUXG output is not allowed to drive both local and general wires.
- (e) If the CYMUXG output requires general wiring, the YBMUX mux is set to 1, and connects it to external slice port YB.
- (f) If the DI input is driven by the GAND output, the CY0G mux is set to PROD.
- (g) If the DI input is driven by the GLUT I0 input, the CY0G mux is set to G1.
- (h) If the DI input is driven by the GLUT I1 input, the CY0G mux is set to G2.
- (i) If the DI input is driven by VCC, the CY0G mux is set to 1.
- (j) If the DI input is driven by GND, the CY0G mux is set to 0.
- (k) If the DI input is driven by external slice port BY, the CY0G mux is set to BY.
- (l) If the DI input is driven by anything else, an error is generated.

8. If FLUT is used:

- (a) The A1, A2, A3, and A4 inputs are connected respectively to external slice inputs F1, F2, F3, and F4.
- (b) If the D output requires general wiring, the FXMUX mux is set to F, and connects it to external slice port X. If port X is already in use by a different net, an error is generated.
- (c) If the D output requires local wiring but is tied to an external net, the FXMUX mux is set to F, and connects it to external slice port X. If port X is already in use by a different net, an error is generated.
- (d) If the FLUT primitive represents a LUT, the F attribute is set to LUT.
- (e) If the FLUT primitive represents a ROM, the F attribute is set to ROM.

9. If GLUT is used:

- (a) The A1, A2, A3, and A4 inputs are connected respectively to external slice inputs G1, G2, G3, and G4.

- (b) If the D output requires general wiring, the GYMUX mux is set to G, and connects it to external slice port Y. If port Y is already in use by a different net, an error is generated.
- (c) If the D output requires local wiring but is tied to an external net, the GYMUX mux is set to G, and connects it to external slice port Y. If port Y is already in use by a different net, an error is generated.
- (d) If the GLUT primitive represents a LUT, the G attribute is set to LUT.
- (e) If the GLUT primitive represents a ROM, the G attribute is set to ROM.

10. If FFX is used:

- (a) The Q output is connected to external slice port XQ.
- (b) If the flip-flop is set for negative-edge triggering, the CLKINV mux is set to CLK_B, otherwise it is set to CLK. If the triggering polarity is different from what already defined, an error is generated.
- (c) If the D input is driven by the external slice output X, the DXMUX mux is set to 1, and connects external slice port DX to the input.
- (d) If the D input is not driven by the external slice output X, the DXMUX mux is set to 0, and connects external slice port BX to the input. If port BX is already in use by a different net, an error is generated.
- (e) If the flip-flop uses a clock, the CK input is connected to external slice port CLK. If port CLK is already in use by a different net, an error is generated.
- (f) If the flip-flop uses a chip-enable, the CE input is connected to external slice port CE. If port CE is already in use by a different net, an error is generated.
- (g) If the flip-flop uses a reset, the SYNC_ATTR attribute is set to SYNC and the SR input is connected to external slice port SR. If the SYNC_ATTR attribute is already set to ASYNC or if the SR port is already in use by a different net, an error is generated.
- (h) If the flip-flop uses a clear, the SYNC_ATTR attribute is set to ASYNC and the SR input is connected to external slice port SR. If the SYNC_ATTR attribute is already set to SYNC or if the SR port is already in use by a different net, an error is generated.
- (i) If the flip-flop uses a set, the SYNC_ATTR attribute is set to SYNC and the REV input is connected to external slice port BY. If the SYNC_ATTR attribute is already set to ASYNC or if the BY port is already in use by a different net, an error is generated.
- (j) If the flip-flop uses a preset, the SYNC_ATTR attribute is set to ASYNC and the REV input is connected to external slice port BY. If the SYNC_ATTR attribute is already set to SYNC or if the BY port is already in use by a different net, an error is generated.
- (k) If the flip-flop mode is specified as a latch, the FFX attribute is set to LATCH, otherwise it is set to FF. If the FFX attribute is already set to the opposite setting, an error is generated
- (l) If the flip-flop does not use a reset, the SRINV mux is set to SR_B, otherwise it is set to SR.

(m) The CEINV mux is set to CE.

11. If FFY is used:

- (a) The Q output is connected to external slice port YQ.
- (b) If the flip-flop is set for negative-edge triggering, the CLKINV mux is set to CLK_B, otherwise it is set to CLK. If the triggering polarity is different from what already defined, an error is generated.
- (c) If the D input is driven by the external slice output Y, the DYMUX mux is set to 1, and connects external slice port DY to the input.
- (d) If the D input is not driven by the external slice output Y, the DYMUX mux is set to 0, and connects external slice port BY to the input. If port BY is already in use by a different net, an error is generated.
- (e) If the flip-flop uses a clock, the CK input is connected to external slice port CLK. If port CLK is already in use by a different net, an error is generated.
- (f) If the flip-flop uses a chip-enable, the CE input is connected to external slice port CE. If port CE is already in use by a different net, an error is generated.
- (g) If the flip-flop uses a reset, the SYNC_ATTR attribute is set to SYNC and the SR input is connected to external slice port SR. If the SYNC_ATTR attribute is already set to ASYNC or if the SR port is already in use by a different net, an error is generated.
- (h) If the flip-flop uses a clear, the SYNC_ATTR attribute is set to ASYNC and the SR input is connected to external slice port SR. If the SYNC_ATTR attribute is already set to SYNC or if the SR port is already in use by a different net, an error is generated.
- (i) If the flip-flop uses a set, the SYNC_ATTR attribute is set to SYNC and the REV input is connected to external slice port BY. If the SYNC_ATTR attribute is already set to ASYNC or if the BY port is already in use by a different net, an error is generated.
- (j) If the flip-flop uses a preset, the SYNC_ATTR attribute is set to ASYNC and the REV input is connected to external slice port BY. If the SYNC_ATTR attribute is already set to SYNC or if the BY port is already in use by a different net, an error is generated.
- (k) If the flip-flop mode is specified as a latch, the FFY attribute is set to LATCH, otherwise it is set to FF. If the FFY attribute is already set to the opposite setting, an error is generated
- (l) If the flip-flop does not use a reset, the SRINV mux is set to SR_B, otherwise it is set to SR.
- (m) The CEINV mux is set to CE.

Appendix B

Design Proposal

B.0 Preamble

This appendix is excerpted from the research proposal. It discusses and evaluates candidate target architectures and implementation tasks, and provides insight into design decisions made. Embedded references have been updated where applicable, to match corresponding sections of this dissertation.

B.1 Overview

This chapter considers an approach to implement an autonomous computing system. It quickly becomes apparent that implementation of the Level 8 system of Chapter 4 would require an entire team of engineers, with expertise ranging from software and embedded systems to artificial intelligence. Each part of the system can be realized at present, but the complete system is still too complex to address in the scope of a single dissertation.

B.2 Component Availability

This section considers the components required by a Level 8 system, and comments upon whether each one is available and has been demonstrated in practice or in concept. In many cases the individual components are already well understood, but their inclusion in a constrained system may introduce some additional complications. The major components are as follows:

Internal Configuration Port: Allows the device to change its own configuration. Already available in some FPGAs and in Cell Matrix devices.

Allocation Map: Map of used and available resources, comparable to the File Allocation Table on a disk. Can be implemented in memory; should suit the target architecture, but be flexible enough for a wide range of 2- and 3-dimensional topologies in general. Closely tied to the internal model.

Internal Model: Model that the system maintains of itself. To be defined; must suit the target architecture, but be flexible enough for a wide range of 2- and 3-dimensional topologies in general.

Placer: Tool to assign design logic components to available resources. Well understood and available to standard tool flows. Embedded placement has been conceptually demonstrated in [91]. Incremental placement has been demonstrated in [49]. Research tools are demonstrated in [48].

Router: Tool to connect logic components as specified by the netlist. Well understood and available to standard tool flows. Embedded routers are not known to exist,¹ and the complexity of devices may result in very large memory footprints. Incremental routing has been demonstrated in [22]. Research tools are demonstrated in [48].

Configuration Generator: Tool to generate a configuration bitstream from the placed and routed design. Available but generally proprietary. Few suitable tools are available outside of standard tool flows, with JBits being the notable exception [25], although limited embedded generators have been demonstrated [93], and some independent work has been performed by [94].

Synthesizer: Tool to convert a hierarchical or behavioral design into a netlist of gates or primitive logic. Available but generally proprietary. Open-source Verilog [95] and VHDL [96] synthesizers merit further investigation.

Mapper: Tool to convert generic netlists into netlists compatible with the target architecture. Available but generally proprietary.²

Monitor: Allows the system to track sensor input and/or internal state. To be defined—system specific.

Sensors: Allow the system to observe itself and to observe its environment. To be defined—system specific. One possible model is the central nervous system, along with the five senses.

Prioritizer: Discriminates between simple observations and conditions that require a response, and assigns priorities when faced with conflicting requirements. To be defined—system specific.

¹JBits and its router were demonstrated running on a LEON core inside an FPGA [92], but this does not constitute *self-routing* because the router did not target the device that it was running on.

²Non-proprietary mapping from structural Java circuits has been demonstrated by JHDL [97], but the JHDL flow does not support standard hardware descriptions languages such as VHDL, Verilog, or SystemC.

Response Library: Provides candidate responses to detected conditions. To be populated and configured by the system designer, and supplemented by peer systems or by other resources similar to public libraries.

Evaluator: Determines how well candidate responses resolve the detected condition. To be defined—system specific. Likely requires some basic artificial intelligence.

Simulator: Simulates the effect of applying a candidate response to the system. To be defined—optional component.

Implementer: Directs the selected response through the synthesizer, mapper, placer, router, and configuration generator, and sends it to the internal configuration port.

Requester: Passes response requests along to neighbors, peers, supervisors, or the system designer. To be defined. Requires a standard protocol for interoperability.

Inference Engine: Extracts information from its rules and from the observed conditions, and formulates candidate response. To be defined. Must understand how to assemble basic building blocks into complete behaviors that address the conditions. Probably requires the presence of the simulator component. This is an artificial intelligence component.

Assessor: Analyzes the suitability of the responses applied for detected conditions, and updates the response library accordingly. To be defined. Must track experiences and determine the desirability or undesirability of responses, learning from mistakes and successes. This is an artificial intelligence component.

It is clear that many of these components are not presently available in a form suitable for inclusion within an incrementally changing system. Almost all of the components have been researched in at least one field of study, but have not been brought together into the framework proposed here.

B.3 Proposal

When considering what has already been accomplished by the research community, one finds examples of Level 0 and Level 1 systems described by [58] and [14, 98], respectively. Level 2 systems have been indirectly suggested [45], but have not actually been implemented. Kubisch et al. describe and implement a rudimentary Level 3 system that has practical appeal, but its granularity is very coarse, and they recognize the limitations that accompany that:

“Developments in dynamical reconfiguration considerably depend on currently available technologies. The previous sections already pointed out the need for more fine grained technologies, but these are not yet ready for use or are not yet considered at all. The finer [the] modifications [that] can be applied to a system, the less software processing is

necessary and the less abstraction levels must be traversed. The fewer [the] layers [that] must be traversed, the faster and more flexible the whole process of self reconfiguration will be.” [38]

The system discussed by Kubisch et al. executes the entire design flow, from synthesis to bitstream generation, each time a change is applied, which means that the response time of the system will be on the order of minutes or hours.³ Its granularity depends on a tile abstraction that incorporates a Network-on-Chip with some kind of configurable functionality. But by the arguments of its designers, it would be desirable to implement a finer grained system, perhaps even at the cost of beginning with a lower level system.

Therefore, this work proposes to 1) explore and describe the methodology necessary to realize a Level 8 system, and to 2) demonstrate the viability of autonomous computing by implementing a fine grained Level 2 system that accepts technology-mapped netlists and implements them within itself.

B.4 Platform Selection

The choice of a platform for proof-of-concept implementation depends on a variety of factors. Various Xilinx FPGA architectures are appealing because of the tools or environments that they provide. Novel architectures like the Cell Matrix are also interesting for their very fine configuration granularity and the very unconventional possibilities that they afford. Altera FPGA architectures also provide flexible tool options that make them very attractive for research [99], but appear to lack an internal configuration port and the ability to accept partial configurations [100], both of which are indispensable for even a Level 0 system.

The POEtic chip is another an intriguing research platform [101, 102], with its attached processor, direct access to configuration bits, and support for dynamic routing. While the device appears to be an excellent platform for the study of evolvable hardware, the configurable fabric is intended primarily as a playpen, with the processor and configurable fabric segregated from each other. As a result, while the processor has full control over the configurable fabric, the fabric lacks the ability to configure itself, making it unable to function as a Level 0 system. Other interesting commercial architectures like the Stretch S5000 [103, 104] or the IPFlex DAPDNA [105] are unable to generate their own configurations, and are therefore not included in the platform selection.

Table B.1 shows the pros and cons of the three most promising platforms: the Cell Matrix, the Virtex-II, and the Virtex-II Pro. Although the Cell Matrix is an interesting architecture, it requires a radically different design methodology and programming model. While some of its features are desirable in autonomous computing, they do not directly support the hypothesis presented here.

³Kubisch et al. do also consider caching previously generated bitstreams, but reverting to a previous bitstream does not facilitate implementing new changes to the system.

Table B.1: Platform considerations

<i>Feature</i>	<i>Cell Matrix</i>	<i>Virtex-II</i>	<i>Virtex-II Pro</i>
Configuration Port	✓ yes	✓ yes	✓ yes
Allocation Map	cells	✓ memory	✓ memory
Internal Model	✓ simple	✗ complex	✗ complex
Incremental Placer	✗ not available	simple	simple
Incremental Router	✗ not available	ADB †	ADB †
Configuration Generator	✓ trivial?	✓ JBits	JBits? ‡
Processor	✗ not available	LEON2 ††	✓ PowerPC or LEON2 ††
Operating System	✗ not available	✓ Linux	✓ Linux
Platform	✓ simulator	DN3000K10	✓ ML310 or XUP
Acceptance	✗ research	✓ mainstream	✓ mainstream
Flexibility	✓ excellent	good	good
Configuration Granularity	✓ cell	frame	frame

† The router provided with ADB needs improved support for Virtex-II and Virtex-II Pro.

‡ JBits 3.0 would need to be extended to support Virtex-II Pro.

†† The LEON2 is a soft core that can be implemented in configurable fabric [106].

Furthermore the Cell Matrix lacks any tool support at present, though work has begun on the development of place and route tools [107]. The lack of tools means that there is presently no easy way to implement a soft processor inside a Cell Matrix, nor is there any kind of operating system support. But most importantly, the Cell Matrix is still a proposed architecture that does not yet exist as an actual integrated circuit. These many factors argue against the Cell Matrix architecture as a candidate platform.

The Virtex-II and Virtex-II Pro platforms remain attractive for implementation purposes. The availability of tools like JBits and ADB means that incremental changes are viable, although the ADB router needs some work for both architectures, and JBits would need to be extended to support Virtex-II Pro. However Virtex-II Pro has established its ability to run Linux on its embedded PowerPC processors, and the fact that the Virtex-II Pro is only one generation behind Xilinx's new Virtex-4 flagship is also appealing. Finally, it is easier to find large amounts of memory on Virtex-II Pro boards than on Virtex-II boards, and that is a significant consideration because both JBits and ADB have large memory footprints.

Discussions are currently underway with Xilinx about the possibility of obtaining bitstream support for Virtex-II Pro, as a minor part of this proposed work. If that option materializes, this work will proceed with the Virtex-II Pro, otherwise it will proceed with the Virtex-II.

B.5 Tasks

A number of individual tasks need to be accomplished in order to demonstrate a Level 2 system. The system needs an incremental placer, an incremental router, a framework for these tools, an internal model of itself, and the ability to a Java Virtual Machine (JVM) under Linux. Each of these tasks is considered in turn.

B.5.1 Incremental Placer

The proposed system will be supplied with technology-mapped netlists, so it must be able to place the netlist logic within its available space. Because this work is not primarily focused on placement research, it requires only a basic placer with no sophistication to speak of.

Of particular interest are the facts that the placement will be constrained and tightly coupled to the system's internal model. Clearly, the system can only place circuitry in areas that are not already in use, and this restriction in fact simplifies the placement problem by reducing the search space as the device fills up. Furthermore, because this system will change while it is in operation, one must pay special attention to keeping the internal model synchronized with the system's changing configuration.

A more sophisticated system in the future might have the ability to relocate parts of itself when appropriate: The concept would be analogous to hard drive defragmentation. At this point in the research, however, a very simple placer will do. The only novelty in this placer will be its dynamic interaction with the system model.

Description: Develop a basic placer that is compatible with a bitstream interface.

Time frame: 1 month

B.5.2 Incremental Router

In addition to placing the netlist logic, the system must also make the connections dictated by the netlist, which is to say that it must be able to route the netlist. One aspect of this is simply to make the connections within the circuit described by the netlist. The other aspect is to be able to connect that circuit to the rest of the system, and therefore to be able to extend existing nets in order to connect to the new circuitry.

The ability to route, unroute, and trace configuration bitstreams has already been demonstrated by ADB. However, the ADB router does not perform well on the Virtex-II and Virtex-II Pro architectures. The current router uses as little domain knowledge as possible, and that approach simply does not work very well with the wiring in Virtex-II class devices. ADB's router must therefore be extended to better understand and support the architecture.

A more sophisticated router in the future would have the ability to use and understand timing information. While it wouldn't be difficult to extend ADB to include timing information, the underlying timing data is not readily available from Xilinx.⁴ At present, the only likely way to obtain timing information for ADB would be to first extract it from existing devices through testing, or from Xilinx timing reports through analysis. In addition, a more sophisticated router would have the ability to consider all of a subcircuit's nets at once, and to prioritize routes based on path length and timing requirements. The ADB wiring data could support such a router, but as with placement, routing is merely a necessary part of a Level 2 system, and not a research priority in the present work.

Description: Extend the ADB router to better support Virtex-II class wiring.

Time frame: 1 month

B.5.3 Virtex-II Pro Bitstream Support

This section pertains to Virtex-II Pro bitstream support. It is superseded by the Device API tool developed with Xilinx, and all details have been removed for intellectual property reasons.

B.5.4 System Model

Unlike many computing systems, the system model in this work plays a primary part in the function of the system. In many cases it is common to abstract architectural details away, in order to work with simplified models. In this case it is necessary to provide a very good model of the system and the changes that it undergoes, in order to provide a higher level of abstraction to clients. This is the essence of the work and the reason for shifting more of the complexity into the system itself, so that the outside world can interact with it through a simpler interface.

The system model and the manner in which it interacts with the system and its environment constitute the main contribution of this work. It is the one part that is mostly unexplored by other research, and also the part that opens the door to higher levels of autonomy. Although the implementation may be fairly simple, a reasonable amount of emphasis will also be placed on theory, because the intent is to spur continued research into autonomous computing systems, and such systems will not necessarily share the same architectural properties as the platform used here.

The system *must* understand every part of itself that it is permitted to change, and furthermore, it may need to understand things about its environment in order to properly interact with that environment. This simply affirms that autonomy requires some form of awareness, and that the awareness must account for the changes that the system undergoes.

⁴Part of the timing information is built into the actual tools instead of being available in separate data files as the wiring model is.

What kind of model should one expect? As noted above, the system must know about the parts of itself that may change, so in a general sense it will have to understand its resources and their capabilities. Because the system will undergo virtual changes and not physical changes, its model of total wiring and logic resources will be static, but its model of available resources will be dynamic. In the proposed implementation, ADB already knows about to total and available wiring resources, so part of the model may simply be maintained by ADB. Perhaps the placer and the logic model can be integrated in a similar manner.

While the precise form of the model has not yet been established, one may expect it to encode its information hierarchically when possible. The high degree of regularity in an FPGA means that many of the logic resources are replicated many times in the device, and instead of modeling each resource individually, it makes sense to maintain a single model of a resource type, and simply keep track of its many instances. Tracking the resources that are in use can be accomplished with a compacted bitmap, as is commonly used by ADB when tracking wiring usage.

Hierarchy also makes it possible to achieve a reasonable level granularity without creating models that are too large to fit inside the system. What level of granularity is necessary? This research will hopefully suggest some useful answers to that question. At present JBits has no model of logic or wiring resource usage, while ADB has an exhaustive model of all wiring usage. And yet ADB's model is not comprehensive because it has no timing information for wires. From a routing standpoint, it would clearly be desirable for the timing information to be available. On the other hand, it is probably completely unnecessary to model the FPGA wires as transmission lines with distributed parasitics, though that level of granularity might be appropriate for certain parts of a system or its environment.

*Description: Create a model of logic and wiring resources inside the system (which may include some resources outside of the FPGA itself), and perhaps some model of the system's environment.
Time frame: 2 months*

The two month time frame may seem inordinately long for something that sounds fairly simple, but because the modeling is the one part that really has not been studied yet, and because it constitutes the main contribution of this work, it deserves proper attention.

B.5.5 Platform Support

The target platform will hopefully be a Virtex-II Pro board running JBits, ADB, and other tools on a Java Virtual Machine under Linux. The ability to run Linux has already been demonstrated on the ML310 [108] and XUP [24] boards, and something comparable should be possible on a DN3000K10 [109] board with a LEON2 [106] core if JBits cannot be extended for Virtex-II Pro. The non-Java parts of the system will have to be compiled for the target platform. The overall development will necessitate the use of the Xilinx Embedded Development Kit (EDK), and some knowledge of Xilinx bus architectures and protocols.

Description: Run JBits, ADB, a JVM, and related tools under Linux on an ML310 or XUP or DN3000K10 development board.

Time frame: 2 months

The two month time frame may seem inordinately long for something that has already been demonstrated within the Configurable Computing Lab and elsewhere, but this author has essentially no experience running Linux on an embedded system, and conversely is well aware of numerous problems that other Xilinx EDK users have encountered when attempting to do so. The estimate is intended to be conservative, and to allow for unforeseen difficulties, and since the actual duration will likely be shorter, the remaining time should be usable for other tasks.

B.5.6 Complete System

The complete system will require some additional functionality, including the ability to parse netlists and coordinate with the model, the placer, and the router. Many of the details are difficult to foresee at present, but they will undoubtedly arise in the course of the work.

This task completes the demonstration of a Level 2 system, that is able to accept technology-mapped netlists and implement them inside itself, while updating its internal model with each incremental change.

Description: Combine all of the components and any auxiliary pieces into a complete Level 2 system.
Time frame: 1 month

B.6 Optional Tasks

The tasks presented in the previous section constitute all that is being officially proposed. This section describes some additional work that may be pursued if time allows and if the tasks seem promising.

The author expects to have roughly two months to pursue one of these tasks, but no commitment is offered in that respect.

It is also possible that some of the tasks described above will present intriguing opportunities, and that any remaining time will be devoted to those things instead of the optional tasks described here.

Application: This task would use a Level 2 system to demonstrate some useful application, and not simply a proof-of-concept platform to accept netlists.

Synthesis: This task would attempt to employ the Icarus Verilog synthesizer [95] or the Alliance VLSI CAD synthesizer [96] or some similar tool to build a Level 3 system on top of the Level

2 system. A Level 3 system could accept abstract behavioral descriptions and not simply fully mapped netlists.

Timing-Driven Routing: This task would supplement or extend ADB with timing information in order to enable timing-driven routing. This is particularly important in cases where existing routes that are known to meet timing requirements are modified. Without timing-driven routing, it is impossible to guarantee that the modified system will continue to meet timing requirements.

Bibliography

- [1] N. Steiner and P. Athanas, "Hardware-software interaction: Preliminary observations," in *Proceedings of the 12th Reconfigurable Architectures Workshop, RAW 2005, (Denver, Colorado), April 4–5*, (Los Alamitos, California), IEEE Computer Society, 2005.
- [2] W. D. Hillis, "New computer architectures and their relationship to physics or why computer science is no good," *International Journal of Theoretical Physics*, vol. 21, no. 3/4, pp. 255–262, 1982.
- [3] J. D. Bekenstein, "Black holes and information theory," *Journal of Contemporary Physics*, vol. 45, pp. 31–43, 2003.
- [4] R. Landauer, "Information is physical," *Physics Today*, vol. 44, pp. 23–29, May 1991.
- [5] J. A. Wheeler, "Information, physics, quantum: The search for links," in *Complexity, Entropy and the Physics of Information* (W. H. Zurek, ed.), pp. 3–28, Redwood City, California: Addison Wesley Publishing Company, 1990. *Proceedings of the 1988 Workshop on Complexity, Entropy, and the Physics of Information held May–June, 1989, in Santa Fe, New Mexico*.
- [6] D. Minic. (Associate Professor of Physics, Virginia Tech), personal communication, August 2005.
- [7] Z. Toroczkai. (Deputy Directory, Center for Nonlinear Studies, Los Alamos National Laboratory), personal communication, July 2005.
- [8] S. Goldstein, M. Budiu, M. Mishra, and G. Venkataramani, "Reconfigurable computing and electronic nanotechnology," in *Proceedings of the 14th IEEE International Conference on Application-Specific Systems, Architectures and Processors, ASAP 2003, (The Hague, Netherlands), June 24–26*, (Los Alamitos, California), IEEE Computer Society, 2003.
- [9] A. G. Ganek and T. A. Corbi, "The dawning of the autonomic computing era," *IBM Systems Journal*, vol. 42, no. 1, 2003.
- [10] DARPA Information Processing Technology Office (IPTO). <http://www.darpa.gov/ipto>.
- [11] N. Steiner and P. Athanas, "Autonomous computing systems: A proposed roadmap," in *Proceedings of the 2007 International Conference on Engineering of Reconfigurable Systems and Algorithms, ERSA 2007, (Las Vegas, Nevada)*, June 25–28, 2007.

- [12] C. Kao, "Benefits of partial reconfiguration," *Xcell*, vol. 4Q05, no. 55, pp. 65–67, 2005.
- [13] G. Brebner, "The swappable logic unit: A paradigm for virtual hardware," in *Proceedings of the 5th Annual IEEE Symposium on Field-Programmable Custom Computing Machines, FCCM 1997, (Napa, California), April 16–18*, (Los Alamitos, California), IEEE Computer Society, 1997.
- [14] K. Baskaran and T. Srikanthan, "A hardware operating system based approach for run-time reconfigurable platform of embedded devices," in *Proceedings of the Sixth Real-Time Linux Workshop, RTL Workshop 2004, (Singapore), November 3–5*, 2004.
- [15] W. D. Hillis, *The Connection Machine*. Cambridge, Massachusetts: MIT Press, 1985.
- [16] D. D. Gajski, N. D. Dutt, A. C.-H. Wu, and S. Y.-L. Lin, *High-Level Synthesis: Introduction to Chip and System Design*. Boston, Massachusetts: Kluwer Academic Publishers, 1992.
- [17] Xilinx, San Jose, California, *Xilinx Introduces Breakthrough Virtex-II Pro FPGAs to Enable a New Era of Programmable System Design*. Press Release, March 4, 2002, http://www.xilinx.com/prs_rls/silicon_vir/0204_v2p_main.html (archive).
- [18] Xilinx, Inc., *Virtex-II Pro and Virtex-II Pro X Platform FPGAs: Complete Data Sheet*, March 2007. <http://direct.xilinx.com/bvdocs/publications/ds083.pdf> (archive).
- [19] Xilinx, Inc., *Virtex-II Platform FPGA Handbook*, December 2001. Revision 1.3.
- [20] Xilinx, Inc., *Libraries Guide*. ISE 8.1i, <http://toolbox.xilinx.com/docsan/xilinx8/books/docs/lib/lib.pdf>.
- [21] comp.arch.fpga, *for all those who believe in (structured) ASICs....* Posted Mar 24 2006, 10:45 am by rickman: <http://www.fpga-faq.com/archives/99425.html#99444>.
- [22] N. J. Steiner, "A standalone wire database for routing and tracing in Xilinx Virtex, Virtex-E, and Virtex-II FPGAs," Master's thesis, Virginia Tech, August 2002. <http://scholar.lib.vt.edu/theses/available/etd-09112002-143335> (archive).
- [23] Xilinx, Inc., San Jose, California, *Virtex-II Pro Development System: Hardware Reference Manual*, March 2005. http://www.digilentinc.com/Data/Products/XUPV2P/XUPV2P_User_Guide.pdf (archive).
- [24] Digilent, Inc., Pullman, Washington, *Virtex-II Pro Development System*. <http://www.digilentinc.com/Products/Detail.cfm?Prod=XUPV2P>.
- [25] S. Guccione, D. Levi, and P. Sundararajan, "JBits: Java based interface for reconfigurable computing," in *Proceedings of the Second Annual Military and Aerospace Applications of Programmable Devices and Technologies Conference, MAPLD 1999, (Laurel, Maryland), September 28–30*, 1999.
- [26] Linux Kernel. <http://www.kernel.org>.

- [27] *uClinux*. <http://www.uclinux.org>.
- [28] *MontaVista Linux*. <http://www.mvista.com>.
- [29] University of Washington (Embedded Research Group), *EMPART Project*. http://www.cs.washington.edu/research/lis/empart/xup_ppc_linux.shtml (archive).
- [30] Brigham Young University (Configurable Computing Laboratory), *Linux on FPGA Project*. <http://splish.ee.byu.edu/projects/LinuxFPGA/configuring.htm> (archive).
- [31] University of Illinois Urbana-Champaign (Soft Systems Lab), *Building a Xilinx ML300/ML310 Linux Kernel*. <http://www.crhc.uiuc.edu/IMPACT/gsrc/hardwarelab/docs/kernel-HOWTO.html> (archive).
- [32] University of California Berkeley (Berkeley Wireless Research Center), *Building a Linux Kernel for BEE2*. <http://bee2.eecs.berkeley.edu/wiki/BEE2wiki.html> (archive).
- [33] W. Klingauf, *Linux on the Xilinx ML300*. www.klingauf.de. <http://www.klingauf.de/v2p/index.phtml> (archive).
- [34] *BusyBox*. <http://www.busybox.net>.
- [35] Dan Kegel, *crosstool*. <http://kegel.com/crosstool>.
- [36] University of York (Real-Time Systems Research Group), *DEMOS Project: Digitally Enhanced Micro Operating System*. <http://www.cs.york.ac.uk/rtslab/demos/amos/xupv2pro>.
- [37] S. Kubisch, R. Hecht, and D. Timmermann, “Design flow on a chip - an evolvable HW/SW platform,” in *Proceedings of the 2nd IEEE International Conference on Autonomic Computing, ICAC 2005, (Seattle, Washington), June 13–16*, (Los Alamitos, California), pp. 393–394, IEEE Computer Society, 2005.
- [38] S. Kubisch, R. Hecht, and D. Timmermann, “Design flow on a chip - an evolvable HW/SW platform.” *Unpublished full version of the ICAC 2005 poster of reference [37]*, June 2005.
- [39] S. Kubisch. (Dipl.-Ing., University of Rostock), personal communication, October 2005.
- [40] R. Hecht, S. Kubisch, A. Herrholtz, and D. Timmermann, “Dynamic reconfiguration with hardwired networks-on-chip on future FPGAs,” in *Proceedings of the 15th International Conference on Field Programmable Logic and Applications, FPL 2005, (Tampere, Finland), August 24–26* (T. Rissa, S. Wilton, and P. Leong, eds.), IEEE, 2005.
- [41] R. Lysecky, F. Vahid, and S. X.-D. Tan, “A study of the scalability of on-chip routing for just-in-time FPGA compilation,” in *Proceedings of the 13th IEEE Symposium on Field-Programmable Custom Computing Machines, FCCM 2005, (Napa, California), April 18–20*, (Los Alamitos, California), pp. 57–62, IEEE Computer Society, 2005.

- [42] Y. Ha, R. Hipik, S. Vernalde, D. Verkest, M. Engels, R. Lauwereins, and H. De Man, “Adding hardware support to the hotspot virtual machine for domain specific applications,” in *Proceedings of the 12th International Conference on Field-Programmable Logic and Applications, FPL 2002, (Montpellier, France), September 2-4, 2002* (M. Glesner, P. Zipf, and M. Renovell, eds.), vol. 2438 of *Lecture Notes in Computer Science*, pp. 1135–1138, Springer, 2002.
- [43] G. B. Wigley, *An Operating System for Reconfigurable Computing*. Ph.D. thesis, University of South Australia, April 2005.
- [44] G. Brebner, “A virtual hardware operating system for the xilinx XC6200,” in *Proceedings of the 6th International Workshop on Field-Programmable Logic, Smart Applications, New Paradigms and Compilers, FPL 1996, (Darmstadt, Germany), September 25–25* (R. W. Hartenstein and M. Glesner, eds.), vol. 1142 of *Lecture Notes in Computer Science*, pp. 327–336, Springer Verlag, 1996.
- [45] G. Brebner, “Circlets: Circuits as applets,” in *Proceedings of the 6th Annual IEEE Symposium on Field-Programmable Custom Computing Machines, FCCM 1998 (Napa, California), April 15–17*, (Los Alamitos, California), IEEE Computer Society, 1998.
- [46] N. Sherwani, *Algorithms for VLSI Physical Design Automation*. Boston: Kluwer Academic Publishers, third ed., 1999.
- [47] E. Keller, “JRoute: A Run-Time Routing API for FPGA Hardware,” in *Parallel and Distributed Processing, Lecture Notes in Computer Science* (J. Rolim *et al.*, eds.), vol. 1800, (Berlin), pp. 874–881, Springer-Verlag, May 2000.
- [48] V. Betz and J. Rose, “VPR: A new packing, placement and routing tool for FPGA research,” in *Proceedings of the 7th International Workshop on Field-Programmable Logic and Applications, FPL 1997, (London), September 1–3* (W. Luk, P. Y. K. Cheung, and M. Glesner, eds.), vol. 1304 of *Lecture Notes in Computer Science*, pp. 213–222, Springer Verlag, 1997.
- [49] J. Ma, *Incremental Design Techniques with Non-Preemptive Refinement for Million-Gate FPGAs*. PhD thesis, Virginia Tech, January 2003.
- [50] A. Poetter, J. Hunter, C. Patterson, P. M. Athanas, B. E. Nelson, and N. Steiner, “JHDLBits: The merging of two worlds,” in *Proceedings of the 14th International Conference on Field Programmable Logic and Applications, FPL 2004, (Leuven, Belgium), August 30–September 1* (J. Becker, M. Platzner, and S. Vernalde, eds.), vol. 3203 of *Lecture Notes in Computer Science*, pp. 414–423, Springer, 2004.
- [51] P. Bellows and B. Hutchings, “JHDL - an HDL for reconfigurable systems,” in *IEEE Symposium on FPGAs for Custom Computing Machines* (K. L. Pocek and J. Arnold, eds.), (Los Alamitos, CA), pp. 175–184, IEEE Computer Society Press, 1998.
- [52] C. E. Shannon, “A mathematical theory of communication,” *Bell System Technical Journal*, vol. 27, pp. 379–423 and 623–656, July and October 1948.

- [53] Cell Matrix Corporation. <http://www.cellmatrix.com>.
- [54] N. J. Macias. (Founder, Cell Matrix Corporation), personal communication, April 2008.
- [55] Carnegie Mellon University, *Claytronics*. <http://www.cs.cmu.edu/~claytronics> (archive).
- [56] C. A. Mead, *Collective Electrodynamics: quantum foundations of electromagnetism*. Cambridge, Massachusetts: MIT Press, 2000.
- [57] International Conference on Autonomic Computing. <http://www.autonomic-conference.org>.
- [58] W. Westfeldt, “New internet reconfigurable logic for creating web-enabled devices,” *XCell*, vol. 1Q99, no. 31, pp. 6–7, 1999.
- [59] Digilent, Inc., Pullman, Washington, *VDEC1 Video Decoder Board*. <http://www.digilentinc.com/Products/Detail.cfm?Prod=VDEC1>.
- [60] opencores.org, *PS2 interface*. <http://www.opencores.org/projects.cgi/web/ps2>.
- [61] opencores.org, *I2C controller core*. <http://www.opencores.org/projects.cgi/web/i2c>.
- [62] A. Chapweske, *The PS/2 Mouse/Keyboard Protocol*. computer-engineering.org. <http://www.computer-engineering.org/ps2protocol> (archive).
- [63] opencores.org, *Radix 4 Complex FFT*. <http://www.opencores.org/projects.cgi/web/cfft/overview>.
- [64] *GCC, the GNU Compiler Collection*. <http://gcc.gnu.org>.
- [65] *OpenSSL*. <http://www.openssl.org>.
- [66] *OpenSSH*. <http://www.openssh.com>.
- [67] *zlib*. <http://www.zlib.net>.
- [68] G. Kroah-Hartman, *udev and devfs — The final word*. Discusses the switch from devfs to udev in the Linux tree. http://www.kernel.org/pub/linux/utils/kernel/hotplug/udev_vs_devfs (archive).
- [69] X. Calbet, “Writing device drivers in Linux: A brief tutorial,” *Free Software Magazine*, April 2006. http://www.freesoftwaremagazine.com/articles/drivers_linux (archive).
- [70] N. Steiner, *Linux 2.4.26 Virtex-II Pro ICAP Driver*. <http://www.ccm.ece.vt.edu/~nsteiner/icap>.
- [71] R. Schoenherr, A. Thamarakuzhi, and P. Li. (Students at University of Reading, University of Connecticut, and Chinese Academy of Sciences), personal communications, 2007–2008.

- [72] comp.arch.fpga, *Correction for hwicap_v1_00_a code*. Posted Sat, 20 Jan 2007 17:15:36 -0500 by Neil Steiner. http://groups.google.com/group/comp.arch.fpga/browse_thread/thread/02e00d4ec0ef783e (archive).
- [73] M. Shelburne. (M.S. Student, Virginia Tech), personal communication, November 2006.
- [74] Sun Microsystems, Inc., *Java*. <http://java.sun.com>.
- [75] M. Barr and J. Steinhorn, “Kaffe, Anyone? Implementing a Java Virtual Machine,” *Embedded Systems Programming*, pp. 34–36, February 1998.
- [76] *Kaffe.org*. <http://www.kaffe.org>.
- [77] The GNU Project, *The GNU Compiler for the Java Programming Language*. <http://gcc.gnu.org/java>.
- [78] The GNU Project, *GNU Classpath*. <http://www.gnu.org/software/classpath>.
- [79] Sun Microsystems, Inc., *Java SE for Embedded*. <http://java.sun.com/javase/embedded>.
- [80] J. Ocheltree, “Trust Your Ears: The Drum Tech Explorations of Jeff Ocheltree.” DVD. Rittor Music, 2004.
- [81] Brigham Young University (Configurable Computing Laboratory), *Java EDIF Home Page*. <http://reliability.ee.byu.edu/edif> (archive).
- [82] Brigham Young University (Configurable Computing Laboratory), *BYUCC Edif Parser*. http://splish.ee.byu.edu/download/edifparser/edif_parser_download.html (archive).
- [83] UCLA (UCLA Compilers Group), *Java Tree Builder*. <http://compilers.cs.ucla.edu/jtb> (archive).
- [84] dev.java.net, *Java Compiler Compiler (JavaCC)*. <https://javacc.dev.java.net>.
- [85] C. Sechen and A. Sangiovanni-Vincentelli, “The TimberWolf placement and routing package,” *IEEE Journal of Solid-State Circuits*, vol. 20, pp. 510–522, April 1985.
- [86] N. Sherwani, *Algorithms for VLSI Physical Design Automation*. Boston: Kluwer Academic Publishers, second ed., 1995.
- [87] Synplicity, Inc., *Synplify Pro FPGA Solution Datasheet*. <http://www.synplicity.com>.
- [88] L. McMurchie and C. Ebeling, “PathFinder: A negotiation-based performance-driven router for FPGAs,” in *Proceedings of the 1995 ACM 3rd International Symposium on Field-Programmable Gate Arrays, FPGA 1995, (Monterey, California), February 12–14*, (New York), pp. 111–117, ACM, 1995.

- [89] M. French and E. Anderson. (Researchers, USC Information Sciences Institute), personal communication, November 2007.
- [90] Wikipedia, *Cybernetics*. <http://en.wikipedia.org/wiki/Cybernetics>.
- [91] N. J. Macias and L. J. K. Durbeck, "Self-assembling circuits with autonomous fault handling," in *Proceedings of the 2002 NASA/DoD Conference on Evolvable Hardware, EH 2002 (Alexandria, Virginia), July 15–18*, 2002.
- [92] S. A. Leon, "A self-reconfiguring platform for embedded systems," Master's thesis, Virginia Tech, August 2001.
- [93] B. Blodget, P. James-Roxby, E. Keller, S. McMillan, and P. Sundrarajan, "A self-reconfiguring platform," in *Proceedings of the 13th International Conference on Field Programmable Logic and Applications, FPL 2003, (Lisbon, Portugal), September 1–3* (P. Y. K. Cheung, G. A. Constantinides, and J. T. de Sousa, eds.), vol. 2778 of *Lecture Notes in Computer Science*, (Berlin), pp. 565–574, Springer Verlag, 2003.
- [94] N. Franklin, "VirtexTools FPGA programming and debugging tools project." <http://neil.franklin.ch/Projects/VirtexTools>.
- [95] M. Baxter, "Open source in electronic design automation," *Linux Journal*, February 2001.
- [96] LIP6, Université Pierre et Marie Curie, "Alliance VLSI CAD System." <http://www-asim.lip6.fr/recherche/alliance> (archive).
- [97] B. L. Hutchings, P. Bellows, J. Hawkins, K. S. Hemmert, B. E. Nelson, and M. Rytting, "A CAD suite for high-performance FPGA design," in *Proceedings of the 7th IEEE Symposium on Field-Programmable Custom Computing Machines, FCCM 1999, (Napa, California), April 21–23*, (Los Alamitos, California), pp. 12–24, IEEE Computer Society, 1999.
- [98] C. Patterson, "A dynamic module server for embedded platform FPGAs," in *Proceedings of the International Conference on Engineering of Reconfigurable Systems and Algorithms, ERSA 2003, (Las Vegas, Nevada), June 23–26* (T. P. Plaks, ed.), pp. 31–40, CSREA Press, 2003.
- [99] S. Malhotra, T. P. Borer, D. P. Singh, and S. D. Brown, "The quartus university interface program: Enabling advanced FPGA research," in *Proceedings of the IEEE International Conference on Field-Programmable Technology, FPT 2004, (Brisbane, Australia), December 6–8*, pp. 225–230, IEEE, 2004.
- [100] Altera Corporation, San Jose, California, *Stratix II Device Handbook*, March 2005.
- [101] Y. Thoma and E. Sanchez, "A reconfigurable chip for evolvable hardware," in *Proceedings of the 2004 Genetic and Evolutionary Computation Conference, Part I, GECCO 2004, (Seattle, Washington), June 26–30* (K. Deb *et al.*, eds.), vol. 3102 of *Lecture Notes in Computer Science*, (Berlin), pp. 816–827, Springer Verlag, 2004.

- [102] Y. Thoma and E. Sanchez, "An adaptive FPGA and its distributed routing," in *Proceedings of the Reconfigurable Communication-centric SoCs Workshop, ReCoSoC 2005, (Montpellier, France)*, pp. 43–51, June 2005.
- [103] Stretch, Inc., Mountain View, California, *S5500 Data Sheet, Version 1.2*, September 2005.
- [104] N. Flaherty, "Reconfigurable, go figure," *Electronic Product Design, August 14*, no. 40814, 2004.
- [105] T. Sato, "DAPDNA-2: A dynamically reconfigurable processor with 376 32-bit processing elements," in *2005 IEEE Hot Chips XVII Symposium (HCS), HOTCHIPS 17, (Stanford, California), August 14–16*, 2005.
- [106] Gaisler Research, *LEON2 Processor User's Manual, Version 1.0.30*, July 2005.
- [107] N. J. Macias. (Founder, Cell Matrix Corporation), personal communication, September–October 2005.
- [108] Xilinx, Inc., San Jose, California, *ML310 User Guide: Virtex-II Pro Embedded Development Platform*, January 2005.
- [109] The Dini Group, La Jolla, California, *DN3000K10 User's Manual Version 1.65*, June 2003.